

## Exercise #3

The aim of this exercise is to implement system calls and invoke it from user-space.

### Task 1

Write a new system call and a user-space process that invokes it. Test correctness and when invoking the system call, it should print pid (process ID) of the invoking process.

**Step 1.** Add the function of system call definition at [/usr/src/linux-4.15.1/kernel](#)

```
#include <linux/linkage.h>
#include <linux/kernel.h>
#include <linux/sched.h>

asmlinkage long my_syscall(void)
{
    struct task_struct *process;

    process = current;

    printk(KERN_EMERG "Hello World! \n");
    printk("** PID_Number: %d ** \n", process->pid);
    return 0;
}
```

Figure 1. my\_syscall.c

**Step 2.** Add the function prototype in the header file in the file [usr/src/linux-4.15.1/include/linux/syscalls.h](#)

```
asmlinkage long sys_mlock2(unsigned long start, size_t len, int flags);
asmlinkage long sys_pkey_mprotect(unsigned long start, size_t len,
                                   unsigned long prot, int pkey);
asmlinkage long sys_pkey_alloc(unsigned long flags, unsigned long init_val);
asmlinkage long sys_pkey_free(int pkey);
asmlinkage long sys_statx(int dfd, const char __user *path, unsigned flags,
                           unsigned mask, struct statx __user *buffer);
asmlinkage long my_syscall(void);
~
#endif
```

Figure 2. syscalls.h

**Step 3.** Create an entry in system call table in the file

[/usr/src/linux-4.15.1/arch/x86/entry/syscalls/syscall\\_64.tbl](#)

```

329    common    pkey_mprotect      sys_pkey_mprotect
330    common    pkey_alloc         sys_pkey_alloc
331    common    pkey_free          sys_pkey_free
332    common    statx              sys_statx
333    common    my_syscall_hello    my_syscall
#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation.
#
512    x32      rt_sigaction        compat_sys_rt_sigaction
513    x32      rt_sigreturn        sys32_x32_rt_sigreturn
514    x32      ioctl               compat_sys_ioctl
515    x32      readv               compat_sys_readv

```

Figure 3. syscall\_64.tbl

The number of 333 is going to be number of our own system call. 'common' means that the system call can run on both 64 and 32 bits' architecture.

**Step 4.** Add the Makefile for compiling our system call function at

[/usr/src/linux-4.15.1/Makefile](#)

```

obj-y    = fork.o exec_domain.o panic.o \
          cpu.o exit.o softirq.o resource.o \
          sysctl.o sysctl_binary.o capability.o ptrace.o user.o \
          signal.o sys.o umh.o workqueue.o pid.o task_work.o \
          extable.o params.o \
          kthread.o sys_ni.o nsproxy.o \
          notifier.o ksysfs.o cred.o reboot.o \
          async.o range.o smpboot.o ucount.o my_syscall.o

```

Figure 4. Makefile (/kernel)

**Step 5.** To create our own kernel images, change the EXTRAVERSION in the main Makefile at [/usr/src/linux-4.15.1](#)

```

# SPDX-License-Identifier: GPL-2.0
VERSION = 4
PATCHLEVEL = 15
SUBLEVEL = 1
EXTRAVERSION = .ownsyscall
NAME = Fearless Coyote

```

Figure 5. Makefile

**Step 6.** Compile the function and build & install the modules

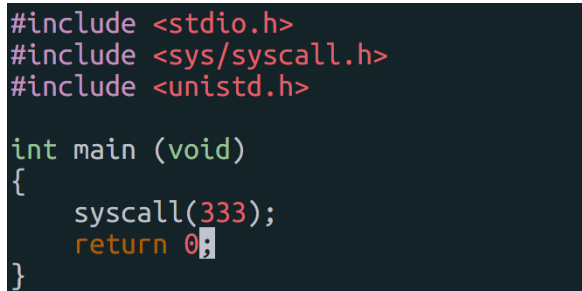
```
sudo make -j 24
sudo make modules -j 24
sudo make modules_install -j 24
sudo make install -j 24
```

-j can compile and build by using several cores. -j 24 means that using 24 CPU cores. Before use this option, check the number of core using `cat /proc/cpuinfo`

**Step 7.** Reboot with our own imange.

```
sudo reboot
```

**Step 8.** Write a simple C application (user level) for calling our own system call.

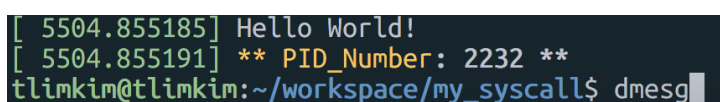


```
#include <stdio.h>
#include <sys/syscall.h>
#include <unistd.h>

int main (void)
{
    syscall(333);
    return 0;
}
```

Figure 6. User level C source file

**Step 9.** Compile by this source file by using gcc and check the result of our own system call by 'dmesg'



```
[ 5504.855185] Hello World!
[ 5504.855191] ** PID_Number: 2232 **
tlimkim@tlimkim:~/workspace/my_syscall$ dmesg
```

Figure 7. dmesg result

## Task 2

Write a Linux kernel module to intercept this new system call and alter its functionality. When the module is removed, the original system call functionality should be restored.

**Step 1.** Write the module source file which works for intercepting our new system call.

```
#include <asm/pgtable.h>
#include <linux/moduleparam.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/unistd.h>
#include <linux/semaphore.h>
#include <asm/cacheflush.h>
#include <asm/set_memory.h> // declared set_memory_rw & set_memory_ro
#include <linux/kallsyms.h>

MODULE_LICENSE("GPL");

void ** sys_call_table;

asmlinkage int (*original_call) (const char*, int, int);

asmlinkage int our_call (const char* file, int flags, int mode)
{
    printk("Intercepted my_syscall \n");
    return 0;
    //return original_call(file, flags, mode);
}

void set_addr_rw(unsigned long addr)
{
    unsigned int level;
    pte_t *pte = lookup_address(addr, &level);

    if (pte->pte & ~_PAGE_RW)
        pte->pte |= _PAGE_RW;
}

static int __init intercept_entry (void)
{
    printk(KERN_ALERT "Module Intercept Inserted \n");

    sys_call_table = (void*)kallsyms_lookup_name("sys_call_table");

    original_call = sys_call_table[333];

    //set_memory_rw((long unsigned int)sys_call_table, 1);
    set_addr_rw(sys_call_table);

    sys_call_table[333] = our_call;
    return 0;
}
```

```
static int __init intercept_entry (void)
{
    printk(KERN_ALERT "Module Intercept Inserted \n");

    sys_call_table = (void*)kallsyms_lookup_name("sys_call_table");
    original_call = sys_call_table[333];

    //set_memory_rw((long unsigned int)sys_call_table, 1);
    set_addr_rw(sys_call_table);

    sys_call_table[333] = our_call;
    return 0;
}

static void __exit intercept_exit (void)
{
    sys_call_table[333] = original_call;

    printk(KERN_ALERT "Module Intercept Removed \n");
}

module_init(intercept_entry);
module_exit(intercept_exit);
```

### 1. '**\_\_init intercept\_entry**'

This function is for starting the module that works main feature. By printing the messages which this module is inserted, it assigned address of "sys\_call\_table" to the variable 'sys\_call\_table'. Our own system call's number is 333, so we assign original system call to variable 'original\_call'. Next, to change this call sys\_call\_table[333] to 'our\_call', we should change the authority of read and write. There are 'set\_memory\_rw' function that change this authority, but there are pointer bugs so create new function named 'set\_addr\_rw'. By using this function, we assign 'our\_call' to 'sys\_call\_table[333]'.

### 2. '**our\_call**'

This function works for intercepting original system call. It returns 0, so the original system call can be ignored by this function.

### 3. '**set\_addr\_rw**'

This function is to find the page table entry for the address and set it writable. It uses 'lookup\_address' function that is find the page table entry for a virtual address. It returns the entry and the level of mapping. Then we change it to writable to assign 'our\_call'.

#### 4. '\_\_exit intercept\_exit'

This function works while we remove module from kernel. When we remove the module, the original system call should do their original function. So we assign 'original\_call' to 'sys\_call\_table[333]'.

Step 2. Create Makefile for compiling the source code

```
obj-m += intercept_syscall.o

all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

~
~
```

Figure 8. Makefile

Step 3. Compile the module and insert the module on kernel

```
sudo make
sudo insmod intercept_syscall.ko
```

Step 4. Check the result of intercepting system call.

```
[ 34.608267] intercept_syscall: loading out-of-tree module taints kernel.
[ 34.608701] intercept_syscall: module verification failed: signature and/or
integrity kernel
[ 34.611601] Module Intercept Inserted
```

Figure 9. Intercepting Module Inserted

```
[ 34.608267] intercept_syscall: loading out-of-tree module taints kernel.
[ 34.608701] intercept_syscall: module verification failed: signature and/or
integrity kernel
[ 34.611601] Module Intercept Inserted
[ 58.020881] Intercepted my_syscall
tlimkim@tlimkim:~/workspace/my_syscall$
```

Figure 10. While calling my\_syscall, Intercepted

```
[ 34.611601] Module Intercept Inserted
[ 58.020881] Intercepted my_syscall
[ 110.786869] Module Intercept Removed
tlimkim@tlimkim:~/workspace/my_syscall$
```

Figure 11. Removed Module

```
[ 34.608267] intercept_syscall: loading out-of-tree module taints kernel.
[ 34.608701] intercept_syscall: module verification failed: signature and
tainting kernel
[ 34.611601] Module Intercept Inserted
[ 58.020881] Intercepted my_syscall
[ 110.786869] Module Intercept Removed
[ 124.700286] Hello World!
[ 124.700292] ** PID_Number: 1859 **
tlimkim@tlimkim:~/workspace/my_syscall$
```

Figure 12. Printing 'my\_syscall' result ordinarily

\*\* Additional things

- In task 2, there should be include one more header file "<linux/mm.h>".

<<References>>

- **Adding system call:**  
<http://harryp.tistory.com/69>,  
<http://lists.kernelnewbies.org/pipermail/kernelnewbies/2013-July/008598.html>  
<https://medium.com/@ssreehari/implementing-a-system-call-in-linux-kernel-4-7-1-6f98250a8c38>
- **Implementing hooking system call:**  
<https://stackoverflow.com/questions/14415561/intercepting-a-system-call>  
<https://stackoverflow.com/questions/2103315/linux-kernel-system-call-hooking-example>