



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Memory Management

**KAIST**

# Memory Management



## ■ Goals

- To provide a convenient abstraction for programming
- To allocate scarce memory resources among competing processes to maximize performance with minimal overhead
- To provide isolation between processes

## ■ Why is it so difficult?

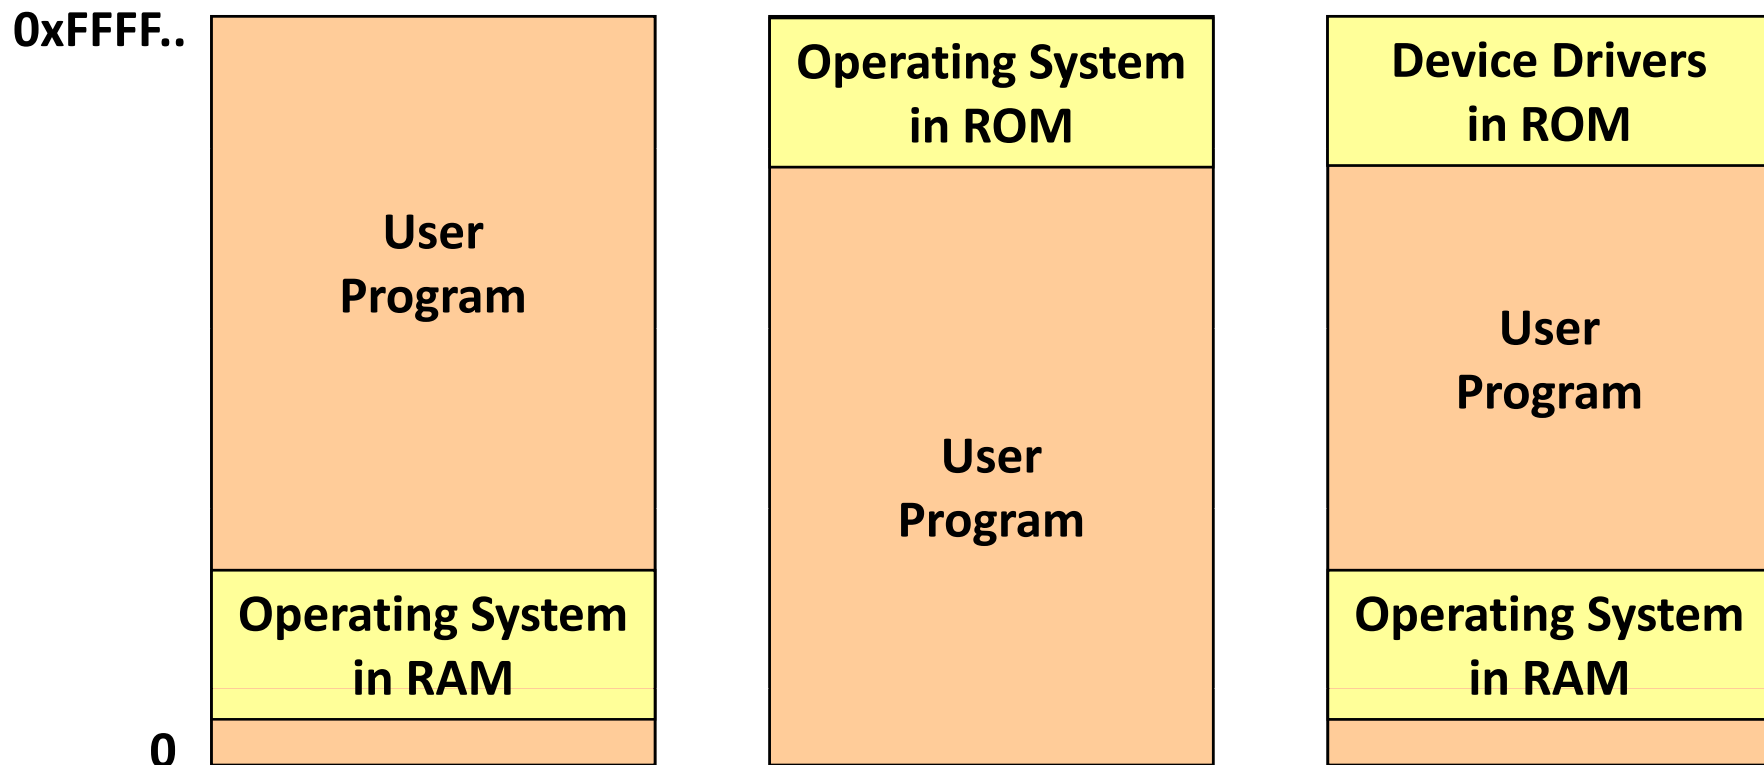
# Hardware Support



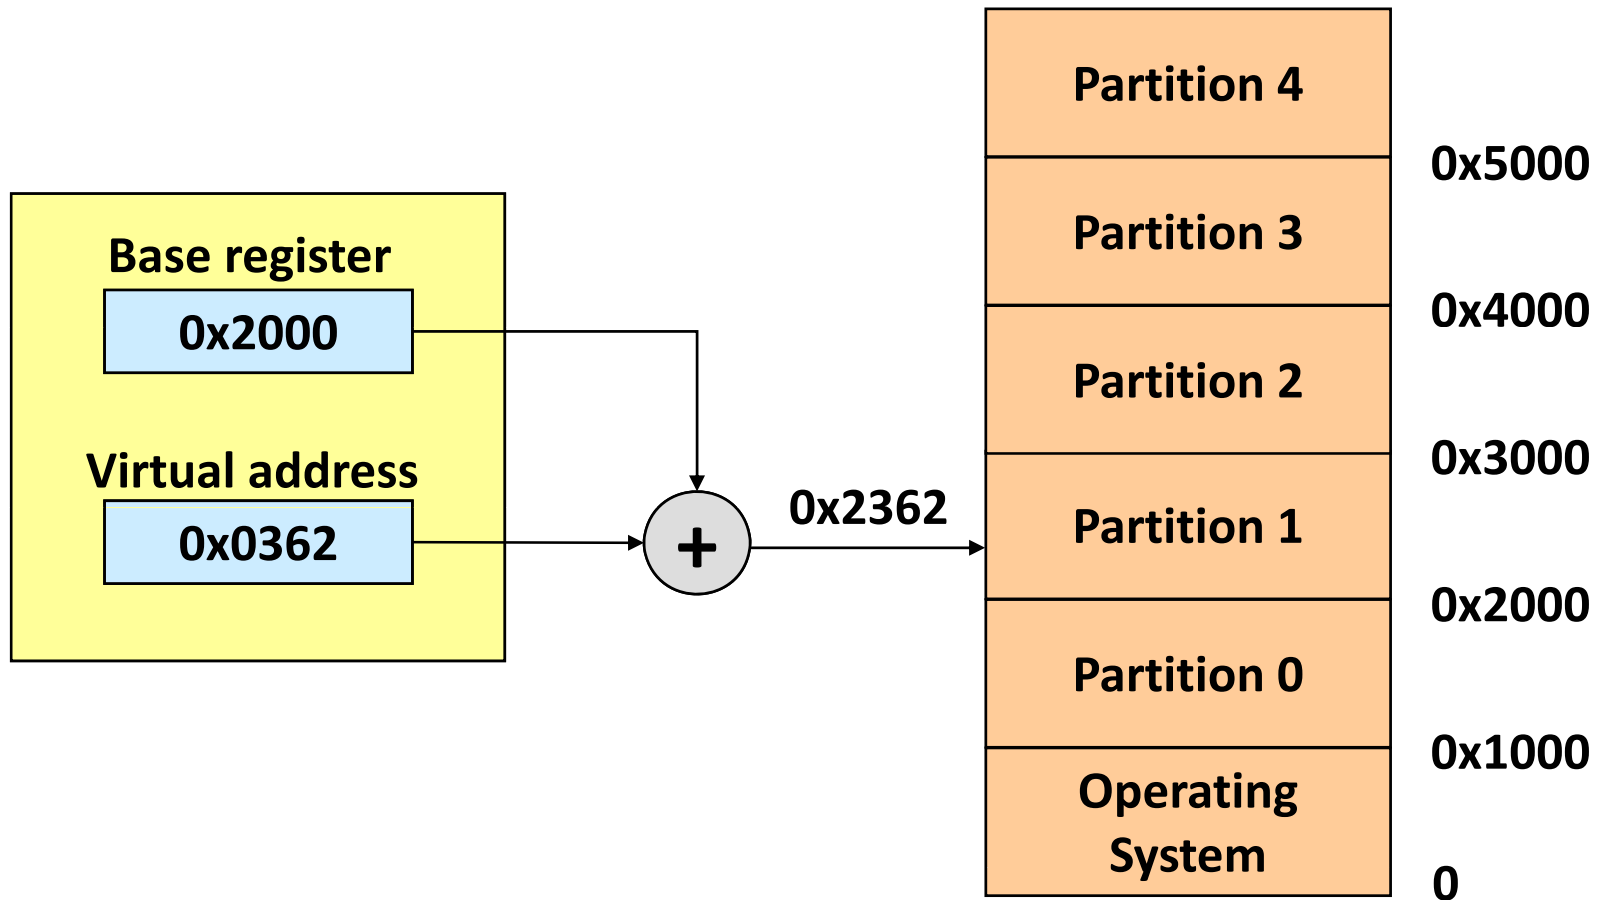
- **None**
  - CPU directly accesses physical memory
- **Memory Protection Unit (MPU)**
  - CPU directly accesses physical memory with memory protection
- **Memory Management Unit (MMU)**
  - Fully implements virtual memory
  - Linux runs with MMU-enabled CPUs

# Single/Batch Programming

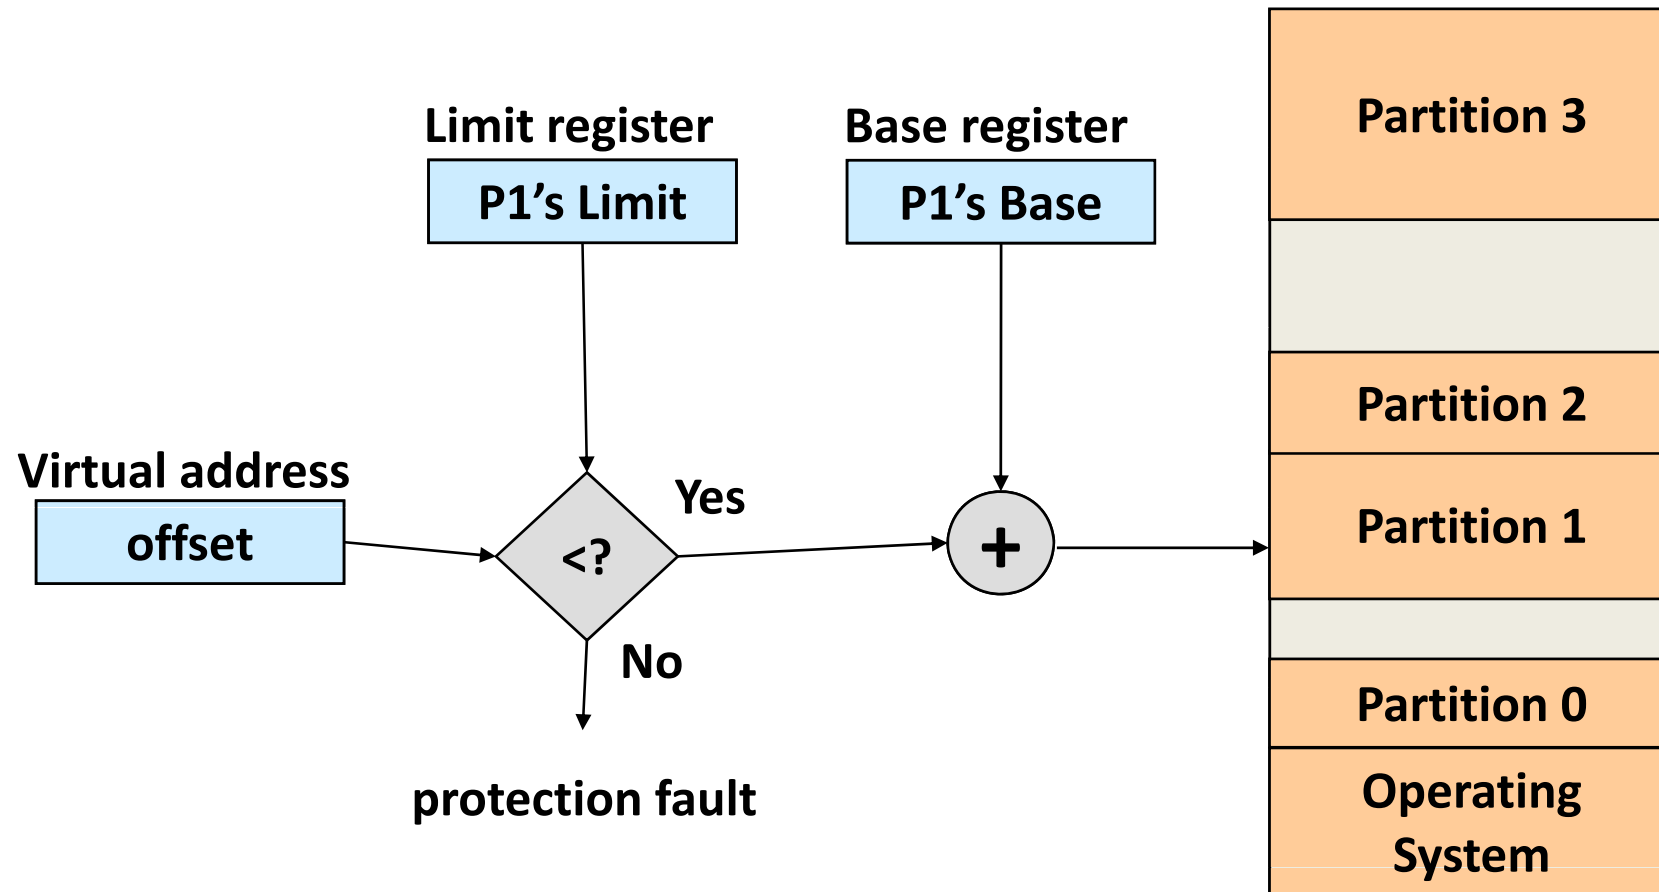
- **An OS with one user process**
  - Programs use physical addresses directly.
  - OS loads job, runs it, unloads it.



# Fixed Partitions



# Variable Partitions





# Virtual Memory



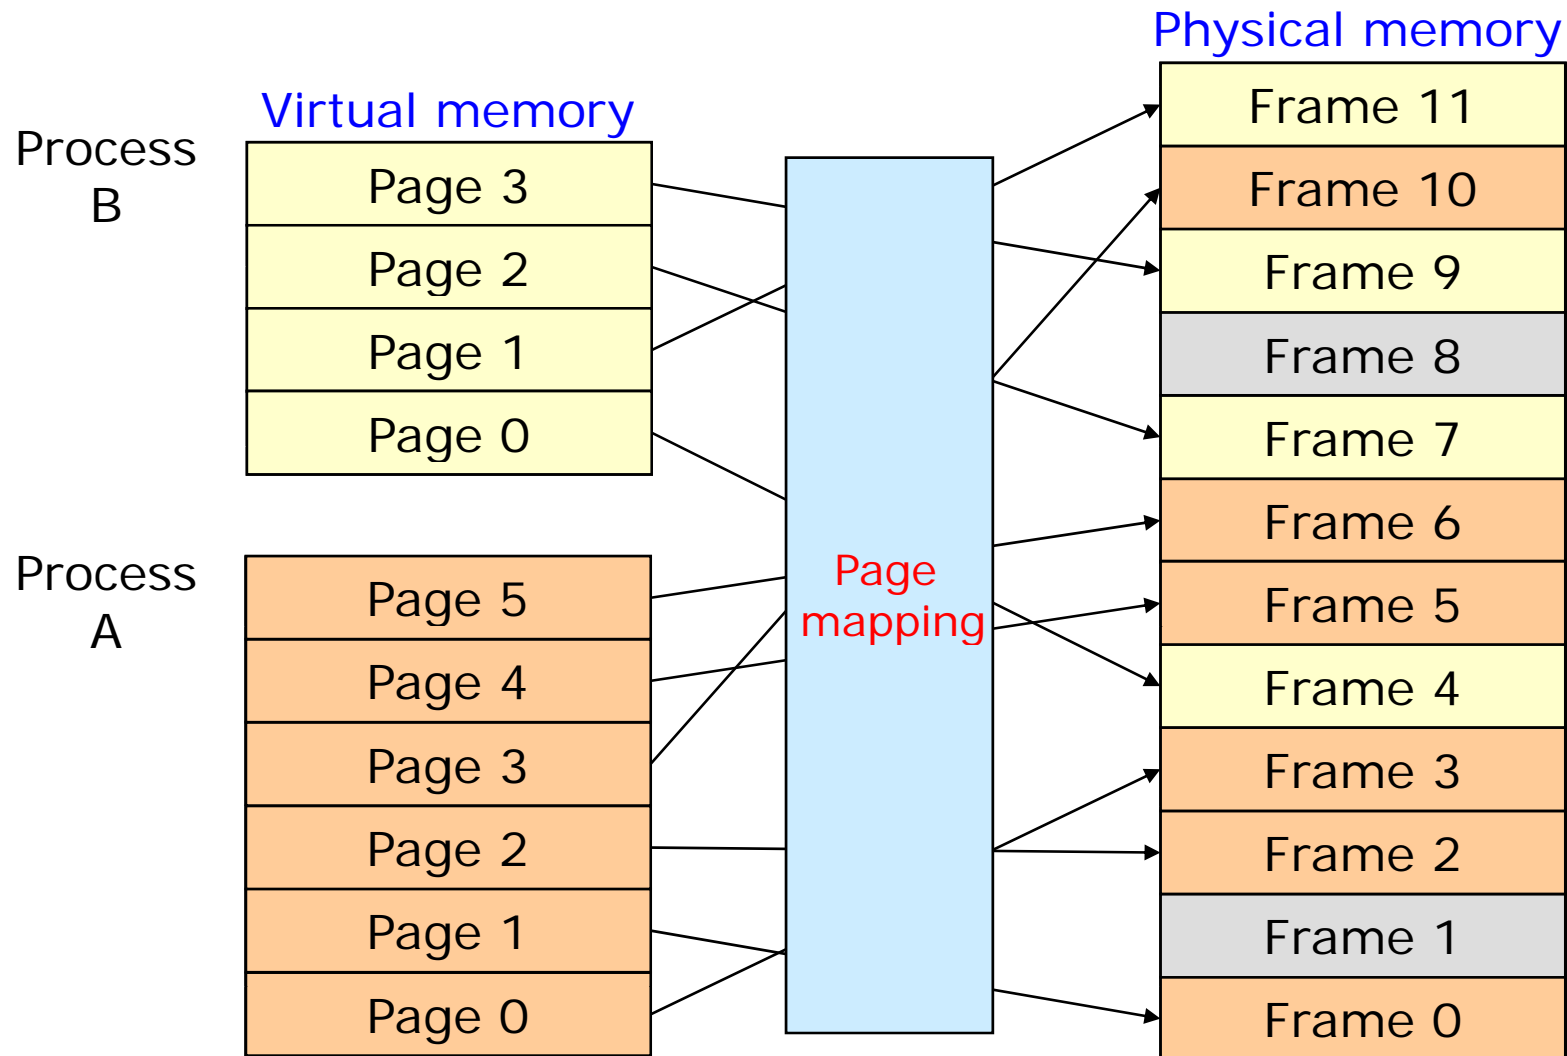
## ■ Why?

- 1.
- 2.
- 3.

## ■ How?

- MMU (Memory Management Unit)
- Address translation
- Demand paging
- Page tables

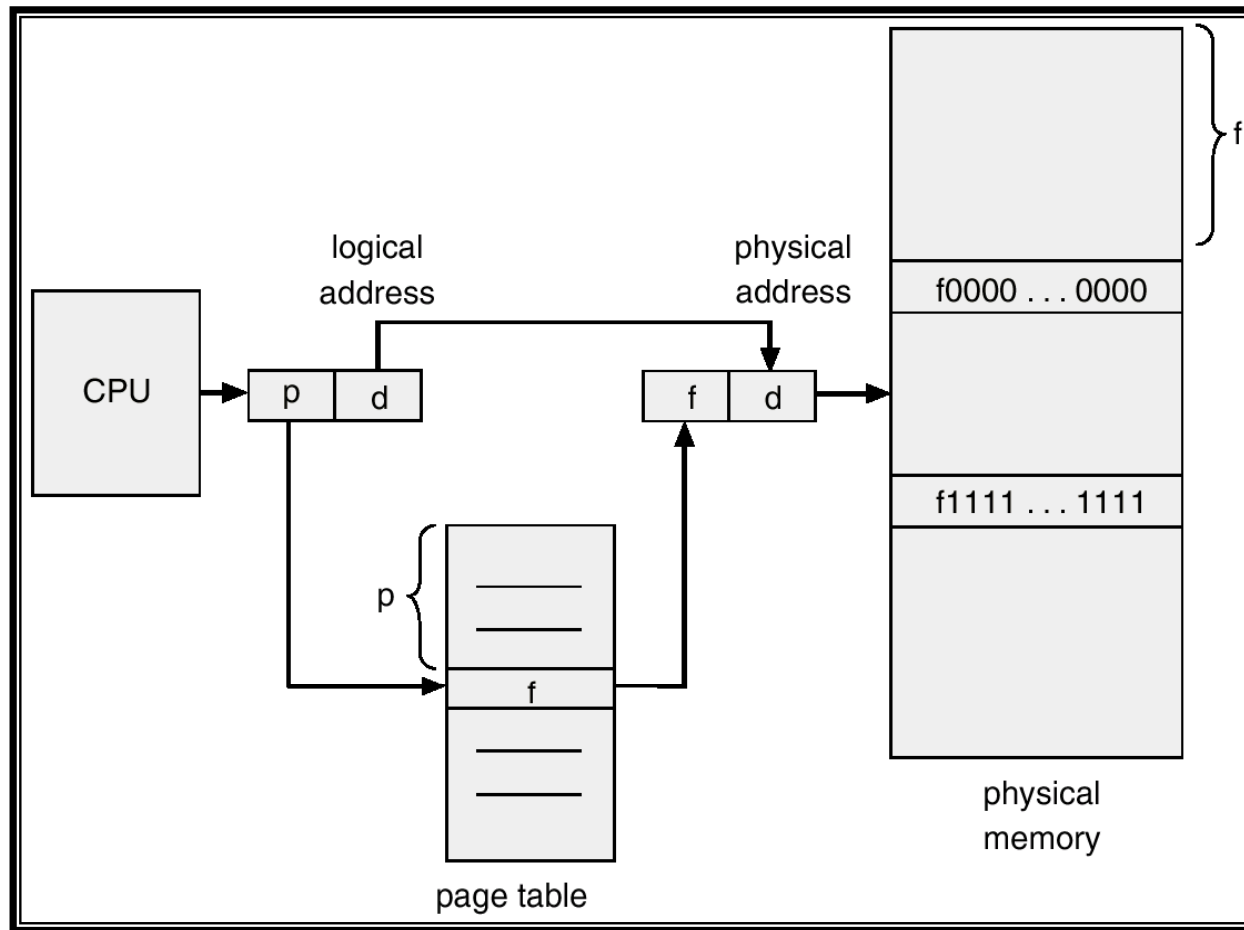
# Paging (1)





# Paging (2)

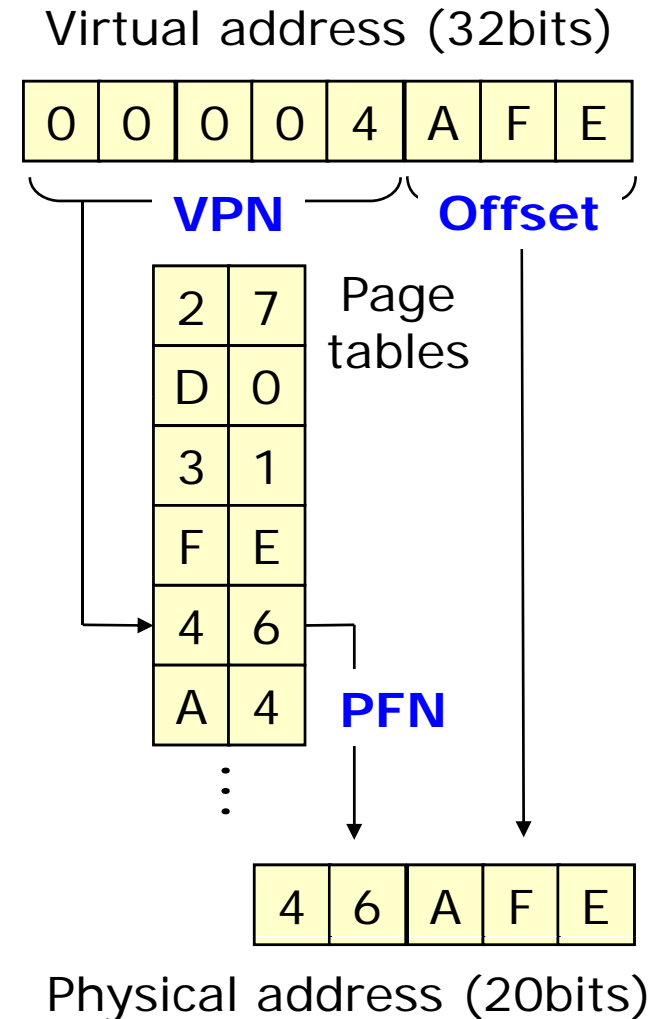
- Address translation architecture



# Paging (3)


## ■ Paging example

- Virtual address: 32 bits
- Physical address: 20 bits
- Page size: 4KB
- Offset: 12 bits
- VPN: 20 bits
- Page table entries:  $2^{20}$



# Paging (4)

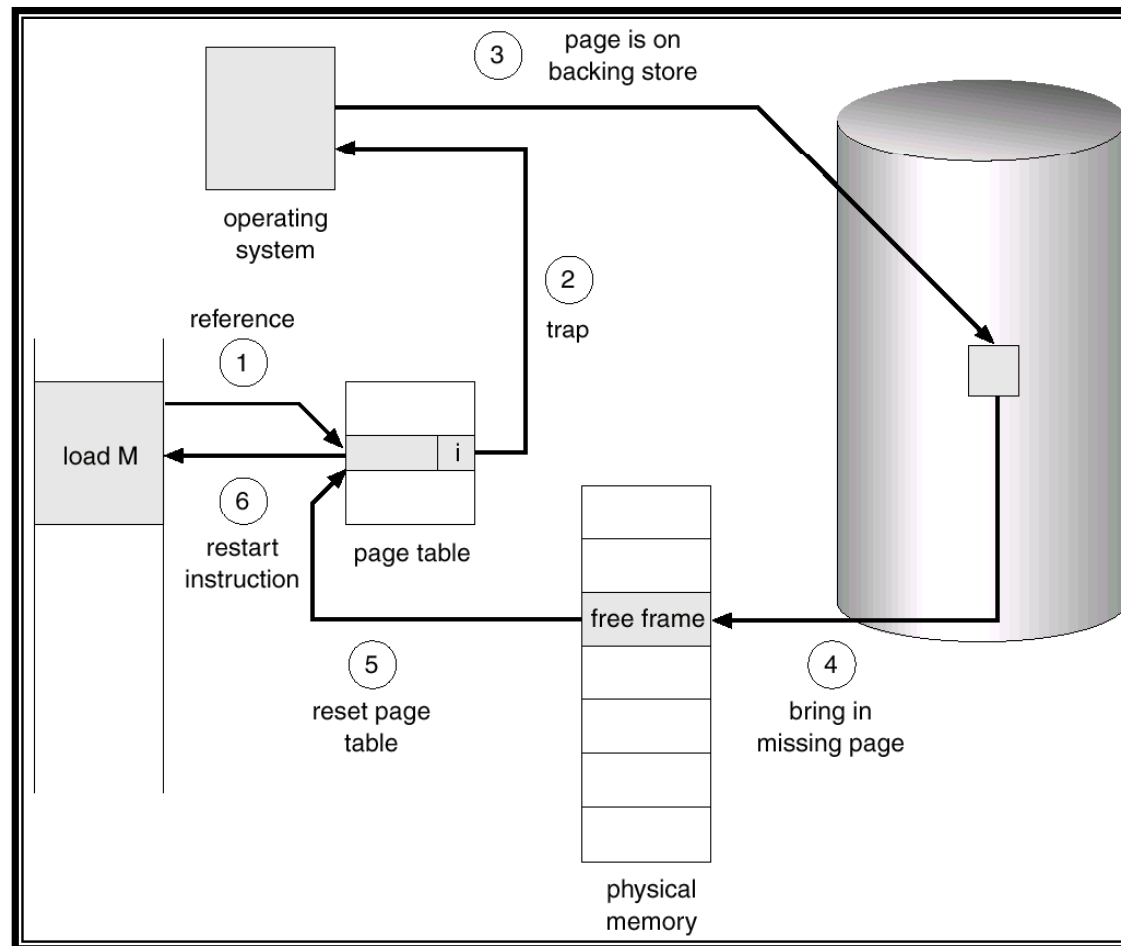
## ■ Page Table Entries (PTEs)



- Valid bit (V) says whether or not the PTE can be used.
  - It is checked each time a virtual address is used.
- Reference bit (R) says whether the page has been accessed.
  - It is set when a read or write to the page occurs.
- Modify bit (M) says whether or not the page is dirty.
  - It is set when a write to the page occurs.
- Protection bits (Prot) control which operations are allowed on the page.
  - Read, Write, Execute, etc.
- Page frame number (PFN) determines physical page.

# Paging (5)

- Handling a page fault



# Problems

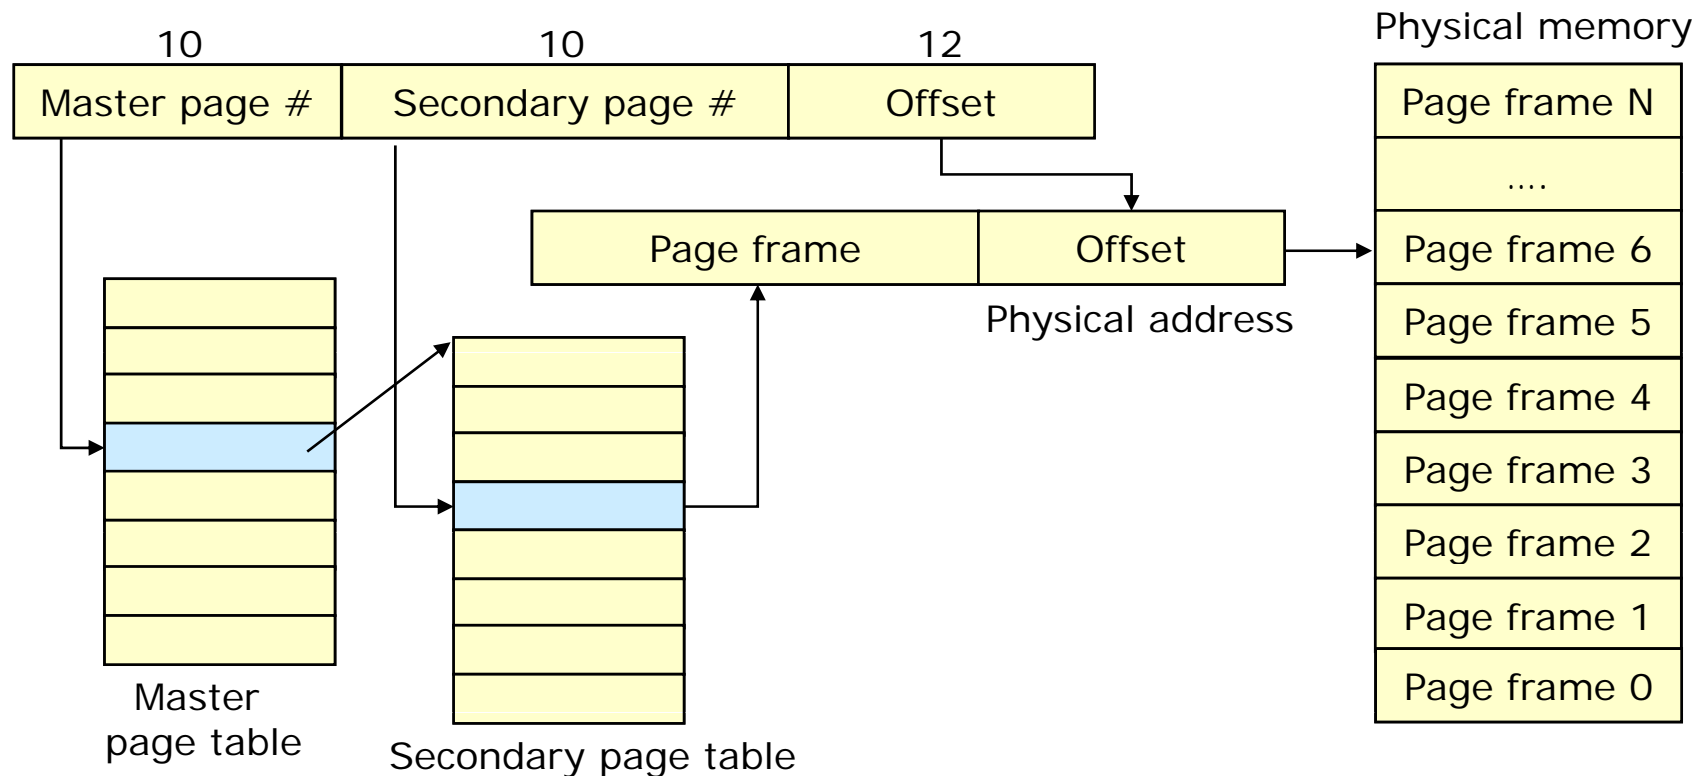


- **Space overhead**
  - Page tables
- **Time overhead**
  - Address translation

# Multi-level Page Tables

## ■ Example: Two-level Page Tables

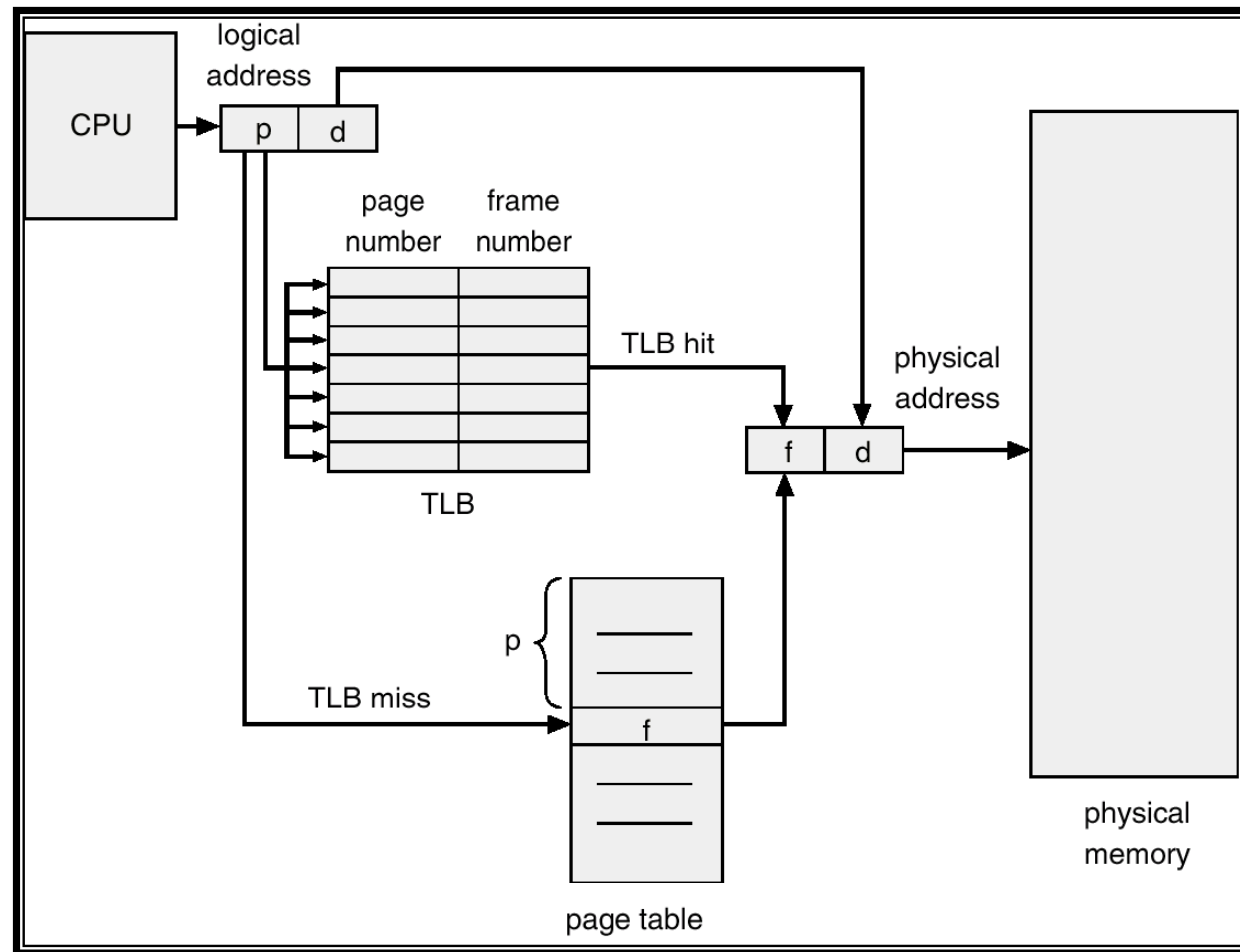
- 32-bit address space, 4KB pages, 4bytes/PTE
- Want master page table in one page





# TLBs

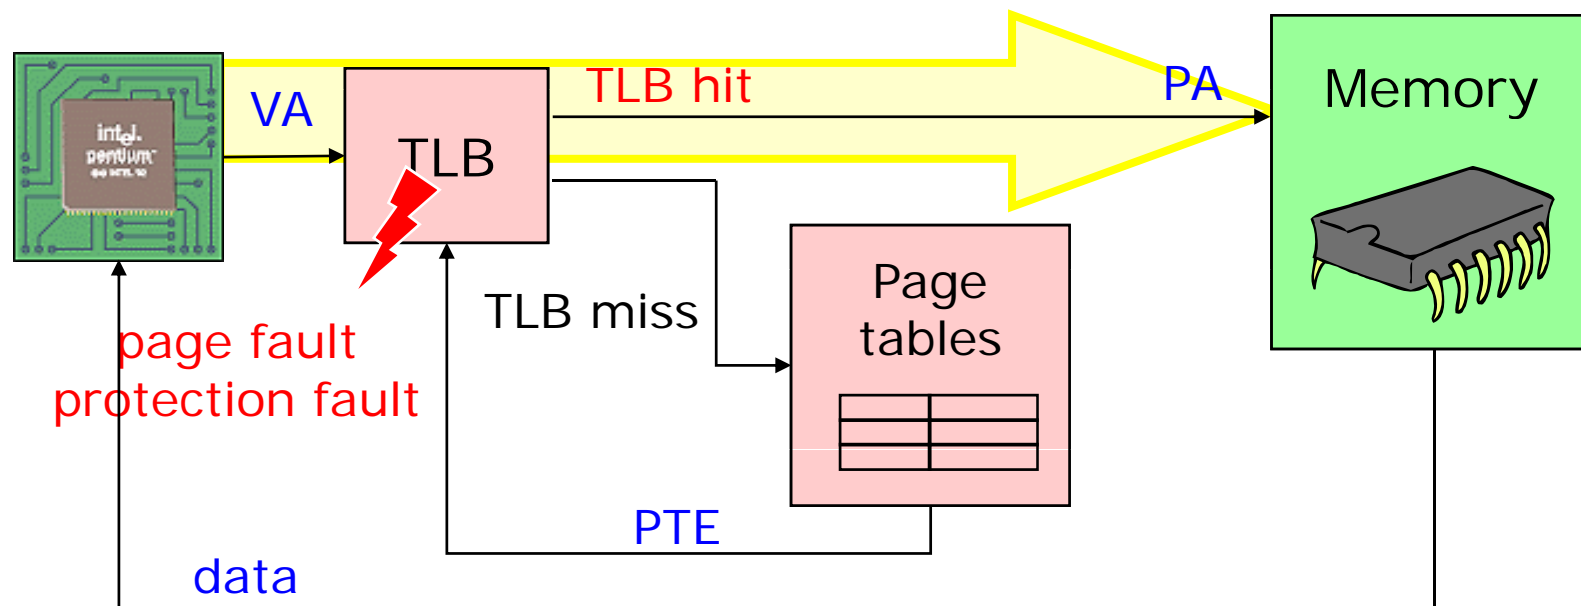
- Address translation with TLB



# Memory Reference

## ■ Situation

- Process is executing on the CPU, and it issues a read to a (virtual) address.



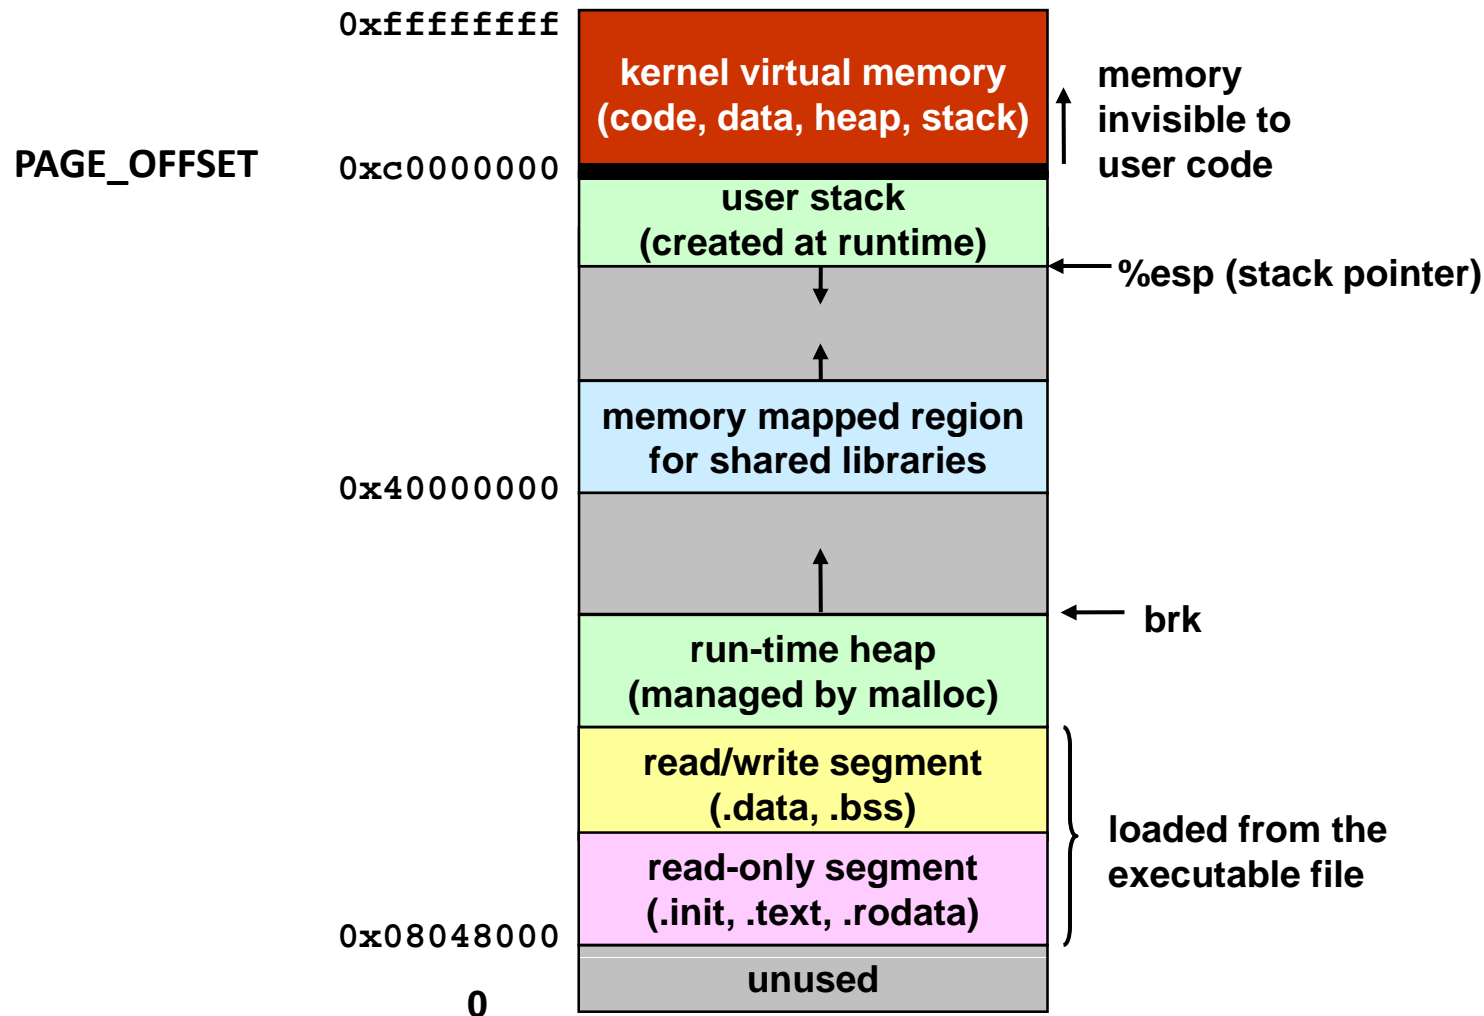


*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Virtual Memory Implementation

**KAIST**

# Process Address Space (1)



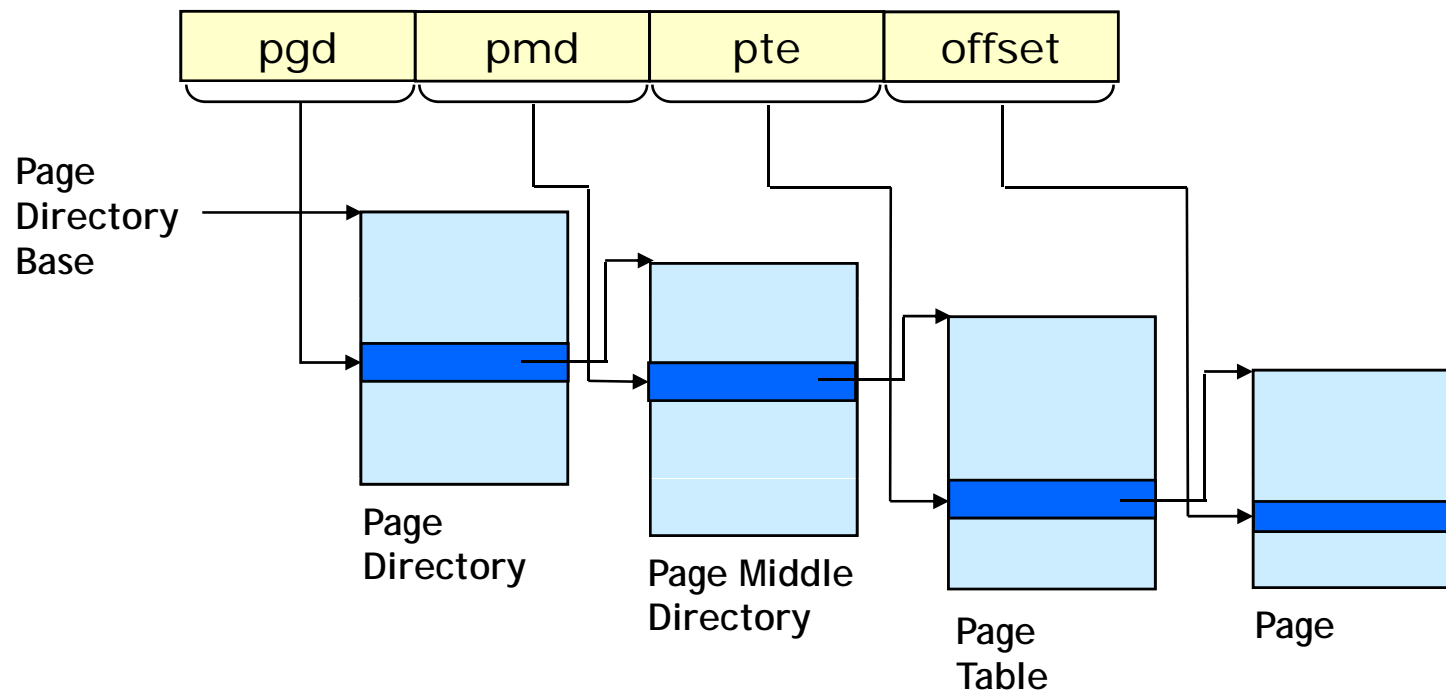
# Process Address Space (2)

## ■ Memory areas

- The intervals of legal addresses in process address space.
- Text section: code
- Data section: initialized global variables
- BSS section: uninitialized global variables
- User-space stack
- Heap
- Shared libraries (text, data, bss for each shlib)
- Memory mapped files
- Shared memory segments

# Process Address Space (3)

- **Paging: three-level address translation**
  - In i386, the size of Page Middle Directory (PMD) is 1, if the physical address extension (PAE) flag is disabled.





# Virtual Memory (1)



## ■ Demand paging

- Physical memory  $\leftrightarrow$  File system
- Pages are backed by files
  - Program code
  - (Initial) program data
  - Memory-mapped files, ...

## ■ Swapping

- Physical memory  $\leftrightarrow$  Swap area
- Anonymous pages
  - Stack, heap, BSS
  - (Written) program data
  - Shared memory, `mmap()` with `MAP_ANON`, ...

# Virtual Memory (2)



## ■ Page cache

- A cache of physical pages
- The page cache holds
  - Pages containing data of regular files
  - Pages containing directories
  - Pages containing data directly read from block device files
  - Pages containing data of user mode processes that have been swapped out on disk
  - Pages belonging to files of special filesystems (e.g., shm)
- Each page included in the page cache contains data belonging to some file.

# Virtual Memory (3)



## ■ Page fault

- Page fault mainly occurs due to
  - Not-present pages
  - Protection violation (especially for copy-on-write)
- Major page fault
  - If the kernel need to access the disk to make the page available.
- Minor page fault
  - If the kernel only need to allocate pages in RAM without reading anything from disk.

# VMA (1)


- **struct vm\_area\_struct** <linux/mm.h>
  - Nonoverlapping regions, each representing a continuous, page-aligned subset of the virtual address space.

struct mm_struct *	vm_mm;	associated mm_struct
unsigned long	vm_start;	VMA start, inclusive
unsigned long	vm_end;	VMA end, exclusive
struct vm_area_struct *	vm_next;	list of VMA's
pgprot_t	vm_page_prot;	access permissions
unsigned long	vm_flags;	VMA flags
struct rb_node	vm_rb;	VMA's node in the tree
struct vm_operations_struct *	vm_ops;	associated ops
struct file *	vm_file;	mapped file, if any.
unsigned long	vm_pgoff;	offset within the file
...		

# VMA (2)

## ■ Organization of VMAs

- Linked list (via  $\text{mm} \rightarrow \text{mmap}$ )
  - Simple and efficient for traversing of all elements
  - Sorted by ascended address (linked via  $\text{vma} \rightarrow \text{vm\_next}$ )
- Red-black tree (via  $\text{mm} \rightarrow \text{mm\_rb}$ )
  - A type of balanced binary tree.
    - » The root & all leaves are black.
    - » Both children of every red node are black.
    - » All paths from any given node to its leaf nodes contain the same number of black nodes.
  - Searching, insertion, deletion:  $O(\log(n))$
  - Used when locating a specific VMA in the address space.





# VMA (3)

## ■ VMA flags

Flag	Description
VM_READ / VM_WRITE / VM_EXEC	Pages can be read from / written to / executed.
VM_SHARED	Pages are shared.
VM_MAYREAD / VM_MAYWRITE / VM_MAYEXEC / VM_MAYSHARE	VM_READ / VM_WRITE / VM_EXEC / VM_SHARE flag can be set.
VM_GROWSDOWN / VM_GROWSUP	The area can grow downward / upward.
VM_SHM	The area is used for shared memory
VM_DENYWRITE / VM_EXECUTABLE	The area maps an unwritable file / an executable file
VM_LOCKED	The pages in this area are locked.
VM_IO	The area maps a device's I/O space.
VM_RESERVED	This area must not be swapped out.
VM_SEQ_READ / VM_RAND_READ	The pages seem to be accessed sequentially / randomly.



# VMA (4)

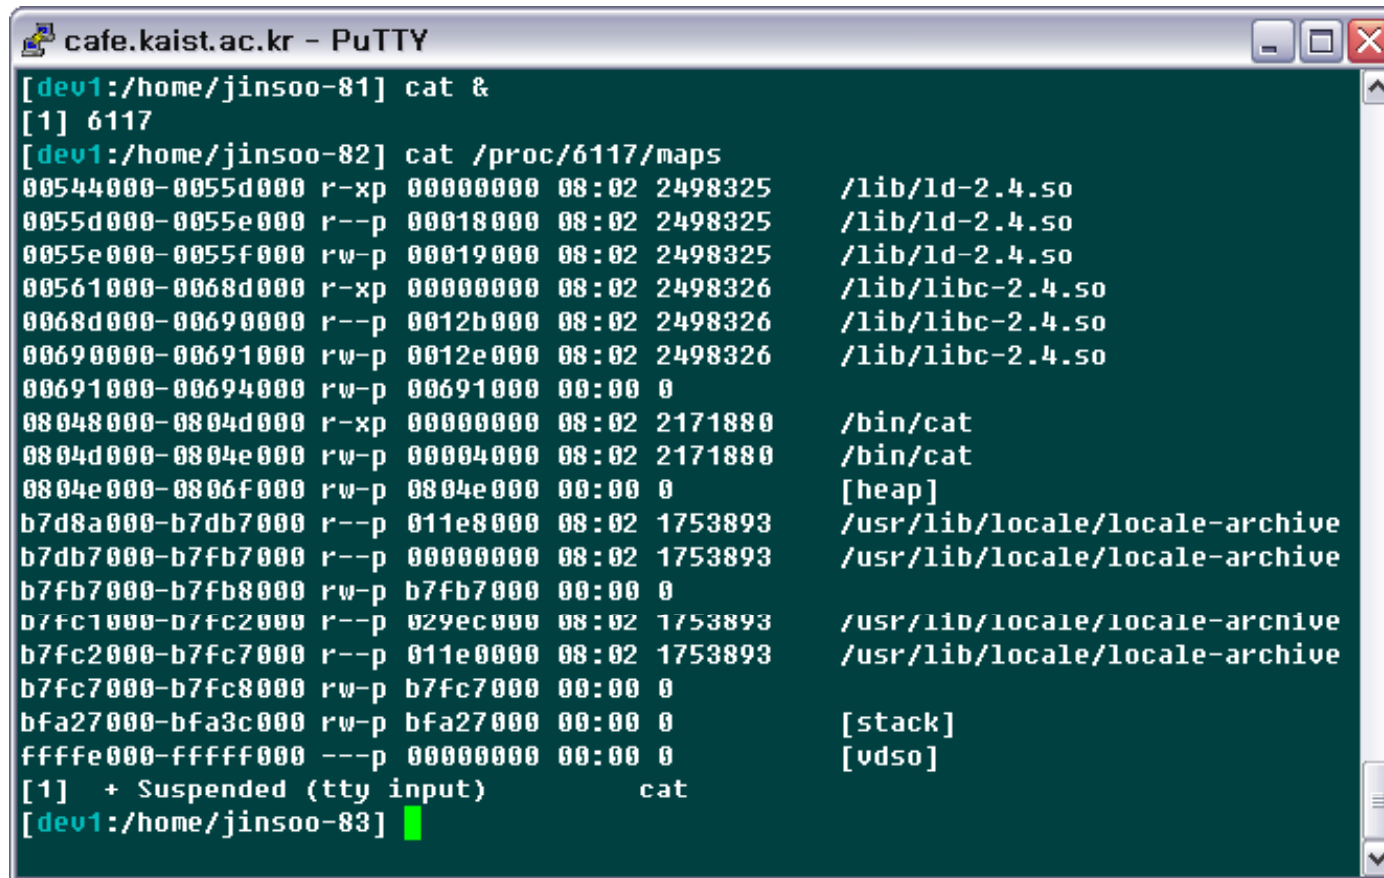


## ■ VMA operations

- `struct vm_operations_struct` `<linux/mm.h>`
- `void open(struct vm_area_struct *area);`
  - Invoked when the given VMA is added to an address space.
- `void close (struct vm_area_struct *area);`
- `struct page *nopcode(struct vm_area_struct *area, unsigned long address, int unused);`
  - Invoked by the page fault handler when a page that is not present in physical memory is accessed.
- `int populate(struct vm_area_struct *area, unsigned long address, unsigned long len, pgprot_t prot, unsigned long pgoff, int nonblock);`

# VMA (5)

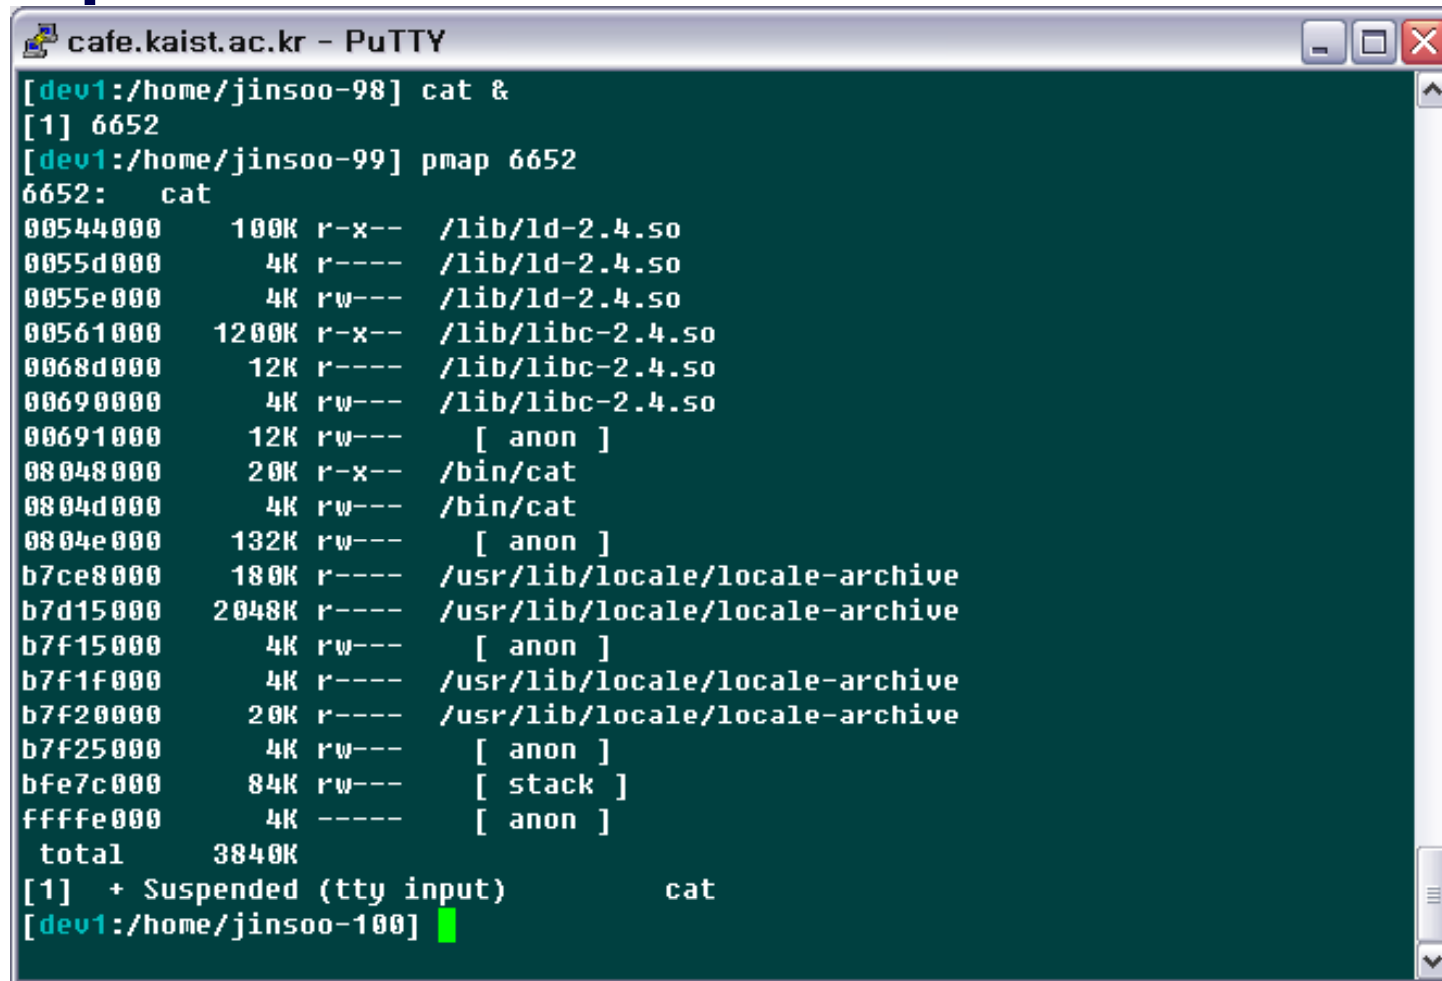
- `/proc/<pid>/maps`
  - start – end permission offset major:minor inode file



```
cafe.kaist.ac.kr - PuTTY
[dev1:/home/jinsoo-81] cat &
[1] 6117
[dev1:/home/jinsoo-82] cat /proc/6117/maps
00544000-0055d000 r-xp 00000000 08:02 2498325 /lib/ld-2.4.so
0055d000-0055e000 r--p 00018000 08:02 2498325 /lib/ld-2.4.so
0055e000-0055f000 rw-p 00019000 08:02 2498325 /lib/ld-2.4.so
00561000-0068d000 r-xp 00000000 08:02 2498326 /lib/libc-2.4.so
0068d000-00690000 r--p 0012b000 08:02 2498326 /lib/libc-2.4.so
00690000-00691000 rw-p 0012e000 08:02 2498326 /lib/libc-2.4.so
00691000-00694000 rw-p 00691000 00:00 0
08048000-0804d000 r-xp 00000000 08:02 2171880 /bin/cat
0804d000-0804e000 rw-p 00004000 08:02 2171880 /bin/cat
0804e000-0806f000 rw-p 0804e000 00:00 0 [heap]
b7d8a000-b7db7000 r--p 011e8000 08:02 1753893 /usr/lib/locale/locale-archive
b7db7000-b7fb7000 r--p 00000000 08:02 1753893 /usr/lib/locale/locale-archive
b7fb7000-b7fb8000 rw-p b7fb7000 00:00 0
b7fc1000-b7fc2000 r--p 029ec000 08:02 1753893 /usr/lib/locale/locale-archive
b7fc2000-b7fc7000 r--p 011e0000 08:02 1753893 /usr/lib/locale/locale-archive
b7fc7000-b7fc8000 rw-p b7fc7000 00:00 0
bfa27000-bfa3c000 rw-p bfa27000 00:00 0 [stack]
ffffe000-ffffff00 ---p 00000000 00:00 0 [vdso]
[1] + Suspended (tty input) cat
[dev1:/home/jinsoo-83]
```

# VMA (6)

- pmap



```
cafe.kaist.ac.kr - PuTTY
[dev1:/home/jinsoo-98] cat &
[1] 6652
[dev1:/home/jinsoo-99] pmap 6652
6652:  cat
00544000    100K r-x--  /lib/ld-2.4.so
0055d000      4K r----  /lib/ld-2.4.so
0055e000      4K rw---  /lib/ld-2.4.so
00561000   1200K r-x--  /lib/libc-2.4.so
0068d000     12K r----  /lib/libc-2.4.so
00690000      4K rw---  /lib/libc-2.4.so
00691000     12K rw---  [ anon ]
08048000     20K r-x--  /bin/cat
0804d000      4K rw---  /bin/cat
0804e000   132K rw---  [ anon ]
b7ce8000    180K r----  /usr/lib/locale/locale-archive
b7d15000   2048K r----  /usr/lib/locale/locale-archive
b7f15000      4K rw---  [ anon ]
b7f1f000      4K r----  /usr/lib/locale/locale-archive
b7f20000     20K r----  /usr/lib/locale/locale-archive
b7f25000      4K rw---  [ anon ]
bfe7c000     84K rw---  [ stack ]
ffffe000      4K ----- [ anon ]
total    3840K
[1] + Suspended (tty input)      cat
[dev1:/home/jinsoo-100]
```

# Memory Descriptor (1)

- **struct mm\_struct** <linux/sched.h>
  - Contains all the information related to the process address space.
  - Threads share a memory descriptor.
  - Doubly linked via the mmlist field.

struct vm_area_struct *	mmap;	list of memory areas (VMAs)
struct rb_root	mm_rb;	red-black tree of memory areas
pgd_t *	pgd;	page global directory
atomic_t	mm_users;	address space users
atomic_t	mm_count;	reference count for mm_struct
int	map_count;	number of memory areas
struct list_head	mmlist;	list of all mm_structs
...		

# Memory Descriptor (2)

## ■ Allocating/freeing a memory descriptor

- `allocate_mm()` <kernel/fork.c>  
= `kmem_cache_alloc(mm_cachep, GFP_KERNEL)`
- `free_mm(mm)`  
= `kmem_cache_free(mm_cachep, (mm))`

```
copy_mm() {                                     <kernel/fork.c>
    ...
    if (clone_flags & CLONE_VM) {                // thread creation
        atomic_inc(&current->mm->mm_users);
        tsk->mm = current->mm;
    } else {                                     // process creation
        tsk->mm = allocate_mm();
        mm_init(tsk->mm);
        ...
    }
}
```



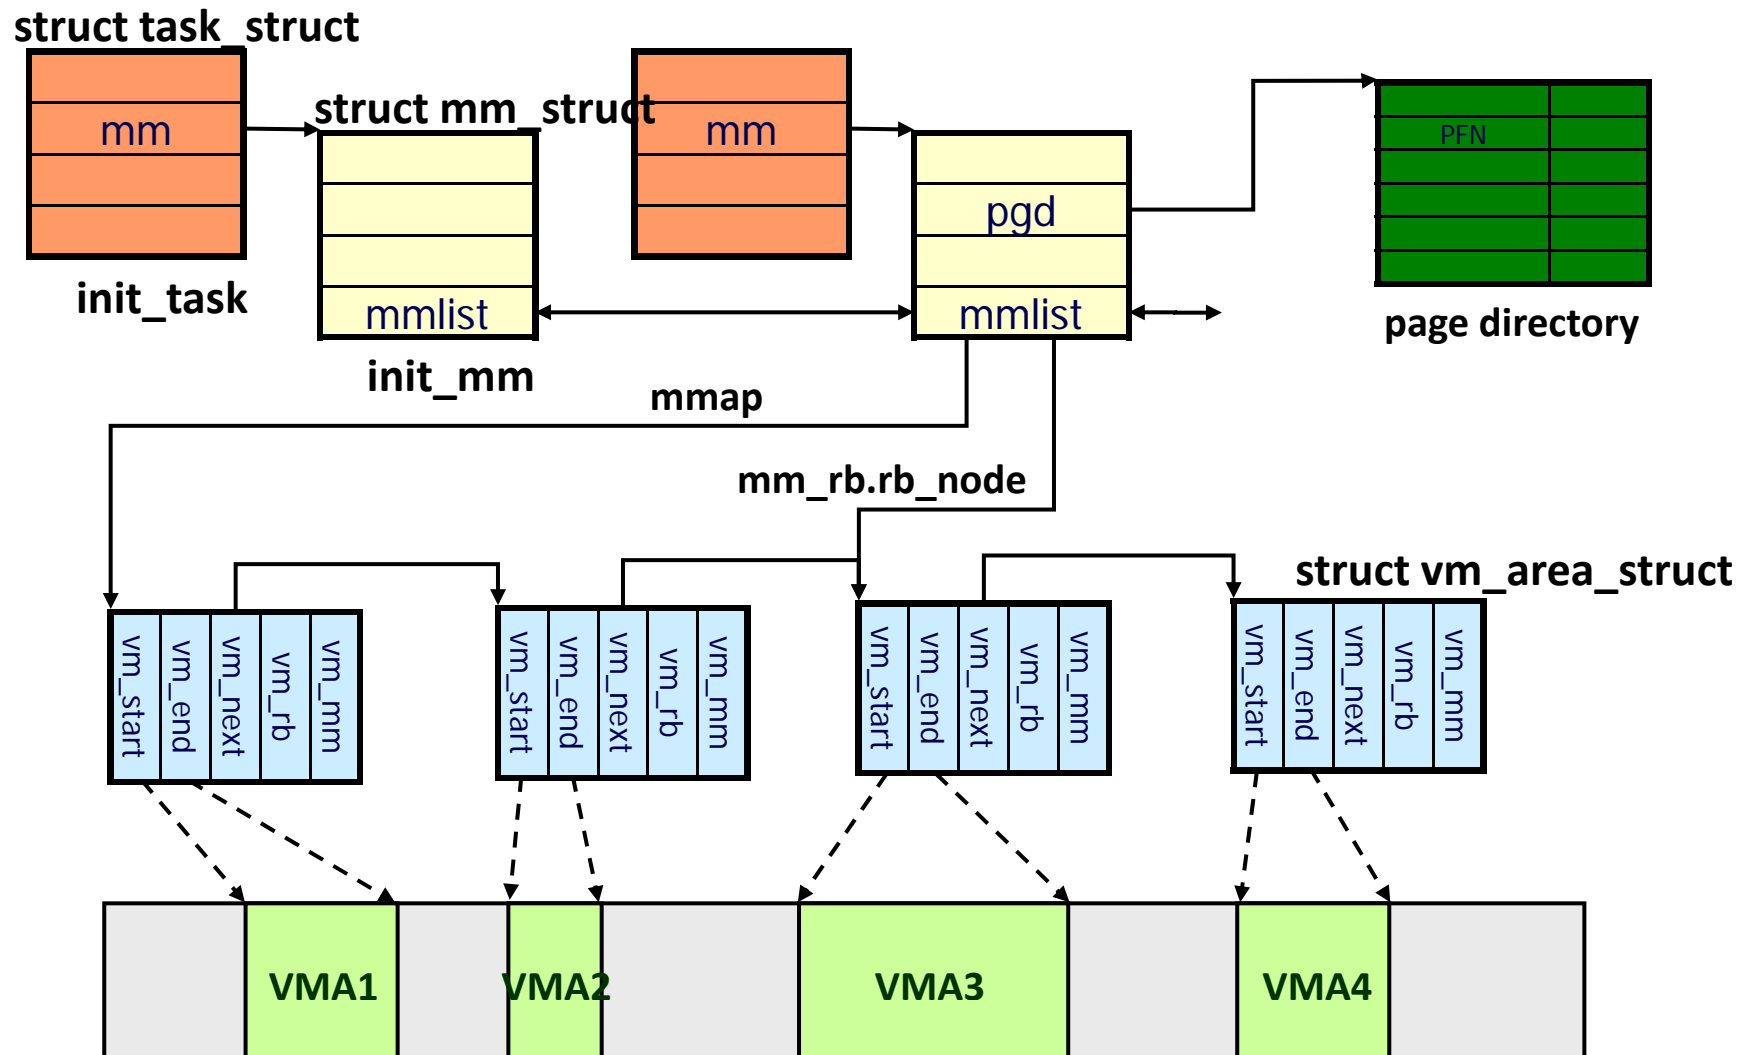
# Memory Descriptor (3)

## ■ Kernel threads

- Kernel threads do not have a process address space and therefore do not have a memory descriptor.
- Kernel threads use the memory descriptor of whatever task ran previously.
  - `task→active_mm`: address space referenced by a process
  - When a kernel thread is scheduled (`task→mm == NULL`), the kernel keeps the previous process's address space loaded.
  - The kernel updates the `active_mm` field.
  - The kernel thread can then use the previous process's page tables as needed.
- Kernel threads use only the information pertaining to kernel memory. (same for all processes)



# Memory Descriptor (4)



# do\_mmap() (1)

- unsigned long **do\_mmap** (struct file \*file, unsigned long addr, unsigned long len, unsigned long prot, unsigned long flag, unsigned long offset);
  - **file and offset:**
    - Specified when the memory region is backed by a file (file-backed mapping)
    - file==NULL and offset==0 for anonymous mapping
  - **addr:** linear address where the search for a free interval must start.
  - **prot:** the access rights of the pages in the region
  - **flag:** the memory region flags

# do\_mmap() (2)



## ■ Protection

- PROT\_READ (= VM\_READ)
- PROT\_WRITE (= VM\_WRITE)
- PROT\_EXEC (= VM\_EXEC)
- PROT\_NONE: none of access rights

## ■ Flags

- MAP\_SHARED: shared among several processes
- MAP\_PRIVATE: private to this process
- MAP\_FIXED: must be exactly the specified address
- MAP\_ANONYMOUS: no file is associated
- MAP\_POPULATE: pre-allocate the page frames

# do\_mmap() (3)

## ■ Implementation for anonymous mapping

- Check the parameters
- Obtain a linear address interval for the new region
  - `get_unmapped_area()`
- Compute VM flags based on `prot` and `flags` parameters
- Locate a VMA structure that precedes the new interval
  - `find_vma_prepare()`
- Check against address space limit
- Check if an old private anonymous mapping can be expanded.
  - They should have exactly the same flags.
  - If so, merge two VMAs using `vma_merge()`

# do\_mmap() (4)



## ■ Implementation (cont'd)

- Allocate a `vm_area_struct` for the new memory region
  - `kmem_cache_alloc()`
- Initialize the `vm_area_struct`
- Insert the new region in the memory region list and red-black tree
  - `vma_link()`
- Increase the accounting information
  - `mm->total_vm`
- If `VM_LOCKED` is set, allocate all the pages of the memory region and lock them in RAM
  - `make_pages_present()`
- Return the linear address of the new memory region



# do\_munmap() (1)

- int **do\_munmap** (struct mm\_struct \*mm,  
                  unsigned long start, size\_t len);
  - mm: the process's memory descriptor
  - start: the starting address of the interval
  - len: the length of the interval
- Phase 1: Scan the list of memory regions owned by the process and unlink all regions included in the linear address space of the process address space.
- Phase 2: update the process page tables and remove the memory regions



# do\_munmap() (2)



## ■ Phase 1 implementation

- Check the parameters
- Locate the first memory region that ends after the linear address interval to be deleted.
  - find\_vma\_prev()
- Split the first memory region if the start address lies inside the region
  - split\_vma()
- Split the last memory region if the linear address ends inside the region
  - split\_vma()

# do\_munmap() (3)



## ■ Phase 2 implementation

- Remove the memory regions included in the linear address interval from the process's linear address space
  - detach\_vmas\_to\_be\_unmapped()
- Clear the page table entries covering the linear address interval and free the corresponding page frames
  - unmap\_region()
- Release the descriptors of the memory regions and adjust the accounting information
  - remove\_vma\_list()

# Page Fault Handling

