




CS632/SEP564: Embedded Operating Systems (Fall 2008)

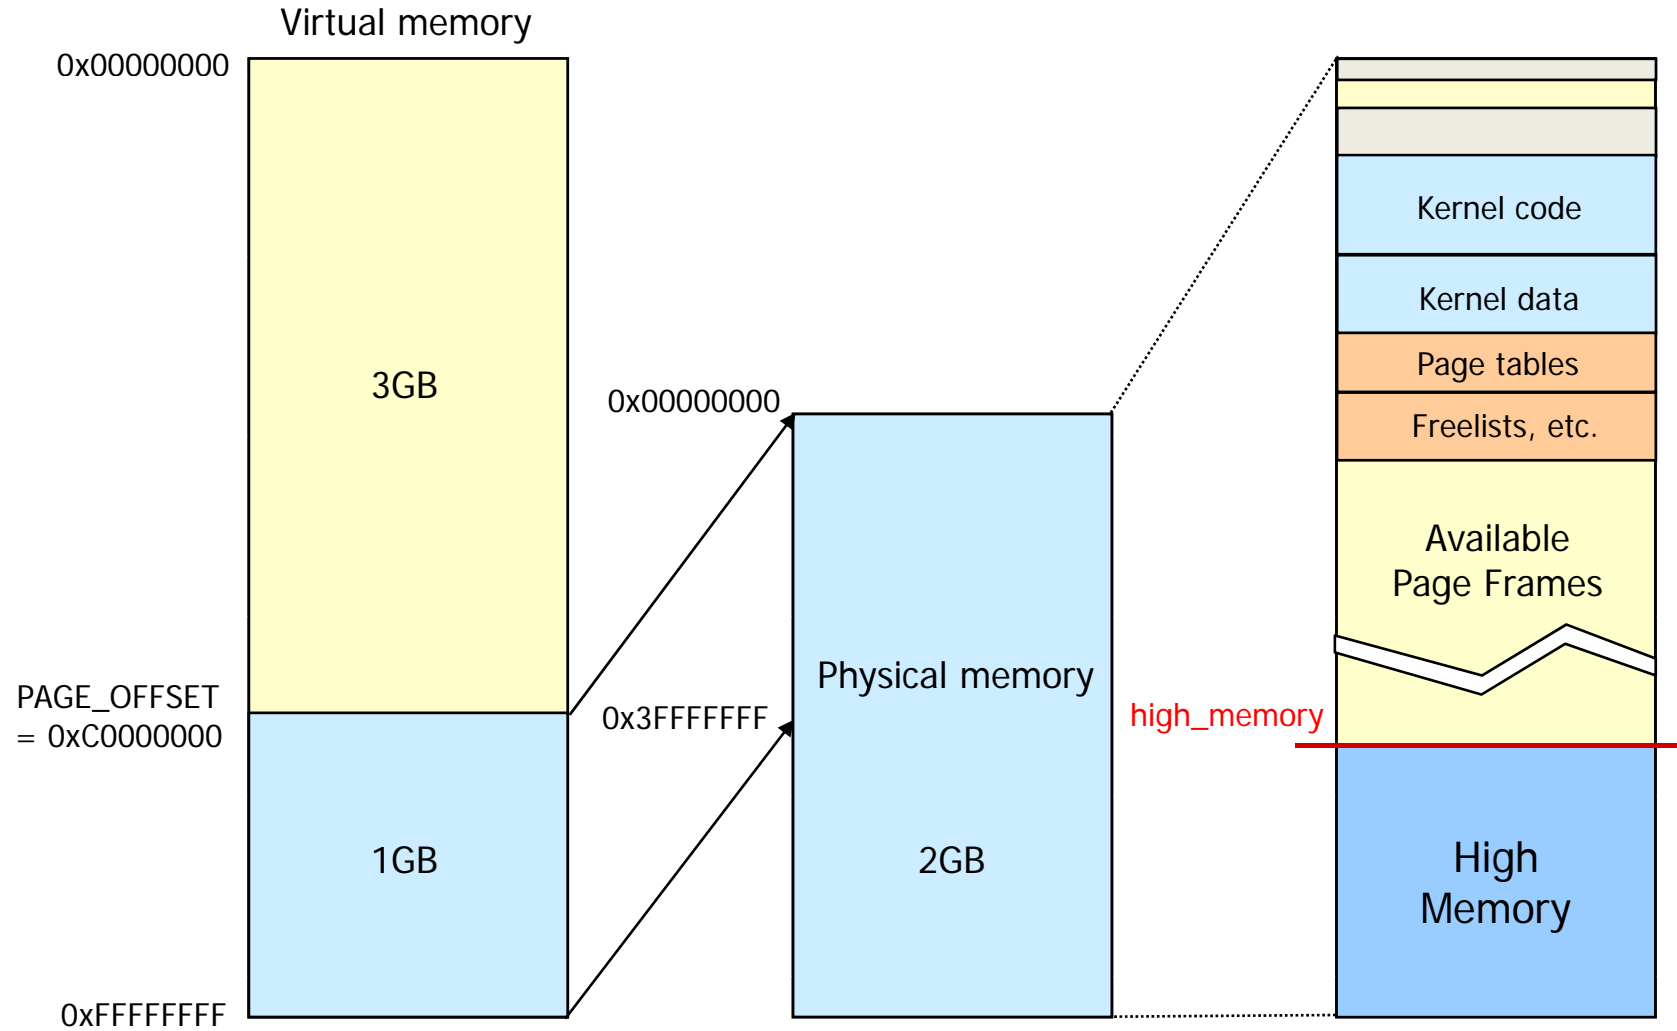
Memory Management

KAIST

The Big Picture



Using Physical Memory



Pages

■ **struct page** <linux/mm.h>

- All page descriptors are stored in the **mem_map** array.
- One entry for every physical page
 - About 1MB for 128MB physical memory with 4KB page size

page_flags_t	flags;	32 flags defined in <linux/page-flags.h>
atomic_t	_count;	The number of references to this page Accessed via page_count().
atomic_t	_mapcount;	Count of PTEs mapped.
unsigned long	private;	Private data
struct address_space *	mapping;	How the page is used?
pgoff_t	index;	Offset within the mapping
struct list_head	lru;	LRU list

Zones (1)



■ Why memory zones?

- Some hardware devices are capable of performing DMA to only certain memory addresses.
- Some architectures are capable of physically addressing larger amounts of memory than they can virtually address.

■ Memory zones in Linux

ZONE_DMA	DMA-able pages	< 16MB
ZONE_NORMAL	Normally addressable pages	16MB – 896MB
ZONE_HIGHMEM	Dynamically addressable pages	> 896MB

Zones (2)

- **struct zone** <linux/mmzone.h>
 - Page frame reclamation is performed on a per zone basis.

spinlock_t	lock;	Spin lock protecting the descriptor
unsigned long	free_pages;	Number of free pages in the zone
unsigned long	pages_min, pages_low, pages_high;	Parameters for page frame reclaiming
struct list_head	active_list;	List of active pages in the zone
struct list_head	inactive_list;	List of inactive pages in the zone
struct pglist_data *	zone_pgdat;	Pointer to the node descriptor
struct page *	zone_mem_map;	Pointer to the first page descriptor
char *	name;	Name of the zone
struct free_area	free_area[];	Buddy map
...		

Zones (3)



■ Reserved page frames

- Memory allocation requests can be satisfied immediately without memory reclaiming.
- Reduces the chance of failure in case of atomic memory allocation requests (GFP_ATOMIC).
- For ZONE_DMA and ZONE_NORMAL:
 - $\text{min_free_kbytes} = \text{sqrt}(\text{directly_mapped_memory} * 16)$ (KB)
 - » 128KB ~ 64MB (512KB for 16MB, 8MB for 4GB)
 - $\text{zone} \rightarrow \text{pages_min} = \text{the zone's contribution to min_free_kbytes} / \text{pagesize};$
 - $\text{zone} \rightarrow \text{pages_low} = \text{zone} \rightarrow \text{pages_min} * 5 / 4;$
 - $\text{zone} \rightarrow \text{pages_high} = \text{zone} \rightarrow \text{pages_min} * 3 / 2;$

Nodes

- **pg_data_t** <linux/mmzone.h>
 - The physical memory is partitioned in several nodes.
 - All node descriptors are stored in a singly linked list, whose first element is pointed to by [pgdat_list](#).

zone_t	node_zones[];	Zones for this node.
zonelist_t	node_zonelists[];	The order of zones that allocations are preferred from.
int	nr_zones;	The number of zones. Not all nodes will have the same number of zones.
struct page *	node_mem_map;	The first page of the struct page array
unsigned long	node_start_pfn;	Index of the first page frame in the node
pg_data_t *	pgdat_next;	Next item in the memory node list
int	node_id;	Node ID
...		

Zone Allocator (1)



■ Requesting page frames

- `struct page *alloc_pages(gfp_t gfp_mask, unsigned int order);`
- `struct page *alloc_page(gfp_t gfp_mask);`
- `unsigned long __get_free_pages(gfp_t gfp_mask, unsigned int order);`
- `unsigned long __get_free_page(gfp_t gfp_mask);`
- `unsigned long get_zeroed_page(gfp_t gfp_mask);`
- `unsigned long __get_dma_pages(gfp_t gfp_mask, int order);`
 - `= __get_free_pages((gfp_mask) | GFP_DMA, (order));`

Zone Allocator (2)



■ Releasing page frames

- void `free_pages`(unsigned long addr, unsigned int order);
- void `__free_pages`(struct page *page, unsigned int order);
- void `free_page`(unsigned long addr);
- void `__free_page`(struct page *page);

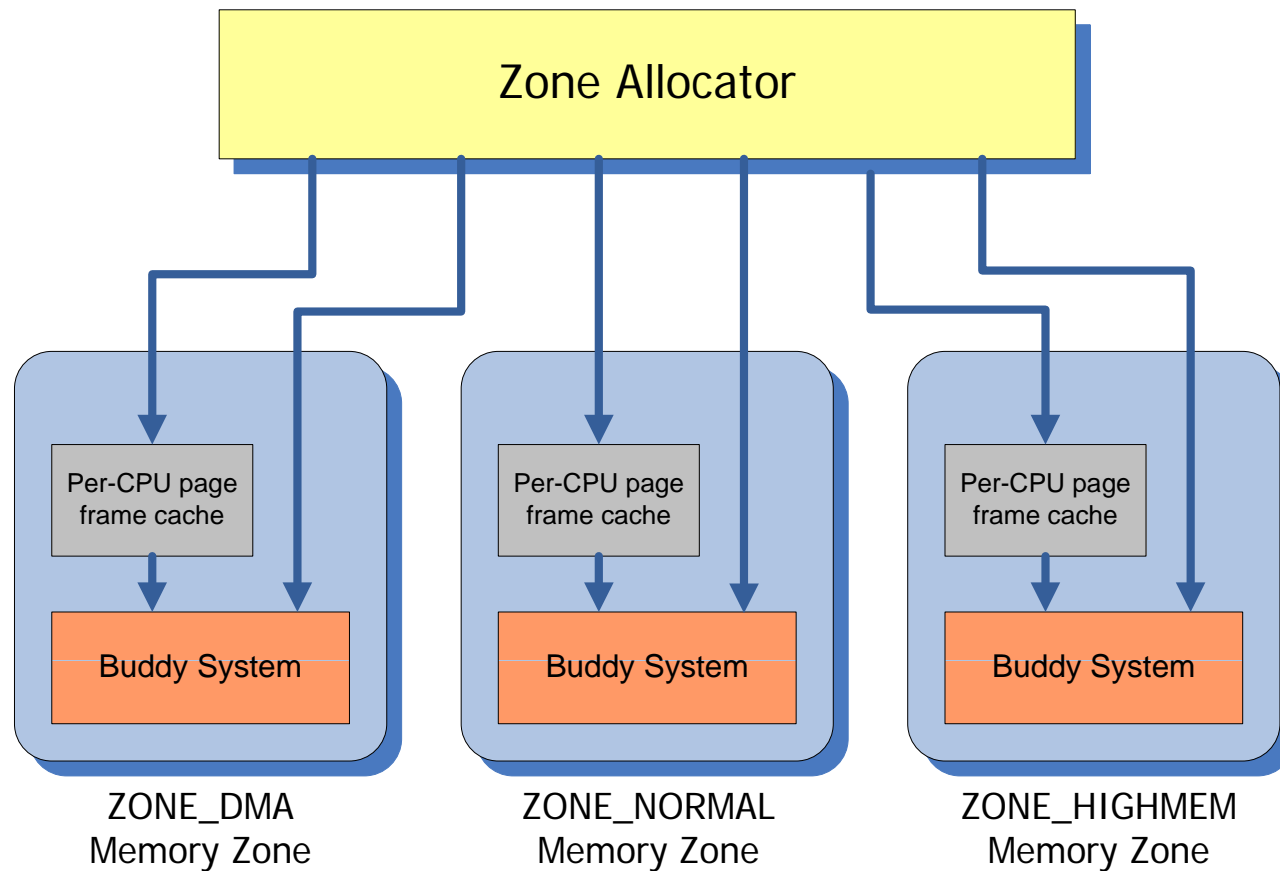
Zone Allocator (3)

- Some flags used to request page frames

Flag	Description
__GFP_DMA	The page frame must belong to the ZONE_DMA zone. Equivalent to GFP_DMA.
__GFP_HIGHMEM	The page frame may belong to the ZONE_HIGHMEM zone.
__GFP_WAIT	Allowed to block the current process waiting for free page frames.
__GFP_HIGH	Allowed to access the pool of reserved page frames.
__GFP_IO	Allowed to perform I/O transfers on low memory pages in order to free page frames.
__GFP_FS	Allowed to perform filesystem-dependent operations
__GFP_ZERO	The page frame returned, if any, must be filled with zeros.
GFP_ATOMIC	__GFP_HIGH
GFP_KERNEL GPF_USER	__GFP_WAIT __GFP_IO __GFP_FS
GFP_HIGHMEM	__GFP_WAIT __GFP_IO __GFP_FS __GFP_HIGHMEM

Zone Allocator (4)

- Components of the zone allocator



Buddy System (1)



■ Buddy algorithm

- Treats physical memory as a collection of 2^n -page-sized blocks aligned on 2^n -page boundaries.
- To allocate a block of a given order,
 - If a block is found at the specified order, it is allocated immediately.
 - If a block of higher order must be used, divide the larger block into two $2^{\text{order}-1}$ blocks, add the lower half to the appropriate freelist, and allocate the memory from the upper half, executing this step recursively.
- When freeing a block,
 - If the block has a free buddy block, combine the two blocks into a single free block; and this process is performed recursively if necessary.

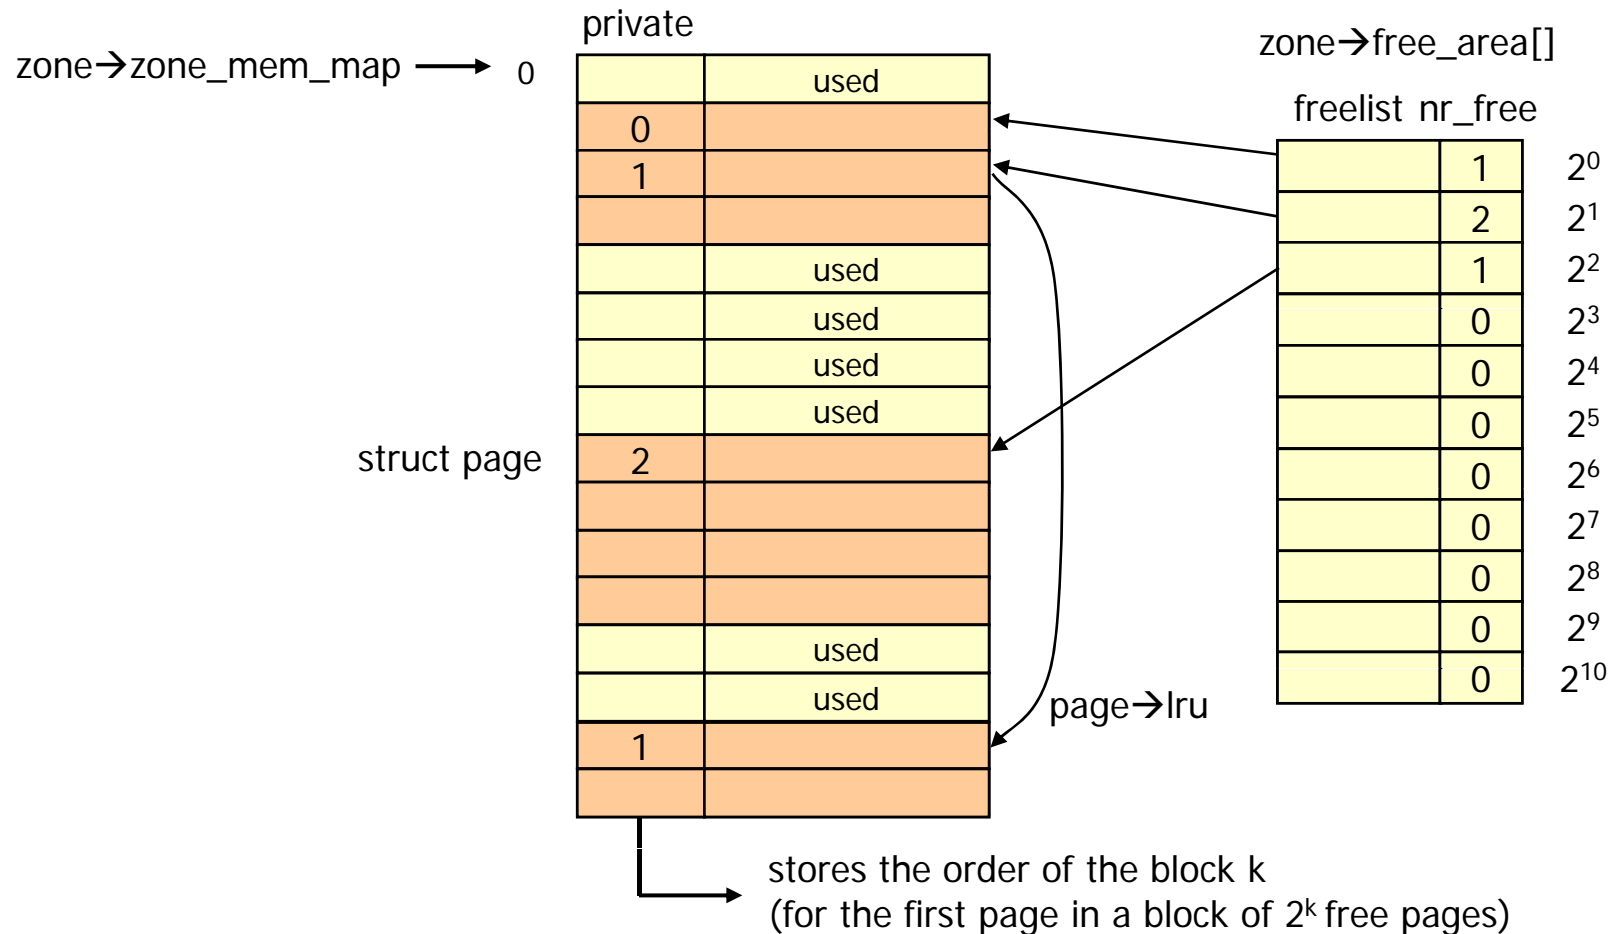
Buddy System (2)



- **Implementation** <mm/page_alloc.c>
 - A different buddy system for each zone.
 - All free page frames are grouped into 11 lists of blocks
 - Groups of 1, 2, 4, 8, 16, 32, 64, 128, 256, 512, and 1024 contiguous page frames
 - struct page ***__rmqueue**(struct zone *zone, unsigned int order);
 - Buddy allocation
 - void **__free_one_page**(struct page *page, struct zone *zone, unsigned int order);
 - Buddy deallocation
 - void **page_is_buddy**(struct page *page, int order);
 - Checks whether the block is free and it's a buddy.

Buddy System (3)

■ Implementation (cont'd)



Per-CPU Page Frame Cache (1)

■ Motivation

- The kernel often requests and releases single page frames.
- Prepare some pre-allocated page frames to be used for single memory requests issued by the local CPU.
- Hot cache
 - Page frames whose contents are likely to be included in the CPU's hardware cache.
 - Taking a page frame from the hot cache is beneficial for system performance.
- Cold cache
 - Useful if the page frame is going to be filled with a DMA operation

Per-CPU Page Frame Cache (2)

■ **struct per_cpu_pages** <linux/mmzone.h>

int	count;	Number of pages frame in the cache
int	low;	Low watermark for cache replenishing
int	high;	High watermark for cache depletion
int	batch;	Number of page frames to be added or subtracted from the cache
struct list_head	list;	List of descriptors of the page frames included in the cache

■ **struct per_cpu_pageset** <linux/mmzone.h>

- struct per_cpu_pageset {
 struct per_cpu_pages pcp[2];
};
- pcp[0]: hot cache, pcp[1]: cold cache

Per-CPU Page Frame Cache (3)

■ Managing per-CPU cache

- If (count < low), replenish the cache by allocating batch single page frames from the buddy system.
- If (count > high), release to the buddy system batch single frames in the cache.
- The kernel always assumes the freed page frame is hot.
- The cold cache is replenished when the low watermark has been reached.
- struct page * buffered_rmqueue (struct zonelist *zonelist, struct zone *zone, int order, gfp_t gfp_flags);
- void free_hot_cold_page (struct page *page, int cold);

Slab Allocator (1)



■ Slab allocator

- Developed by Sun Microsystems for Solaris 2.4 in 1994.
- The kernel functions tend to request memory areas of the same type repeatedly.
- The slab allocator does not discard the objects that have been allocated and then released but saves them in memory.
- When a new object is then requested, it can be taken from memory without having to be reinitialized.
- The slab allocator works on top of the buddy system.

Slab Allocator (2)



■ Caches

- `struct kmem_cache` <mm/slab.c>
- The slab allocator groups objects into caches.
- Each cache is a “store” of objects of the same type.
- General caches
 - `kmem_cache`: objects are the cache descriptors of the remaining caches used by the kernel.
 - General-purpose caches for 13 geometrically distributed sizes
 - » 32, 64, 128, 256, 512, 1K, 2K, 4K, 8K, 16K, 32K, 64K, 128K
 - » Two caches for each size: one for normal, the other for DMA
 - Initialized in `kmem_cache_init()`
- Specific caches
 - For specific objects used frequently by the kernel.
 - Created by `kmem_cache_create()`

Slab Allocator (3)

cafe.kaist.ac.kr - PuTTY

names_cache	7	7	4096	1	1 : tunables	24	12	8 : slabdata	7	7	0
idr_layer_cache	299	319	136	29	1 : tunables	120	60	8 : slabdata	11	11	0
buffer_head	104419	146736	52	72	1 : tunables	120	60	8 : slabdata	2038	2038	0
mm_struct	164	180	448	9	1 : tunables	54	27	8 : slabdata	20	20	0
vm_area_struct	12173	12364	88	44	1 : tunables	120	60	8 : slabdata	281	281	0
fs_cache	159	236	64	59	1 : tunables	120	60	8 : slabdata	4	4	0
files_cache	149	189	448	9	1 : tunables	54	27	8 : slabdata	21	21	0
signal_cache	193	220	384	10	1 : tunables	54	27	8 : slabdata	22	22	0
sighand_cache	193	198	1344	3	1 : tunables	24	12	8 : slabdata	66	66	0
task_struct	221	228	1360	3	1 : tunables	24	12	8 : slabdata	76	76	0
anon_vma	5342	5588	12	254	1 : tunables	120	60	8 : slabdata	22	22	0
pgd	145	145	4096	1	1 : tunables	24	12	8 : slabdata	145	145	0
size-131072(DMA)	0	0	131072	1	32 : tunables	8	4	0 : slabdata	0	0	0
size-131072	0	0	131072	1	32 : tunables	8	4	0 : slabdata	0	0	0
size-65536(DMA)	0	0	65536	1	16 : tunables	8	4	0 : slabdata	0	0	0
size-65536	1	1	65536	1	16 : tunables	8	4	0 : slabdata	1	1	0
size-32768(DMA)	0	0	32768	1	8 : tunables	8	4	0 : slabdata	0	0	0
size-32768	0	0	32768	1	8 : tunables	8	4	0 : slabdata	0	0	0
size-16384(DMA)	0	0	16384	1	4 : tunables	8	4	0 : slabdata	0	0	0
size-16384	8	16	16384	1	4 : tunables	8	4	0 : slabdata	8	16	0
size-8192(DMA)	0	0	8192	1	2 : tunables	8	4	0 : slabdata	0	0	0
size-8192	230	230	8192	1	2 : tunables	8	4	0 : slabdata	230	230	0
size-4096(DMA)	0	0	4096	1	1 : tunables	24	12	8 : slabdata	0	0	0
size-4096	163	163	4096	1	1 : tunables	24	12	8 : slabdata	163	163	0
size-2048(DMA)	0	0	2048	2	1 : tunables	24	12	8 : slabdata	0	0	0
size-2048	634	642	2048	2	1 : tunables	24	12	8 : slabdata	321	321	0
size-1024(DMA)	0	0	1024	4	1 : tunables	54	27	8 : slabdata	0	0	0
size-1024	340	340	1024	4	1 : tunables	54	27	8 : slabdata	85	85	0
size-512(DMA)	0	0	512	8	1 : tunables	54	27	8 : slabdata	0	0	0
size-512	600	656	512	8	1 : tunables	54	27	8 : slabdata	82	82	0
size-256(DMA)	0	0	256	15	1 : tunables	120	60	8 : slabdata	0	0	0
size-256	844	930	256	15	1 : tunables	120	60	8 : slabdata	62	62	0
size-128(DMA)	0	0	128	30	1 : tunables	120	60	8 : slabdata	0	0	0
size-128	2018	2250	128	30	1 : tunables	120	60	8 : slabdata	75	75	0
size-64(DMA)	0	0	64	59	1 : tunables	120	60	8 : slabdata	0	0	0
size-32(DMA)	0	0	32	113	1 : tunables	120	60	8 : slabdata	0	0	0
size-64	5069	5251	64	59	1 : tunables	120	60	8 : slabdata	89	89	0
size-32	3273	3616	32	113	1 : tunables	120	60	8 : slabdata	32	32	0
kmem_cache	150	150	128	30	1 : tunables	120	60	8 : slabdata	5	5	0

[dev1:/home/jinsoo-22]

Slab Allocator (4)



■ Slabs

- `struct slab` `<mm/slab.c>`
- Each slab consists of one or more contiguous page frames.
- Each slab contains both allocated and freed objects.
- External slab descriptor
 - Slab descriptor is stored outside the slab. (in general cache)
- Internal slab descriptor
 - Slab descriptor is stored inside the slab.
 - When the object size is smaller than 512B, or when internal fragmentation leaves enough space.

Slab Allocator (5)



■ Objects

- Objects consist of a set of data structures and methods called the constructor and the destructor.
- Each object has a short descriptor of type `kmem_bufctl_t` <mm/slab.c>
- An object descriptor is meaningful only when the object is free.
 - It contains the index of the next free object in the slab.
 - The last element is marked by `BUFCTL_END` (0xffff).
- External object descriptors vs. internal object descriptors

Slab Allocator (6)



■ Slab coloring

- Objects that have the same offset within different slabs will, with a relatively high probability, end up mapped in the same cache line.
- The slab allocator assigns different arbitrary values called **colors** to the slabs.
- Slabs having different colors store the first object in different memory locations, while satisfying the alignment constraint.
- Coloring works only when there is enough space inside the slab.

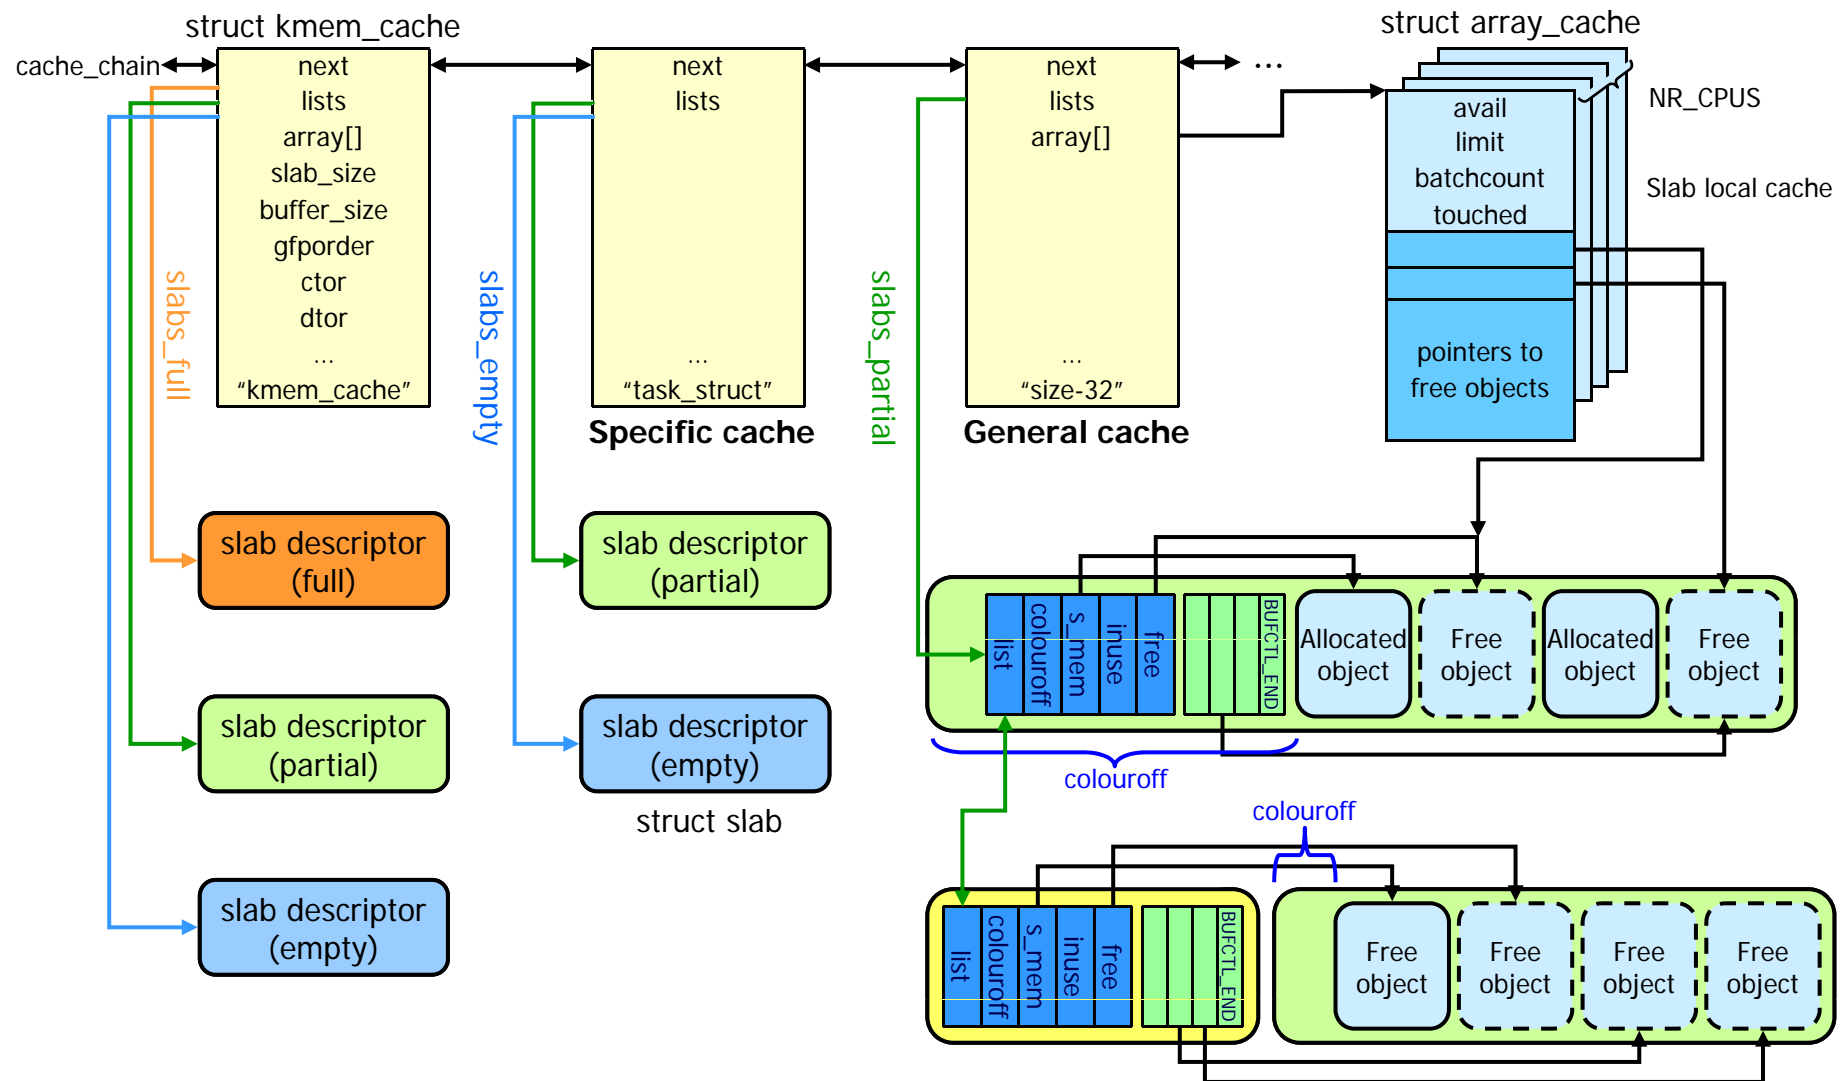
Slab Allocator (7)



■ Slab local cache

- `struct array_cache` <mm/slab.c>
- Per-CPU data structure consisting of a small array of pointers to freed objects.
- Reduce spin lock contention among processors and make better use of the hardware caches.
- The cache size depends on the size of the objects
 - 1 ~ 120 pointers
- Local caches are refilled or emptied from the slab if needed.

Slab Allocator (8)



Slab Allocator (9)



■ Using specific caches

- `struct kmem_cache *kmem_cache_create(const char *name, size_t size, size_t align, unsigned long flags, void (*ctor)(void *, struct kmem_cache *, unsigned long), void (*dtor)(void *, struct kmem_cache *, unsigned long));`
- `void *kmem_cache_alloc(struct kmem_cache *cachep, int flags);`
- `void kmem_cache_free(struct kmem_cache *cachep, void *objp);`

■ Using general caches

- `void *kmalloc(size_t size, int flags);`
- `void kfree(const void *objp);`

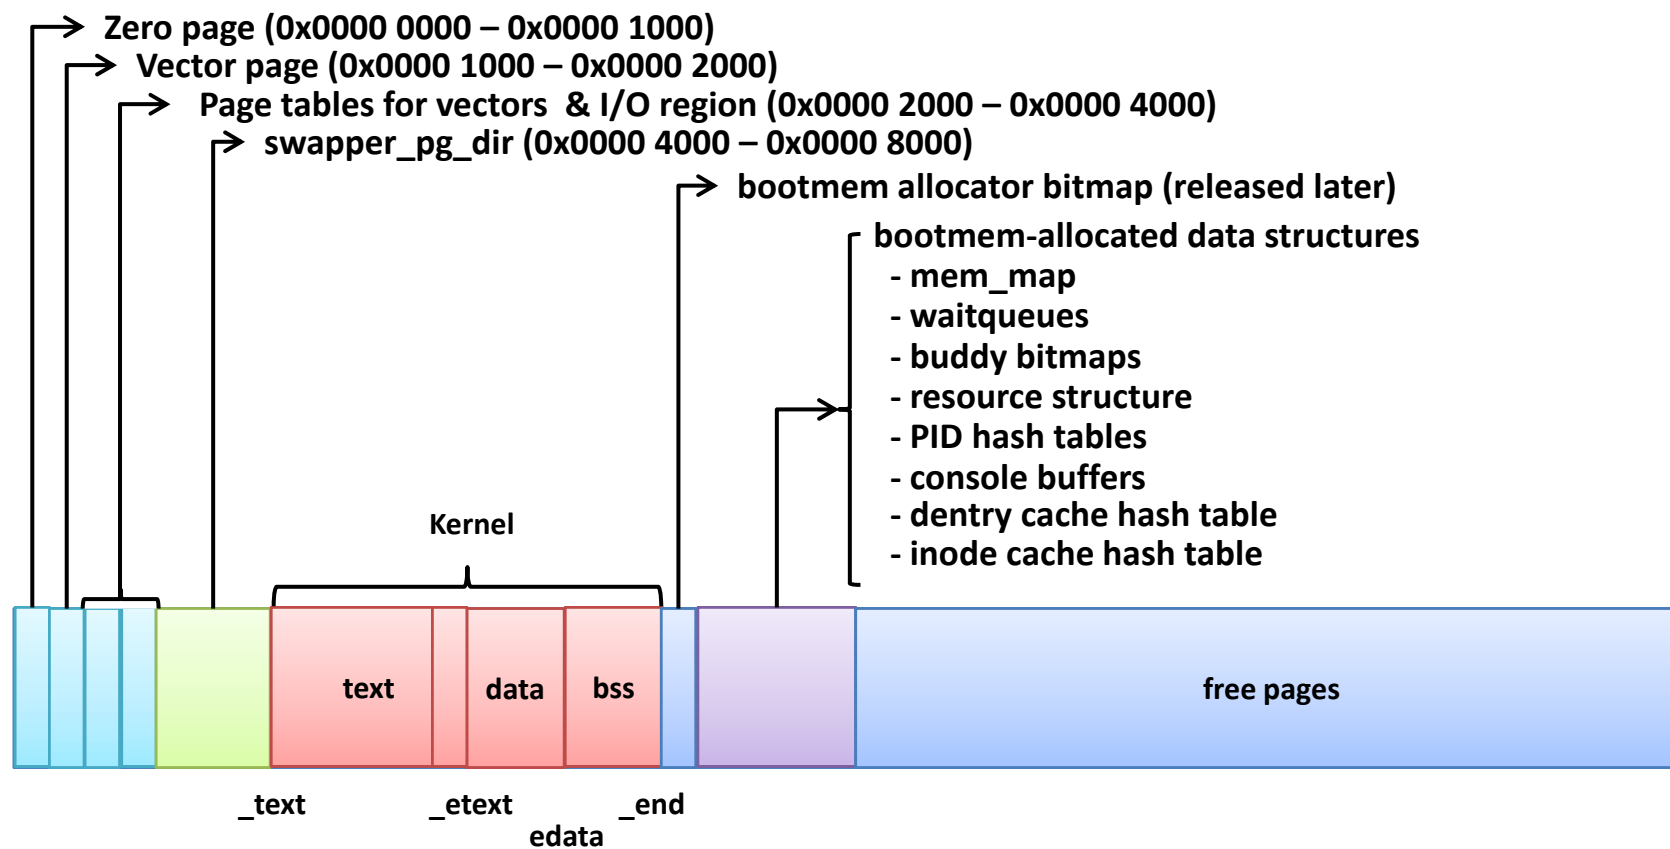
Noncontiguous Areas



- **Noncontiguous memory areas**
 - Virtually contiguous, but physically noncontiguous
 - Need to modify kernel page tables
 - Can make use of high memory page frames
- **Using a noncontiguous memory area**
 - void ***vmalloc**(unsigned long size);
 - void **vfree**(void *addr);

Where Are My Pages? (1)

- After `mem_init()`;
 - Linux kernel 2.6.10 on ARM9



Where Are My Pages? (2)

- **Free pages**

- Buddy
- PCP (Per-CPU Page frame cache)

- **Slab**

- For kernel data structures
- Specific & general

- **Vmalloc**

- **Kernel stacks**

- **Page tables**

- Kernel page tables
- User page tables

Where Are My Pages? (3)



- **Anonymous pages**
 - Private vs. shared
- **File-mapped pages**
 - Page caches
 - Buffer caches
 - Mapped pages (private vs. shared)
- **Modules**
- **Device drivers**
 - DMA areas, buffers, etc.
- **Other page-level data structures**
 - PID bitmap, pipes, etc.

Summary



■ Memory allocation in the kernel

- Page-aligned
 - `alloc_pages()`
- Specific caches from slab
 - `kmem_cache_alloc()`
- General caches from slab
 - `kmalloc()`
- From non-contiguous memory area
 - `vmalloc()`