



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Flash Memory

**KAIST**

# Memory Types




## FLASH

- High-density
- Low-cost
- High-speed
- Low-power
- High reliability

## DRAM

- High-density
- Low-cost
- High-speed
- High-power



## EPROM

- Non-volatile
- High-density
- Ultraviolet light for erasure

## EEPROM

- Non-volatile
- Lower reliability
- Higher cost
- Lowest density
- Electrically byte-erasable

## ROM

- High-density
- Reliable
- Low-cost
- Suitable for high production with stable code

*Source: Intel Corporation.*

# Flash Memory Characteristics

## ■ Erase-before-write

- Read
- Write or Program – change state from 1 to 0
- Erase – change state from 0 to 1

## ■ Bulk Erase

- Program unit:
  - NOR: byte or word
  - NAND: sector or page
- Erase unit: block

1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---

↓ write  
(program)

1	1	0	1	1	0	1	0
---	---	---	---	---	---	---	---

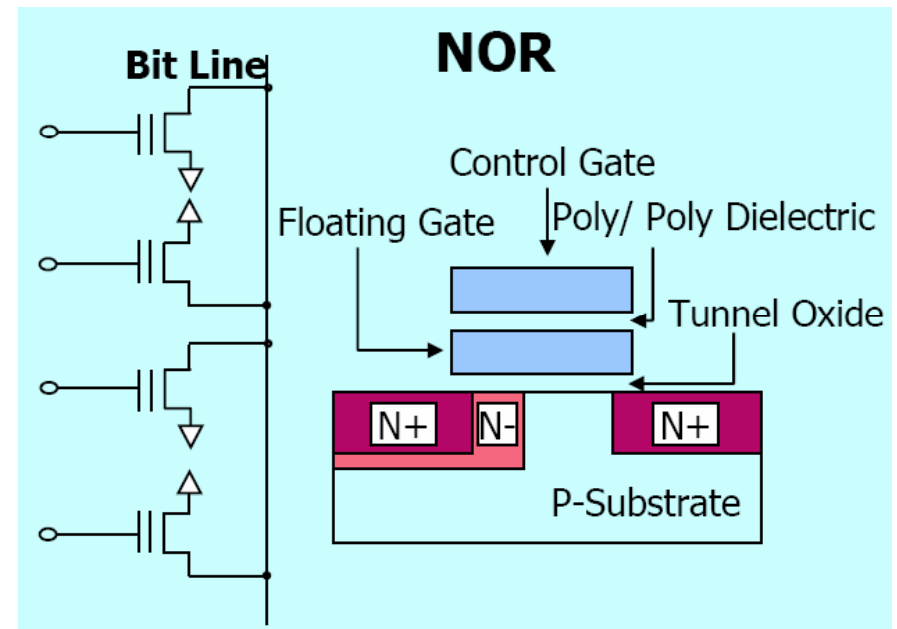
↓ erase

1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---

# NOR Flash

## ■ NOR Flash

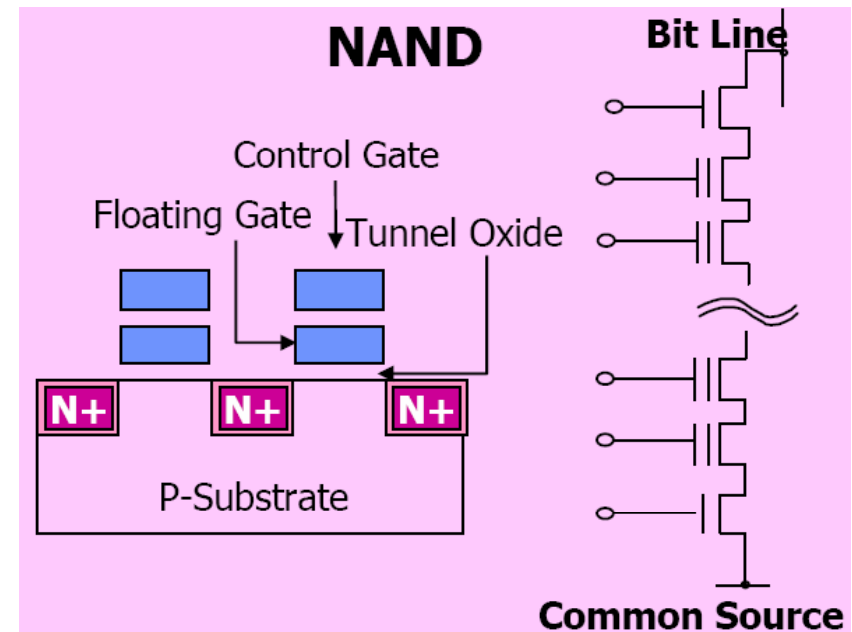
- Random, direct access interface
- Fast random reads
- Slow erase and write
- Mainly for code storage
- Intel, Spansion, STMicro, ...




# NAND Flash

## ■ NAND Flash

- I/O mapped access
- Smaller cell size
- Lower cost
- Smaller size erase blocks
- Better performance for erase and write
- Mainly for data storage
- Samsung,  
Toshiba,  
Hynix, ...



# NOR vs. NAND Flash (1)



Source: Toshiba



# NOR vs. NAND Flash (2)

## Mass Storage-NAND



### Memory Cards

(mobile computers)

### Solid-State Disk

(rugged & reliable storage)



### Digital Camera

(still & moving pictures)



### Voice/Audio Recorder

(near CD quality)

- Low Cost and High Density
- Good P/E Cycling Endurance

## Code Memory-NOR



### BIOS/Networking

(PC/router/hub)



### Telecommunications

(switcher)

### Cellular Phone

(code & data)



### POS / PDA / PCA

(code & data)

- Fast Random Access
- XIP

Source: Samsung Electronics



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# NAND Flash Memory


**KAIST**



# NAND Technology (1)

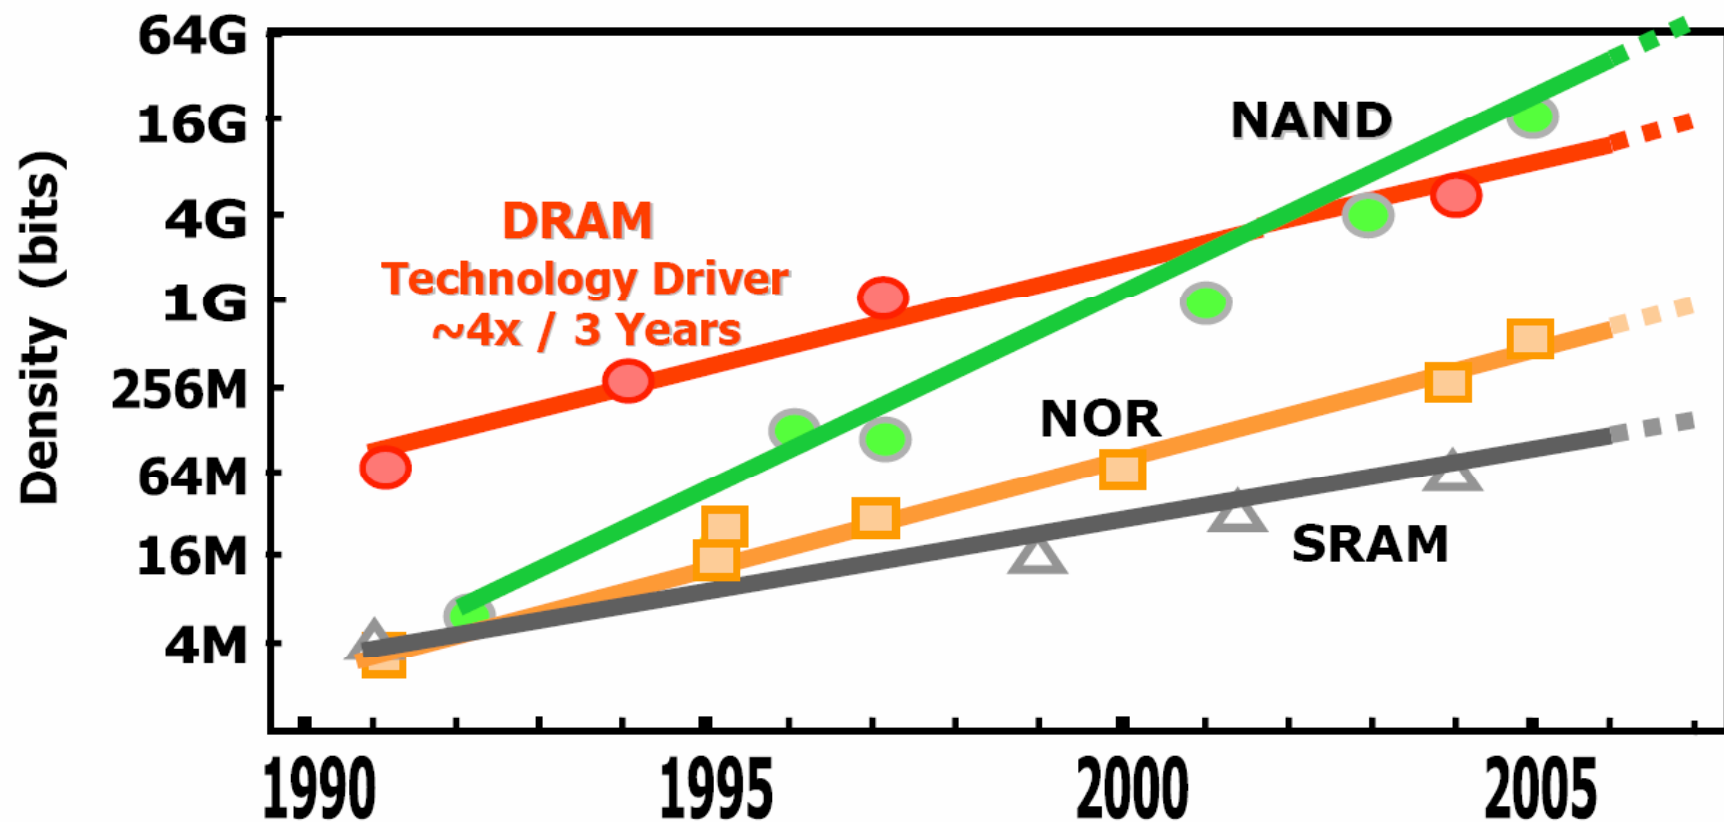
## ■ Hwang's Law

- The density of the top-of-the-line flash memory chips will double every 12 months



# NAND Technology (2)


## ■ Density Growth



Source: Samsung Electronics

# NAND Flash Architecture (1)

## ■ 2Gb NAND Flash Device Organization




Source: Micron Technology, Inc.

# NAND Flash Architecture (2)


## ■ NAND Flash Interface

Pin Name	Pin Function
$I/O_0 \sim I/O_7$	<b>DATA INPUTS/OUTPUTS</b> Command, address, or data
CLE	<b>COMMAND LATCH ENABLE</b>
ALE	<b>ADDRESS LATCH ENABLE</b>
$\overline{CE}$	<b>CHIP ENABLE</b>
$\overline{RE}$	<b>READ ENABLE</b>
$\overline{WE}$	<b>WRITE ENABLE</b>
$\overline{WP}$	<b>WRITE PROTECT</b>
$R/\overline{B}$	<b>READY/BUSY OUTPUT</b>
VCC	<b>POWER</b>
VSS	<b>GROUND</b>



# NAND Operations (1)

## ■ READ




Source: Micron Technology, Inc.



# NAND Operations (2)

## ■ PROGRAM




Source: Micron Technology, Inc.



# NAND Operations (3)

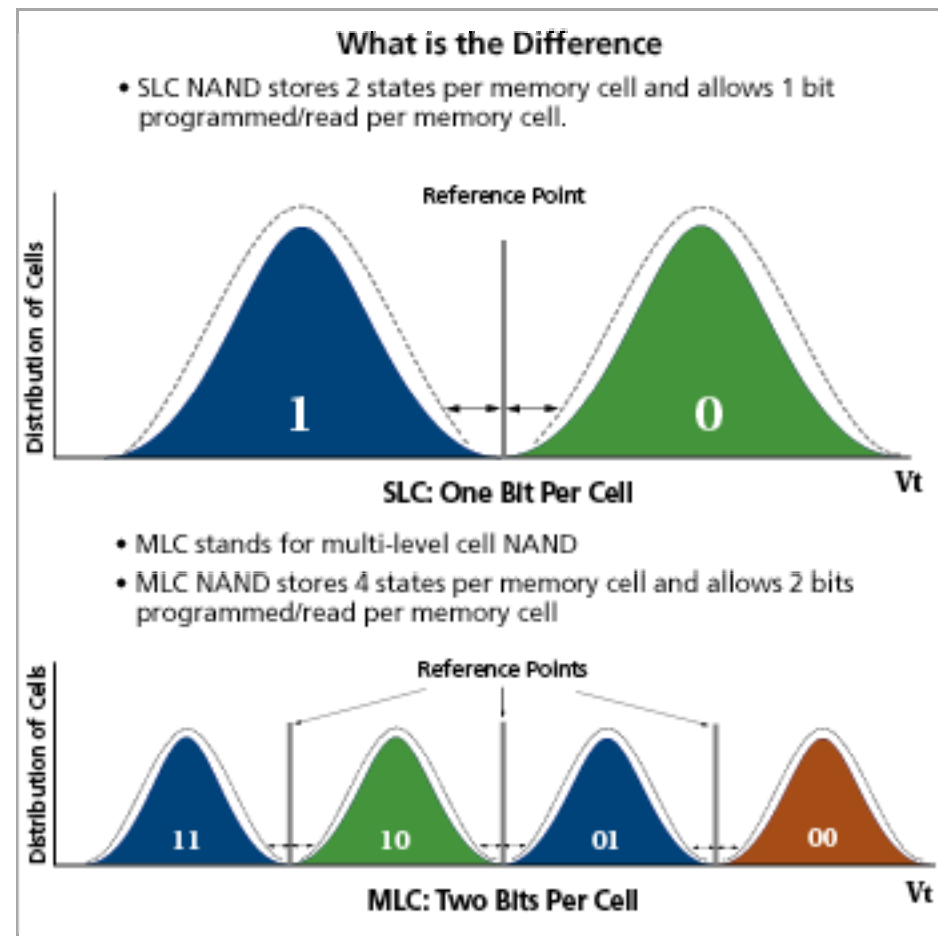
## ■ BLOCK ERASE



Source: Micron Technology, Inc.

# NAND Flash Types (1)

- **SLC NAND Flash**
  - Small block ( $\leq 1\text{Gb}$ )
  - Large block ( $\geq 1\text{Gb}$ )
- **MLC NAND Flash**



Source: Micron Technology, Inc.

# NAND Flash Types (2)


	SLC NAND <sup>1</sup> (small block)	SLC NAND <sup>2</sup> (large block)	MLC NAND <sup>3</sup>
Page size (Bytes)	512+16	2,048+64	4,096+128
Pages / Block	32	64	128
Block size	16KB	128KB	512KB
$t_R$	15 $\mu$ s (max)	20 $\mu$ s (max)	50 $\mu$ s (max)
$t_{PROG}$	200 $\mu$ s (typ) 500 $\mu$ s (max)	200 $\mu$ s (typ) 700 $\mu$ s (max)	600 $\mu$ s (typ) 1,200 $\mu$ s (max)
$t_{BERS}$	2 ms (typ) 3 ms (max)	1.5 ms (typ) 2 ms (max)	3 ms (typ)
NOP	1 (main), 2 (spare)	4	1
Endurance Cycles	100K	100K	10K
ECC (per 512Bytes)	1 bit ECC 2 bits EDC	1 bit ECC 2 bits EDC	4 bits ECC 5 bits EDC

<sup>1</sup> Samsung K9F1208X0C (512Mb) <sup>2</sup> Samsung K9K8G08U0A (8Gb) <sup>3</sup> Micron Technology Inc.


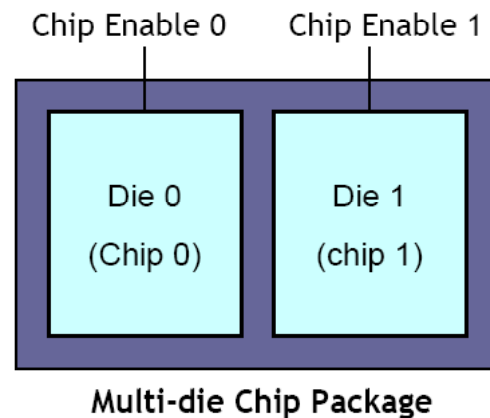
# NAND Flash Types (3)

## ■ Extended NAND Flash Architecture

### Multi-plane



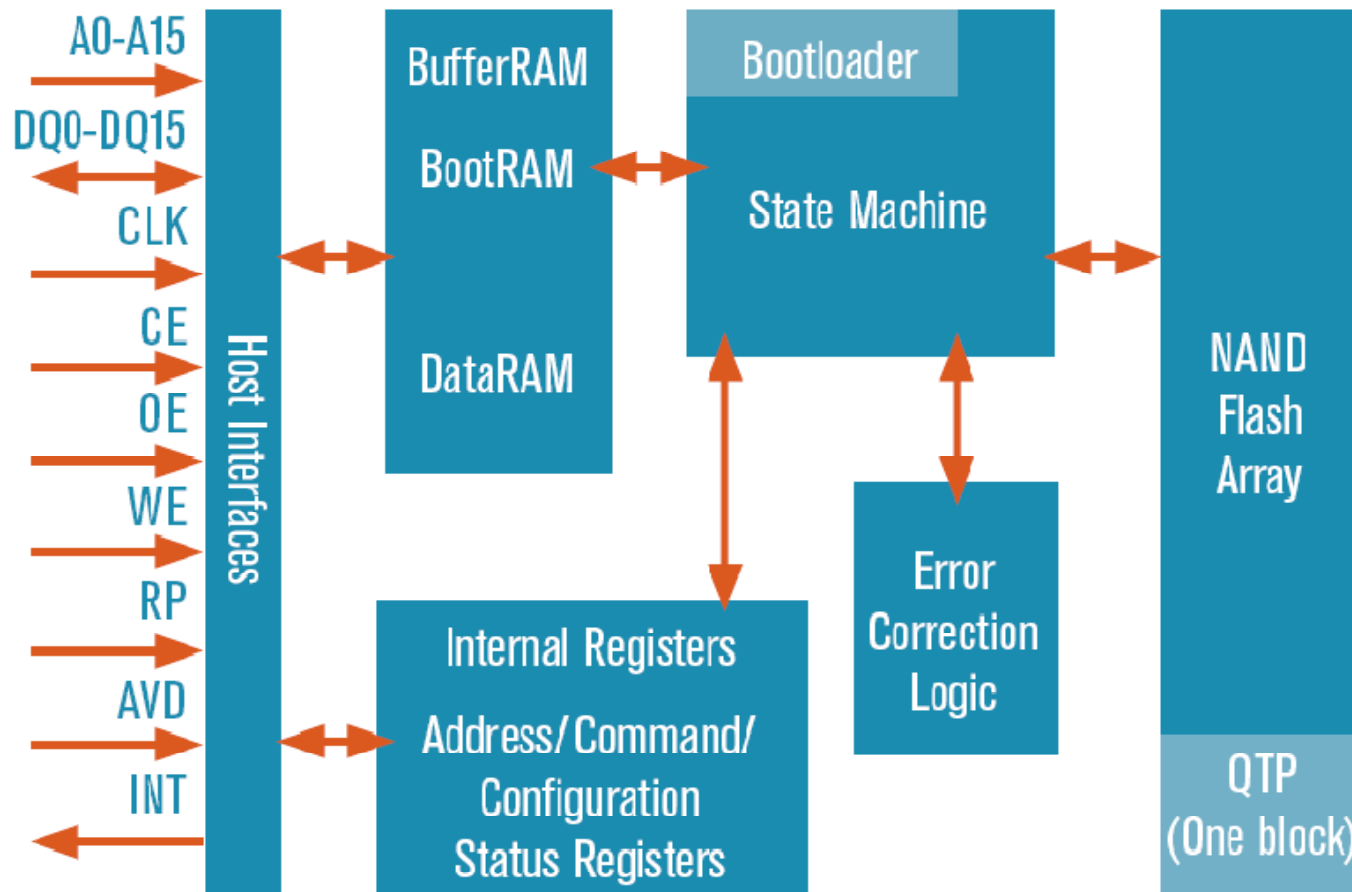
### Multi-die



Source: Zeen Info. Tech.

# NAND Flash Types (4)

- Samsung OneNAND™

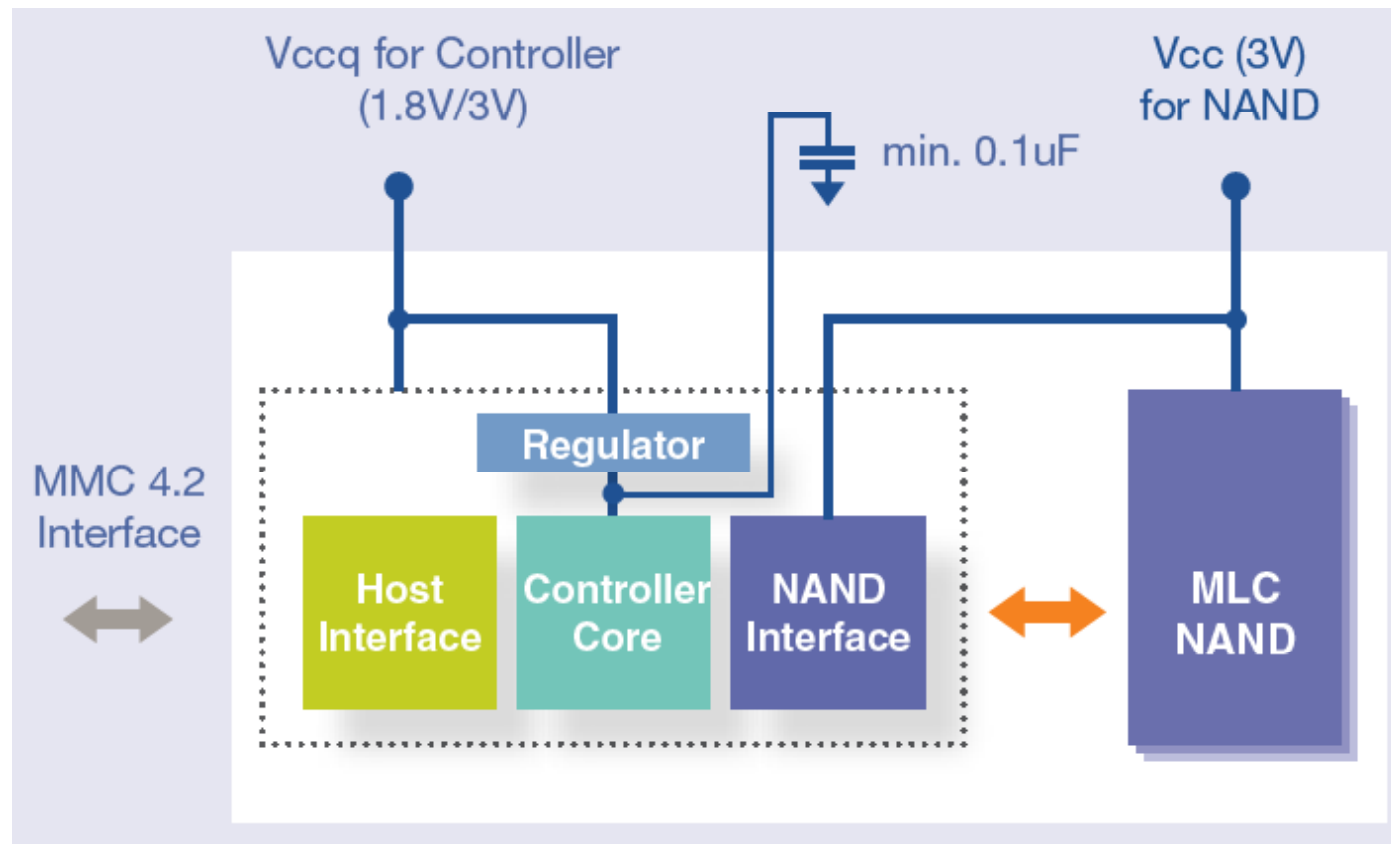


Source: Samsung's OneNAND ebrochure



# NAND Flash Types (5)

- Samsung moviNAND™

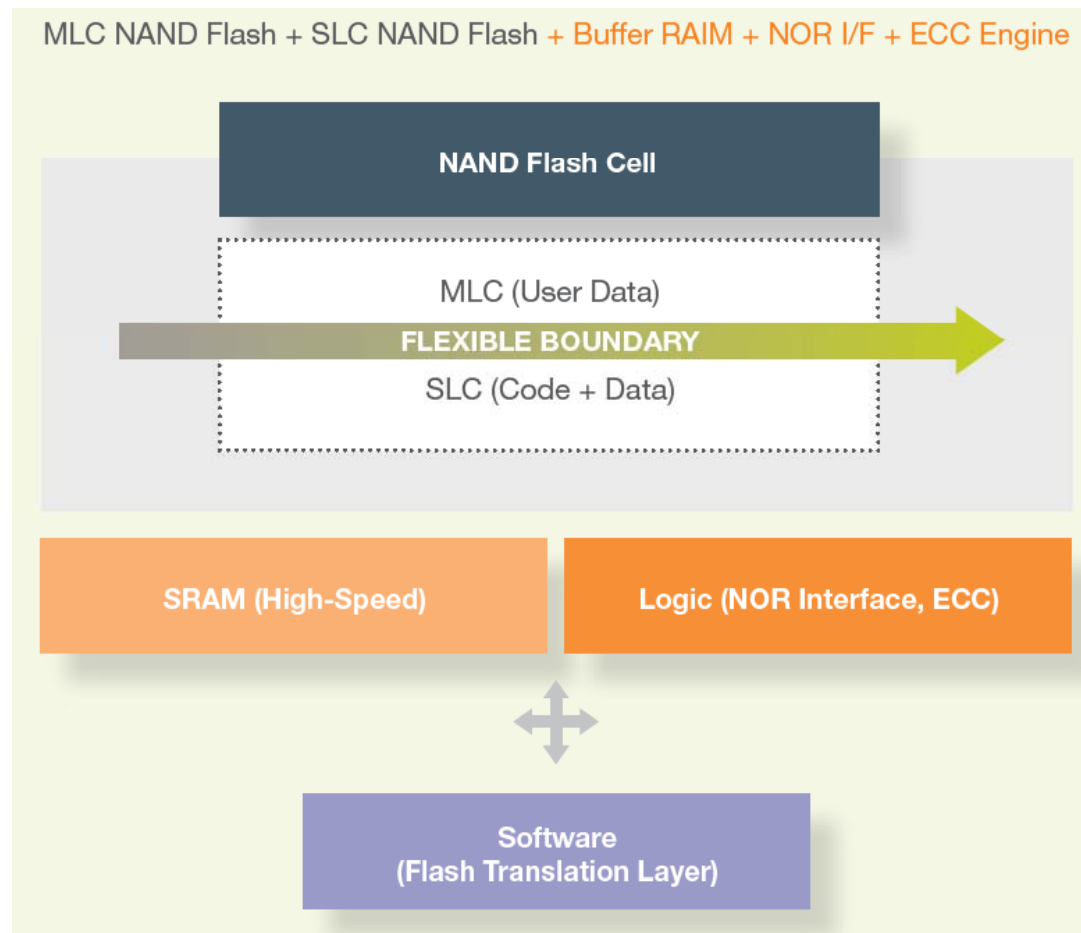


Source: Samsung Fusion Memory



# NAND Flash Types (6)

- Samsung Flex-OneNAND™



Source: Samsung Fusion Memory

# NAND Constraints (1)



## ■ Bit Errors

- Error correction codes (ECC) in spare area

## ■ Bad Blocks

- Factory-marked bad blocks
- Run-time bad blocks
- Bad block management

## ■ Limited Program/Erase Cycles

- 100K for SLCs
- 10K for MLCs
- Wear-leveling

# NAND Constraints (2)



## ■ NOP

- Partial-page programming
- 1 / sector for most SLCs (4 for 2KB page)
- 1 / page for most MLCs

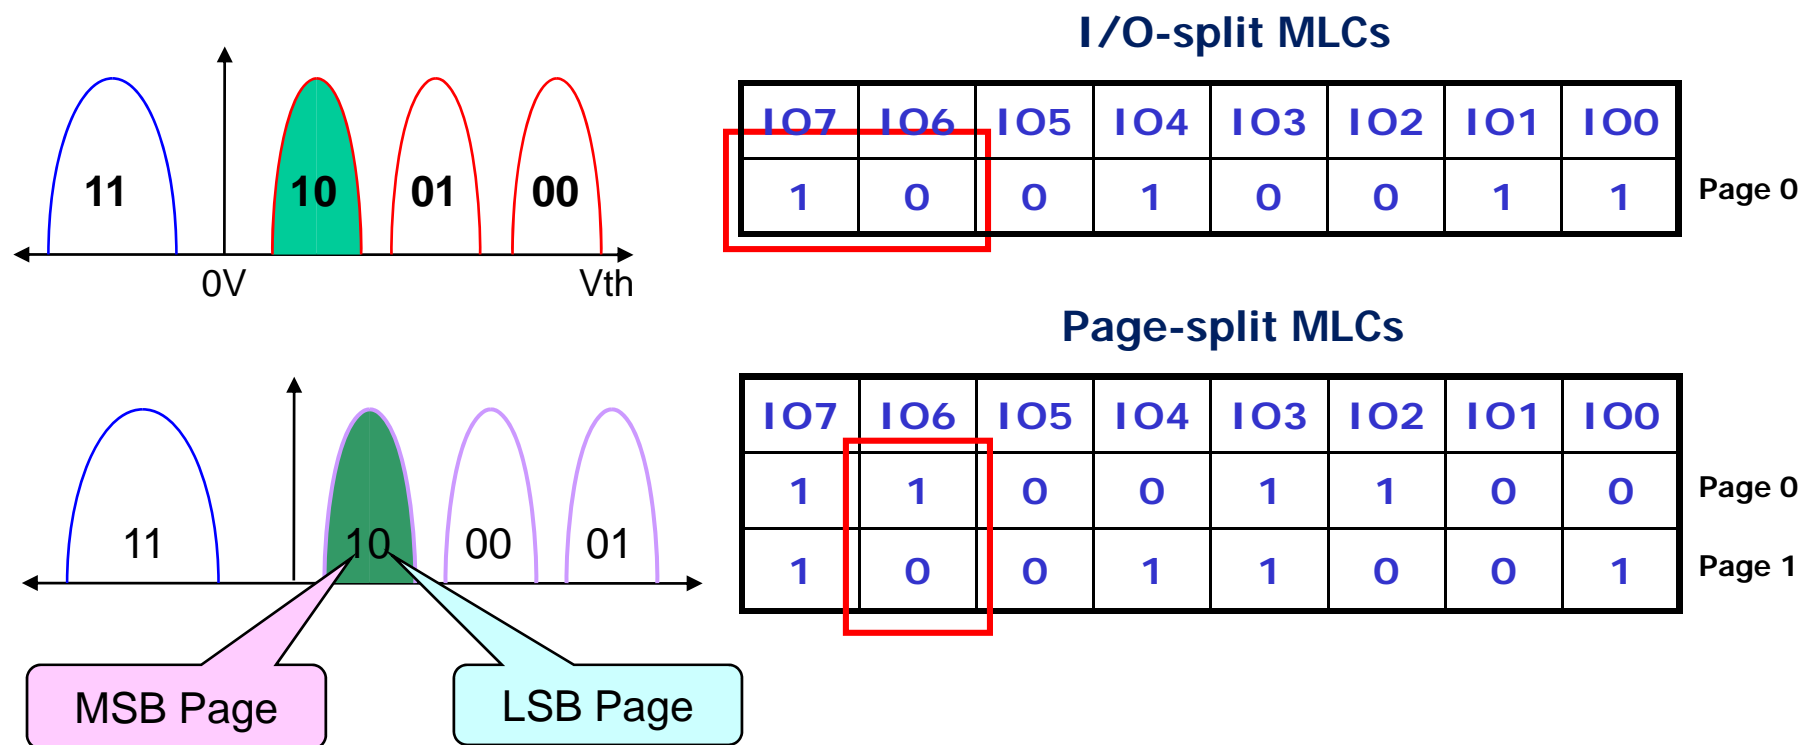
## ■ Sequential Page Programming

- For large block SLCs and MLCs
- From page 0 to page 63 for SLCs
- From page 0 to page 127 for MLCs

# NAND Constraints (3)

## ■ Pair-page Programming in MLCs

- Performance difference
- Interference



# Beauty and the Beast



- **NAND Flash memory is beauty.**
  - Small, light-weight, robust, low-cost, low-power non-volatile device
- **NAND Flash memory is a beast.**
  - Much slower program/erase operations
  - No in-place-update
  - Erase unit > write unit
  - Limited lifetime (10K~100K program/erase cycles)
  - Bad blocks, ...
- **Software support for NAND flash memory is very important for performance & reliability.**



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Memory Architectures for Mobile Devices

**KAIST**



# Requirements

## Code



Mobile



Consumer Electronics



Networking

Read	Writes	Density	Reliability
Fast Random	Medium	Small – Medium	No bad bits

## Data



Cards



MP3



USB Drives

Read	Writes	Density	Reliability
Fast Sequential	Fast	Large	Bad bits allowed

Source: “Non-Volatile Memories”, Intel Corp.


# NOR XIP

## ■ Pros

- Simple, easy to design
- Execute-In-Place (XIP)
- Predictable read latency
- Code + Data in NOR
- Firmware upgrades

## ■ Cons

- Slow read speed
- Much slower write speed
- The high cost of NOR



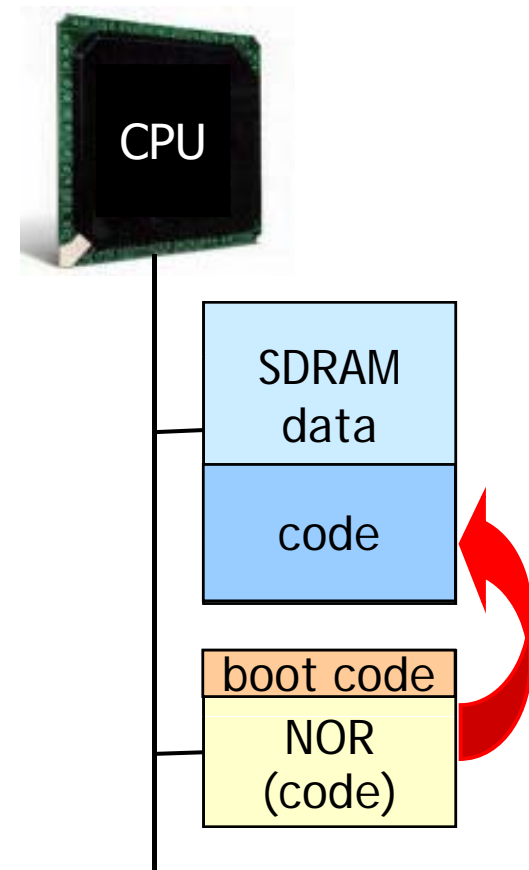
# NOR Shadowing

## ■ Pros

- Faster read and write
- Easy boot-up
- Use a relatively pricey NOR only to boot up the system
- Code can be compressed

## ■ Cons

- Larger DRAM needed
- Require more design time
- Not energy efficient




# NAND Shadowing

## ■ Pros

- Faster read and write
- Cost effective
- NAND for both code and data storage

## ■ Cons

- Require a special boot mechanism
- Extensive ECC for NAND
- Larger DRAM needed
- Require more design time
- Not energy efficient




# Hybrid NAND Shadowing

## ■ Pros

- Much faster read and write speed
- ECC embedded
- Cost effective
- NAND for both code and data storage

## ■ Cons

- Larger DRAM needed
- Not energy efficient






# NAND Demand Paging

## ■ Pros

- Less DRAM required
- Low cost
- Energy efficient
- NAND for both code and data storage

## ■ Cons

- Require MMU-enabled CPU
- Unpredictable read latency
- Complex to design and test









*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Flash Translation Layers (FTL)

**KAIST**

# FTL (1)

- Flash Cards Internals



# FTL (2)

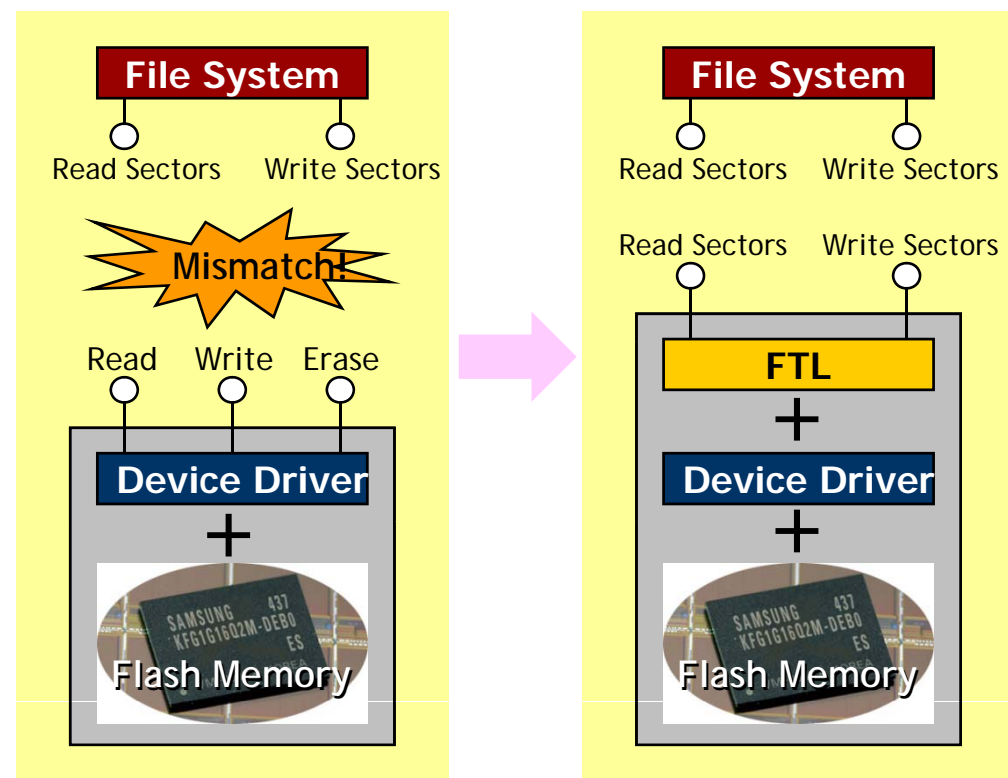
## ■ FTL

- A software layer to make NAND flash fully emulate traditional block device (e.g., disks)

## ■ Why FTL?

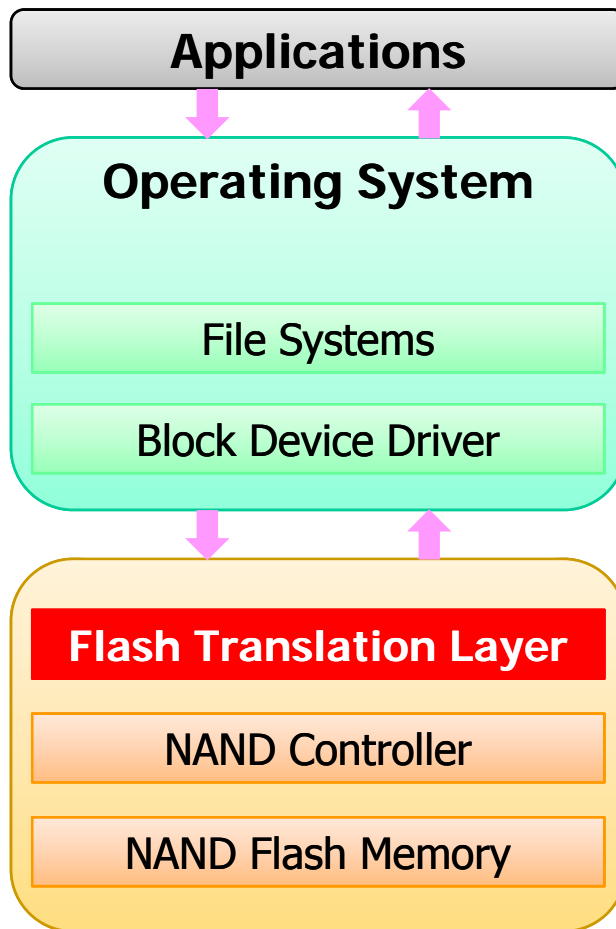
- No in-place-update
- Bulk erase
- Asymmetry in read and write speeds

Source: Zeen Info. Tech.

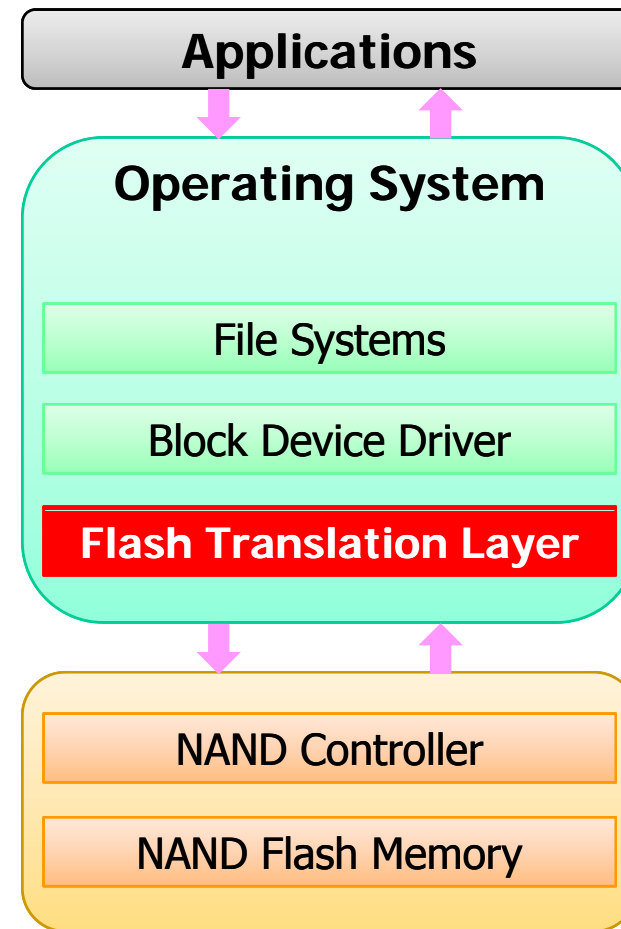


# FTL (3)

## Flash Cards, SSDs



## Embedded Flash Storage



# FTL (4)



## ■ For Performance

- Sector mapping (address translation)
- Garbage collection

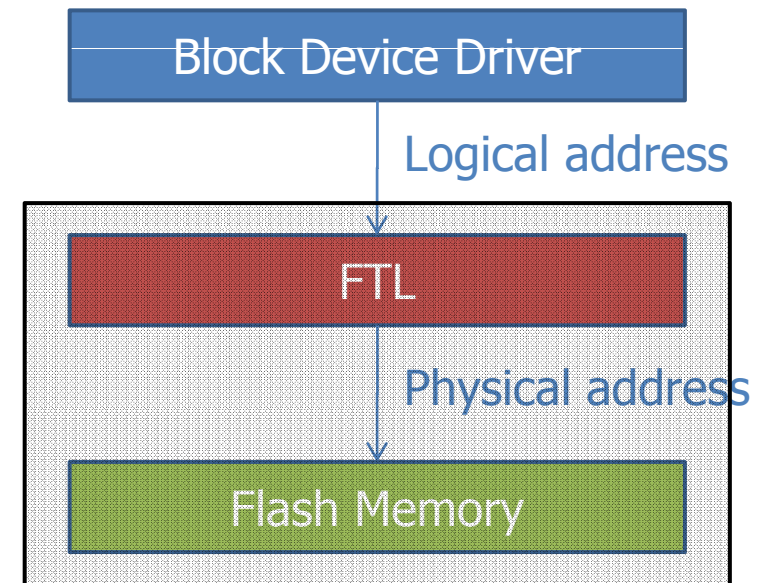
## ■ For Reliability

- Power-off recovery
- Wear-leveling
- Bad block management
- Error correction code (ECC)



# Address Translation

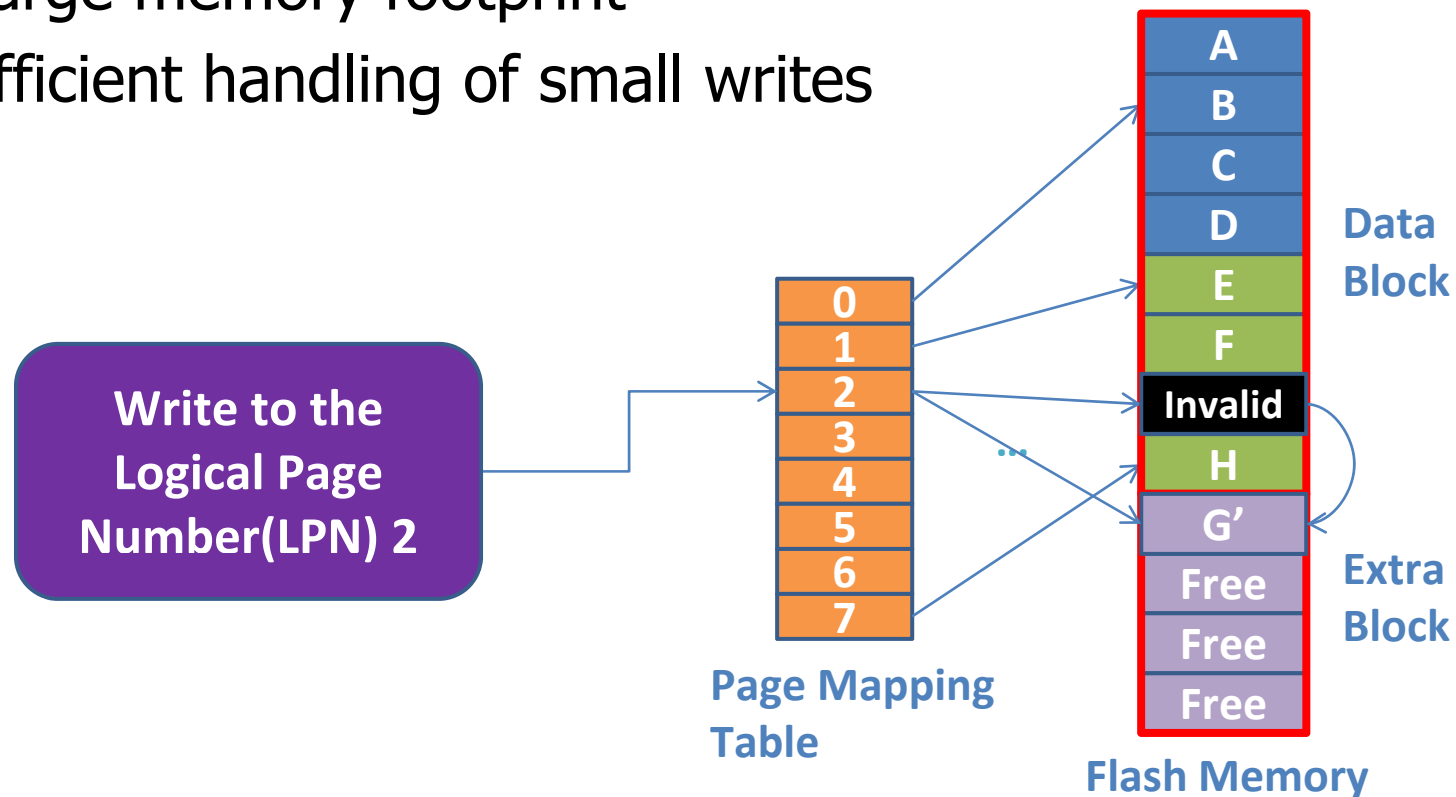
- **Mapping granularity**
  - Page mapping
  - Block mapping
  - Hybrid mapping
- **Block organization**
  - In-place
  - Out-of-place
- **Managing mapping info.**
  - Per block
  - Map block



# Page Mapping

## ■ Characteristics

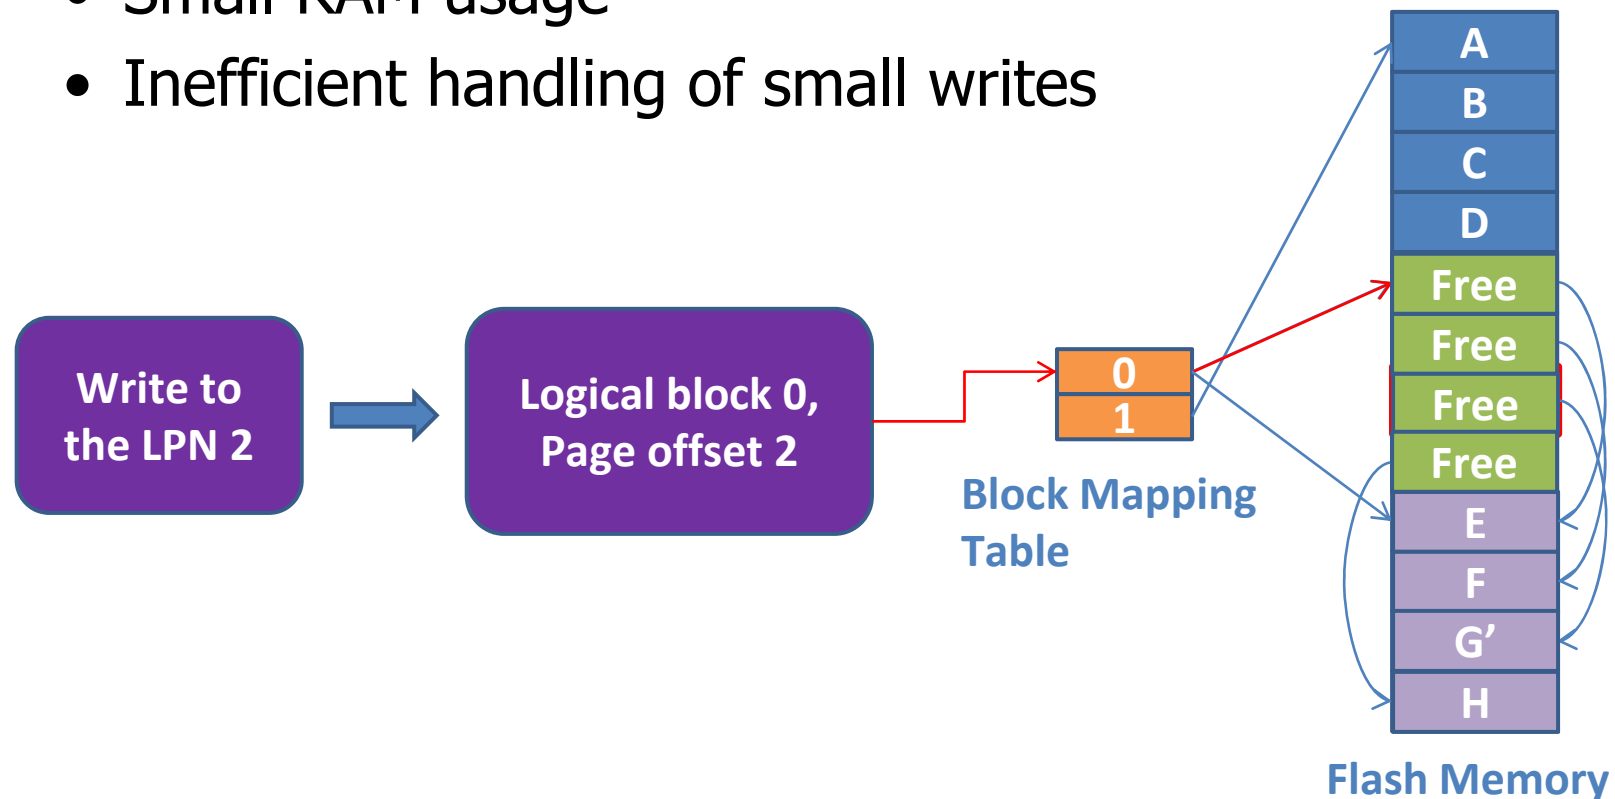
- Each table entry maps one page
- Large memory footprint
- Efficient handling of small writes



# Naïve Block Mapping

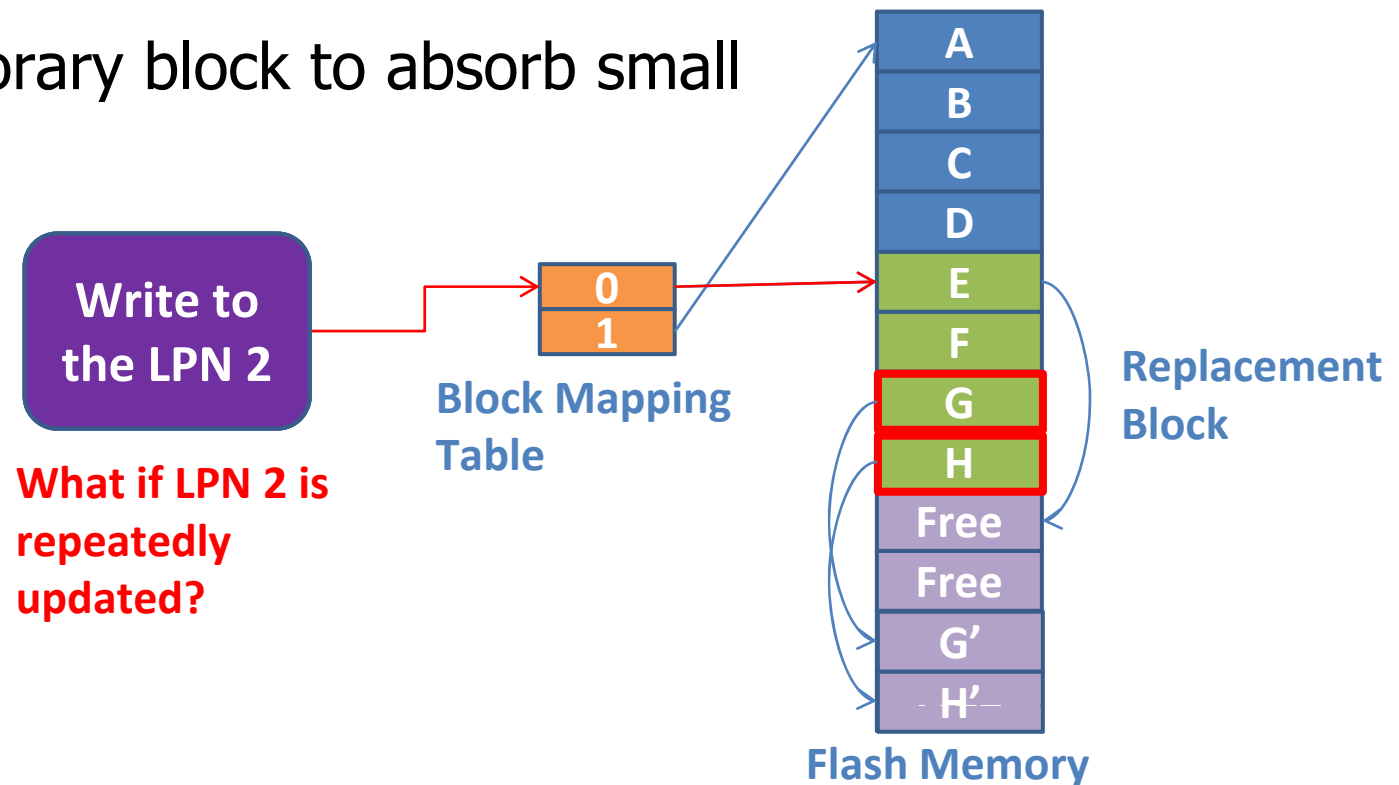
## ■ Characteristics

- Each table entry maps one block
- Small RAM usage
- Inefficient handling of small writes



# Replacement Block Scheme

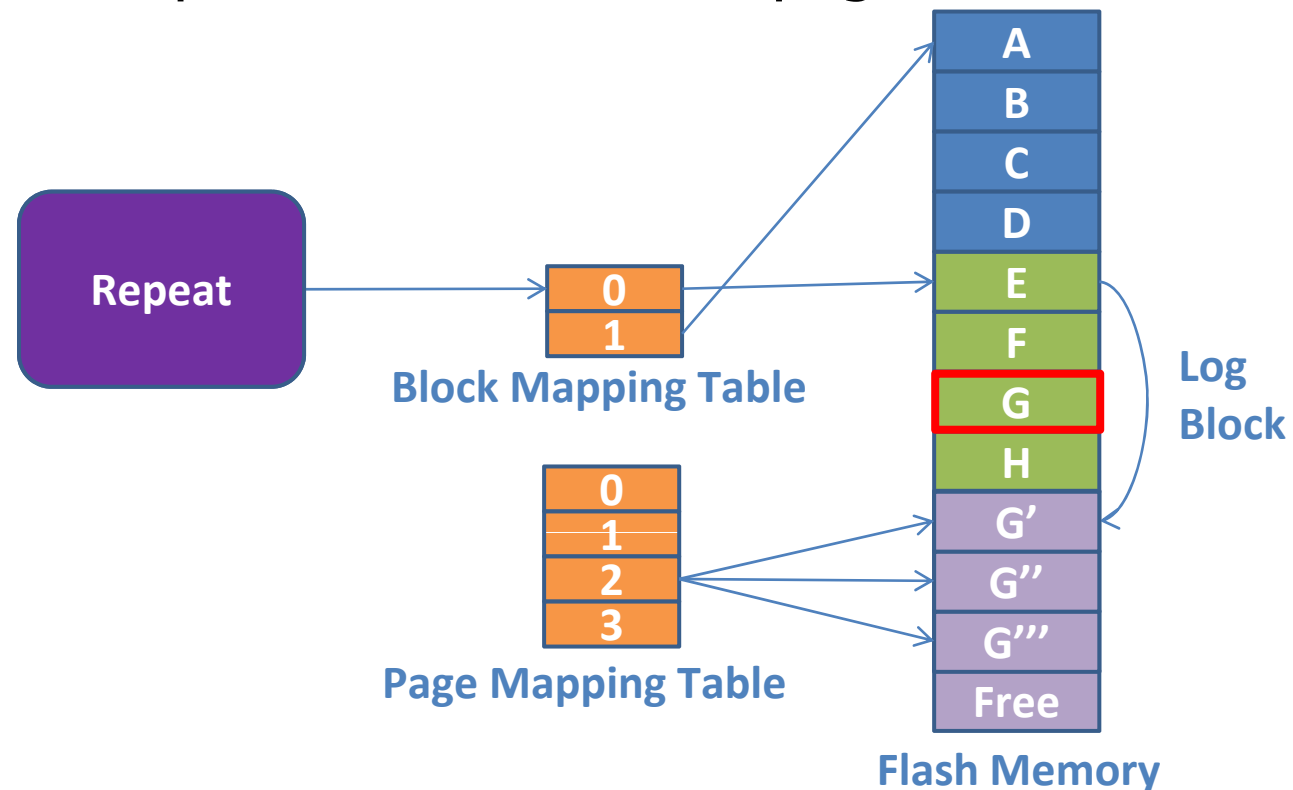
- **Data Block**
  - Storage for an ordinary data
- **Replacement Block**
  - A temporary block to absorb small writes



# Log Block Scheme (1)

## ■ Log Block

- A temporary storage for small writes
- Incremental updates from the first page






# Log Block Scheme (2)

## ■ Block Merge

- When there is no free log block to allocate.
- Allocate a free block and copy all the up-to-date pages



# FAST

## ■ Fully Associative Sector Translation (FAST)

- Sharing log blocks


## ■ Pros

- Increase log block utilization
- Reduces the number of erase operations

## ■ Cons

- Increased merge time

Flash Memory




# Superblock Scheme (1)

## ■ Superblock

- N logically adjacent blocks
- Page-level address translation within a superblock

## ■ Update Block

- A superblock shares log blocks
- More than one log blocks can be allocated to a superblock.




# Superblock Scheme (2)

## ■ Merge Operation

- Gathers hot pages into the same block

Superblock

Log  
Block



# Comparison



		Replacement block scheme	Log block scheme	FAST	Superblock scheme
<b>Data blocks</b>	Terminology	Data blocks	Data blocks	Data blocks	D-blocks
	Management scheme	In-place	In-place	In-place	Out-of-place
	The degree of sharing	1	1	1	N
<b>Update blocks</b>	Terminology	Replacement blocks	Log blocks	Random/ sequential log blocks	U-blocks
	Management scheme	In-place	Out-of-place	Out-of-place	Out-of-place
	The degree of sharing	1	1	P	N

N: superblock size, P: the number of pages in a block



# Performance (1)





## ■ Simulation Environment

- 4GB flash memory
  - Large block SLC NAND (2KB page, 128KB block)
- FTL schemes
  - Naïve block mapping
  - Replacement block
  - Log block
  - Superblock
- Workload
  - Trace from PC using NTFS


# Performance (2)

- Extra Erase and Write Operations
  - 256 extra blocks



# Performance (3)

## ■ Garbage Collection Overhead



# Bad Block Management (1)

## ■ Bad Blocks

- Factory-marked bad blocks
  - Marked “non-FF” at the first byte in the spare area (either 1st or 2nd page of every initial bad block)
  - < 2% of entire blocks
  - The first block is guaranteed to be good.
  - Should not erase or program factory-marked bad blocks.
- Run-time bad blocks
  - Bit errors
  - Endurance cycle
  - PROGRAM/ERASE failures

# Bad Block Management (2)

## ■ Issues

- Bad block table organization
  - Space
  - Bad block #, Remapped block #
  - Very sparse
- Fast lookup
  - Time
- Number of reserved blocks



# Wear-Leveling (1)



## ■ Wear-leveling

- A software technique for balancing the erase counts of physical blocks to fully utilize the lifetime of NAND flash.

	With Wear-Leveling	Without Wear-Leveling
<b>Feature</b>	Maintain erase counts in both flash and RAM	Randomly select a block at run-time
<b>Pros</b>	Maximize the lifetime up to spec. value	No performance overhead
<b>Cons</b>	Performance degradation RAM overhead	Excessive wear-out in the specific area depending on the workload

# Wear-Leveling (2)



## ■ Naïve Wear-leveling

- Use blocks in a round-robin fashion
- Performance!
- JFFS

## ■ Sophisticated Wear-leveling

- Keep erase count per block (in spare area)
- Dynamic mapping
- Consider erase counts during garbage collection
- Block swapping

# Wear-Leveling (3)



## ■ Hot-Cold Swapping

- Swap if  $\max(EC_i) - \min(EC_i) > Threshold$

## ■ CAT (Cost, Age, Times)

- Erase the block with minimizes the following score:

$$Cleaning\ Cost_i \times \frac{1}{Age_i} \times Number\ of\ Cleaning_i = \frac{\mu_i \times \varepsilon_i}{(1 - \mu_i) \times a_i(t)}$$

- $\mu_i$  : utilization (the percentage of valid data)
- $a_i(t)$  : the elapsed time since the block was erased
- $\varepsilon_i$  : the erase count

# Wear-Leveling (4)



## ■ Issues

- Maintaining erase counts
  - Per block
  - Separate area
- Memory footprint
  - Min/Max erase count
  - All the erase counts
  - Building in-memory data structure
- Performance degradation
  - Overhead due to wear-leveling

# Error Correction Codes

- ECC
  - Hardware vs. Software

Error Correction Level	Bits Required in the NAND Flash Spare Area		
	Hamming	Reed-Solomon	BCH
1	13	18	13
2	N/A	36	26
3	N/A	54	39
4	N/A	72	52
5	N/A	90	65
6	N/A	108	78
7	N/A	126	91
8	N/A	144	104
9	N/A	162	117
10	N/A	180	130

*Source: Micron Technology, Inc.*



# Power-Off Recovery



## ■ Power-Off Recovery

- For every modifying operations: write, reclaim
- Atomicity (transactional operation)
- Scan vs. checkpointing
  - Per Block
  - Map Block
- File system level vs. FTL level
- System-wide power-management scheme is necessary
  - On detecting power-off, system should notify devices drivers to safely close the current operation
- Super capacitor?

# Summary (1)



- **Small memory footprint**
  - Directly related to the SRAM size (\$\$\$)
- **Handling random writes**
- **Garbage collection**
  - Minimize valid copy during garbage collection
  - Wear-leveling
- **Wear-leveling**
- **Power-off recovery**
- **Real-time performance**

# Summary (2)



## ■ Physical Constraints

- Large block NAND
  - Sequential page programming
  - NOP 4
- MLC (Multi-Level Cell) NAND
  - NOP 1
  - Less endurance cycles ( $\sim 10K$ )
  - More bit errors
  - Most of spare area is dedicated to ECC/EDC
  - Pair-page programming