



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# File Systems

**KAIST**

# Storage: A Logical View

- Abstraction given by block device drivers:




- **Operations**

- Identify(): returns N
- Read(start sector #, # of sectors)
- Write(start sector #, # of sectors)

*Source: Sang Lyul Min (Seoul National Univ.)*

# File System Basics

- **File system: A mapping problem**
  - $\langle \text{filename, data, metadata} \rangle \rightarrow \langle \text{a set of blocks} \rangle$



# Filesystems in Linux (1)

## ■ Disk-based filesystems

- Ext2/3, ReiserFS, JFS (IBM), XFS (SGI)
- Unix variants: SYSV (System V, Coherent, Xenix), UFS (BSD, Solaris, NEXTSTEP), MINIX, Veritas VxFS (SCO Unixware)
- Microsoft: FAT (MS-DOS), VFAT (Windows 95 and later), NTFS (Windows NT 4 and later)
- HPFS (IBM OS/2), HFS (Apple Macintosh), ADFS (Amiga), ADFS (Acorn)
- ISO9660 CD-ROM filesystem, UDF (Universal Disk Format) DVD filesystem

# Filesystems in Linux (2)



## ■ Network filesystems

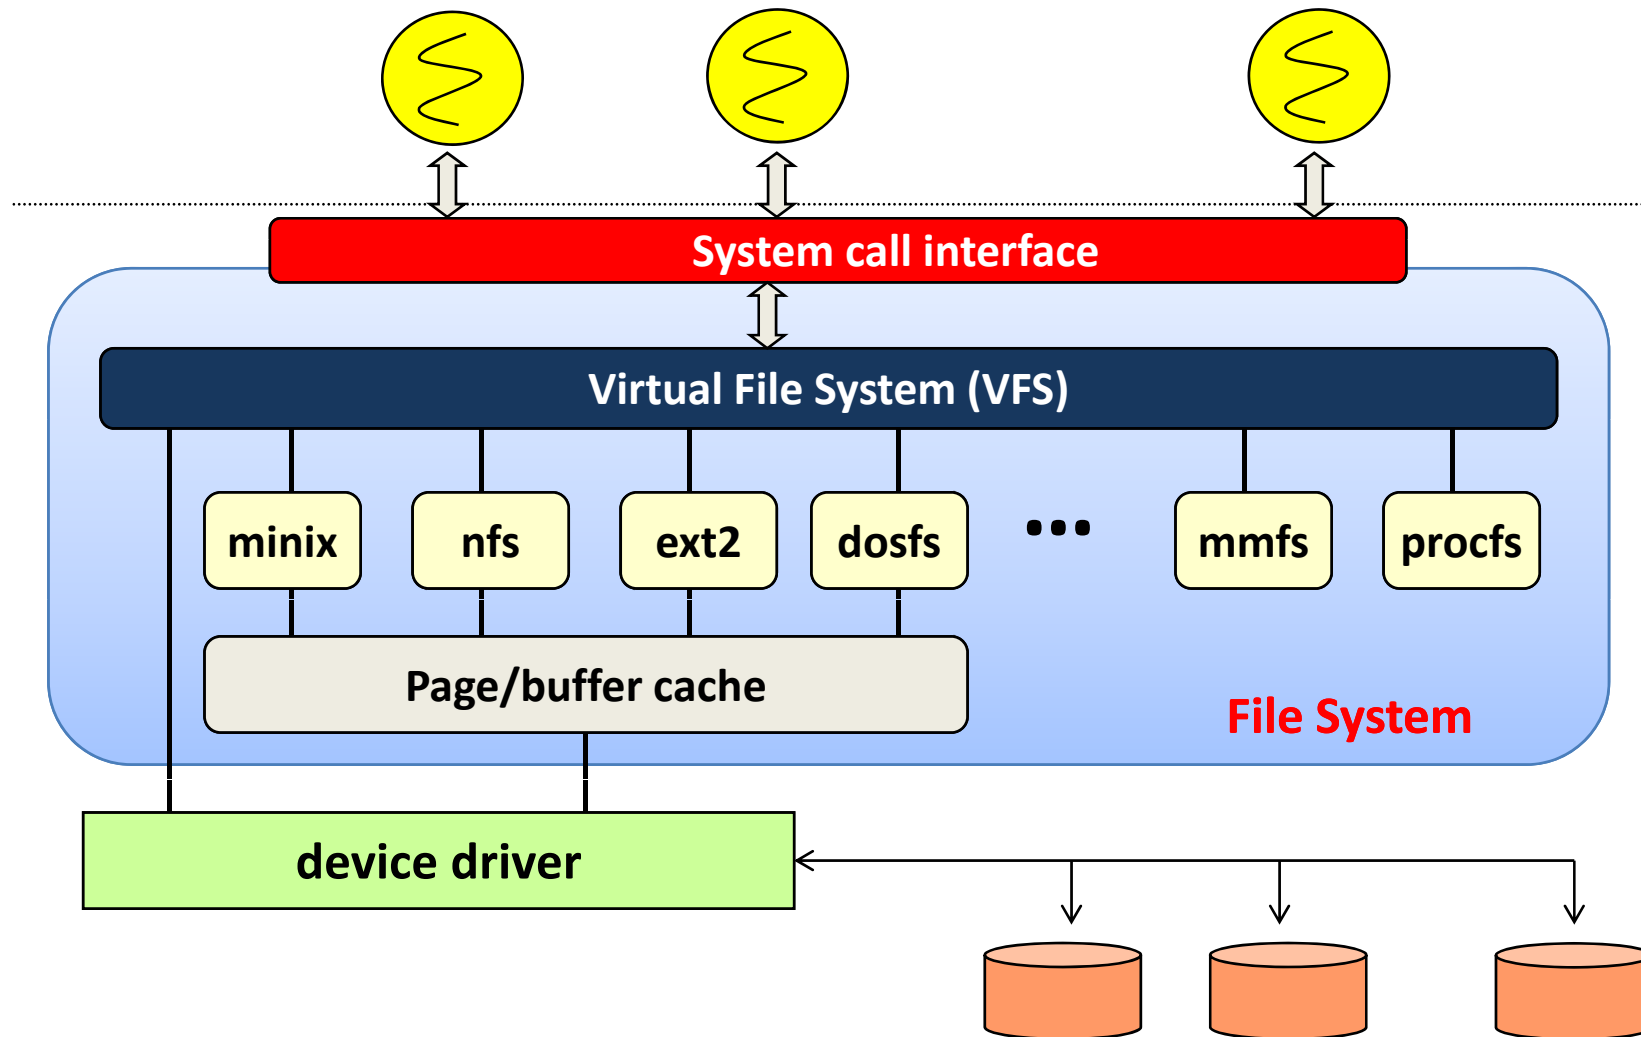
- NFS (Network File System)
- CIFS (Common Internet File System)
- AFS (Andrew File System)
- Coda
- NCP (Novell's NetWare Core Protocol)

## ■ Special filesystems

- /proc
- /dev
- ...

## ■ Flash filesystems

# File System Internals







*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Virtual File System (VFS)

**KAIST**

# VFS (1)



## ■ Virtual File System

- Manages kernel-level file abstractions in one format for all file systems.
- Receives system call requests from user-level (e.g., open, write, stat, etc.)
- Interacts with a specific file system based on mount point traversal.
- Receives requests from other parts of the kernel, mostly from memory management.
- Translates file descriptors to VFS data structures (such as vnode).



# VFS (2)

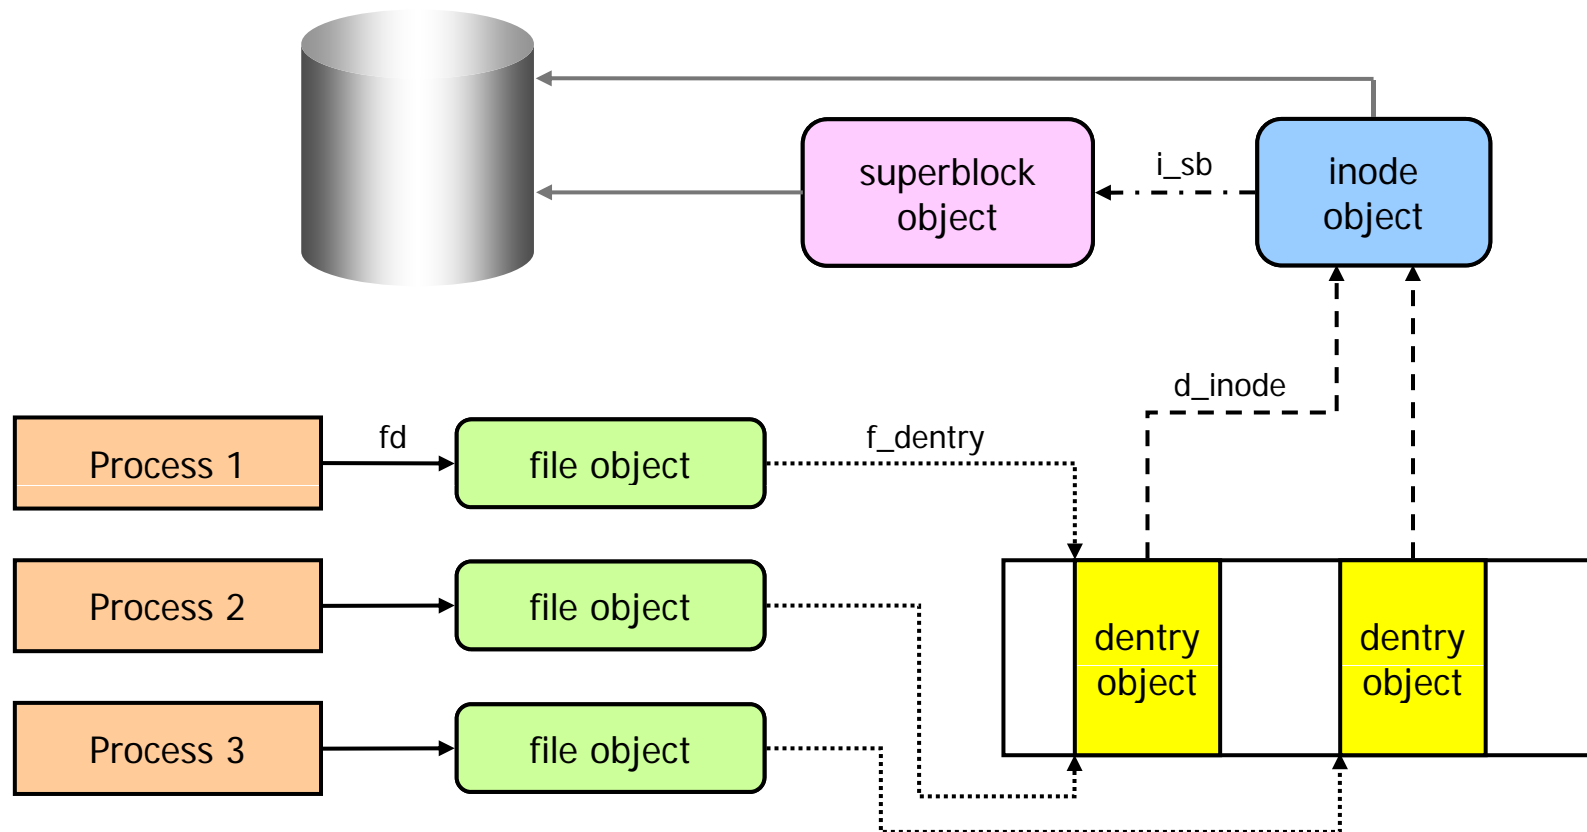


## ■ Linux: VFS common file model

- The superblock object
  - stores information concerning a mounted file system.
- The inode object
  - stores general information about a specific file.
- The file object
  - stores information about the interaction between an open file and a process.
- The dentry object
  - stores information about the linking of a directory entry with the corresponding file.
- In order to stick to the VFS common file model, in-kernel structures may be constructed on the fly.

# VFS (3)

- Interaction b/w processes and VFS objects



# Superblock Object (1)

## ■ The superblock object

- Store information describing the specific filesystem
- Usually corresponds to the filesystem superblock or filesystem control block, which is stored in a special sector on disk.
- Non disk-based filesystems generate the superblock on the fly and store it in memory
  - /proc

# Superblock Object (2)

- **struct super\_block** <linux/fs.h>
  - The global variable super\_blocks points to the list of superblock objects.

struct list_head	s_list;	Superblock objects are doubly linked
dev_t	s_dev;	Device identifier
unsigned long	s_blocksize;	Block size in bytes
struct super_operations *	s_op;	Superblock methods
struct dentry *	s_root;	Dentry object of the filesystem's root directory
void *	s_fs_info;	Filesystem-specific superblock information
unsigned char	s_dirt;	Flag indicating whether the superblock is updated
struct list_head	s_inodes;	List of all inodes
struct list_head	s_dirty;	List of modified inodes
struct list_head	s_files;	List of file objects
...		

# Superblock Object (3)

## ■ **struct super\_operations** <linux/fs.h>

- struct inode \***alloc\_inode**(struct super\_block \*sb);
  - Allocate space for an inode object including the space required for filesystem-specific data
- void **destroy\_inode**(struct inode \*inode);
- void **read\_inode**(struct inode \*inode);
  - Fill the fields of the inode object with the data on disk.
- void **write\_inode**(struct inode \*inode, int sync);
  - Update a filesystem inode synchronously or asynchronously
- void **put\_inode**(struct inode \*inode);
  - Invoked when the inode is released -- its reference counter is decreased.



# Superblock Object (4)

## ▪ struct super\_operations (cont'd)

- void **drop\_inode**(struct inode \*inode);
  - Invoked when the last user releases the inode.
- void **delete\_inode**(struct inode \*inode);
  - Invoked when the inode must be destroyed.
- void **put\_super**(struct super\_block \*sb);
  - Release the superblock object (on unmount)
- void **write\_super**(struct super\_block \*sb);
  - Update a filesystem superblock
- void **statfs**(struct super\_block \*sb, struct kstatfs \*buf);
  - Return statistics on a filesystem
- ...



# Inode Object (1)



## ■ The inode object

- Represent all the information needed by the kernel to manipulate a file or directory.
  - For UNIX filesystems, it is read from the on-disk inode.
  - Otherwise, the filesystem must fill the inode object from whatever stored on the disk. (e.g., FAT)
- The inode is unique to the file and remains the same as long as the file exists.

# Inode Object (2)



## ■ struct inode <linux/fs.h>

unsigned long	i_ino;	Inode number
umode_t	i_mode	File type and access rights
unsigned int	i_nlink;	The number of hard links
uid_t	i_uid;	Owner ID
gid_t	i_gid;	Group ID
loff_t	i_size;	File length in bytes
struct timespec	i_atime;	Time of last file access
struct timespec	i_mtime;	Time of last file write
struct timespec	i_ctime;	Time of last inode change
struct inode_operations *	i_op;	Inode operations
struct file_operations *	i_fops;	Default file operations
struct super_block *	i_sb;	Pointer to superblock object
struct list_head	i_sb_list;	Pointers for the list of inodes of the superblock
struct list_head	i_dentry;	The head of the list of dentry objects referencing this
...		

# Inode Object (3)



- **struct inode\_operations** <linux/fs.h>
  - int **create**(struct inode \*dir, struct dentry \*dentry, int mode, struct nameidata \*nd);
  - struct dentry \***lookup**(struct inode \*dir, struct dentry \*dentry, struct nameidata \*nd);
  - int **link**(struct dentry \*old, struct inode \*dir, struct dentry \*dentry);
  - int **unlink**(struct inode \*dir, struct dentry \*dentry);
  - int **symlink**(struct inode \*dir, struct dentry \*dentry, const char \*symname);
  - int **mkdir**(struct inode \*dir, struct dentry \*dentry, int mode);
  - int **rmdir**(struct inode \*dir, struct dentry \*dentry);

# Inode Object (4)



## ▪ struct inode\_operations (cont'd)

- int `mknod`(struct inode \*dir, struct dentry \*dentry,  
int mode, dev\_t rdev);
- int `rename`(struct inode \*old\_dir, struct dentry \*old\_dentry,  
struct inode \*new\_dir, struct dentry \*new\_dentry);
- ...

# File Object (1)

## ■ The file object

- In-memory representation of a file opened by a process
- File objects have no corresponding image on disk.
- Multiple file objects can exist for the same file.

## ■ **struct file** <linux/fs.h>

struct list_head	f_list;	Pointers for generic file object list
struct dentry *	f_dentry;	dentry object associated with this file
struct file_operations *	f_op;	Pointer to file operation table
atomic_t	f_count;	File object's reference counter
unsigned int	f_flags;	Flags specified when opening the file
mode_t	f_mode;	Process access mode
loff_t	f_pos;	Current file offset (file pointer)
struct address_space *	f_mapping;	Pointer to file's address space object
...		



# File Object (2)



## ▪ **struct file\_operations** <linux/fs.h>

- int **open**(struct inode \*inode, struct file \*file);
- int **llseek**(struct file \*file, loff\_t offset, int origin);
- ssize\_t **read**(struct file \*file, char \*buf, size\_t count, loff\_t \*pos);
- ssize\_t **write**(struct file \*file, const char \*buf, size\_t count, loff\_t \*pos);
- int **readdir**(struct file \*file, void \*dirent, filldir\_t filldir);
- int **mmap**(struct file \*file, struct vm\_area\_struct \*vma);
- int **flush**(struct file \*file);
- int **fsync**(struct file \*file, struct dentry \*dentry, int datasync);
- int **release**(struct inode \*inode, struct file \*file);
- ...



# Dentry Object (1)



## ■ The dentry object

- Directory implementations differ among filesystems.
- Once a directory entry is read into memory, it is transformed by VFS into a dentry object.
- The kernel creates a dentry object for every component to its corresponding inode.
  - /tmp/test: three dentry objects (/ , /tmp , /tmp/test)
- Dentry objects have no corresponding image on disk.
- Dentry objects are stored in a slab allocator cache.
  - `kmem_cache_alloc()`, `kmem_cache_free()`

# Dentry Object (2)



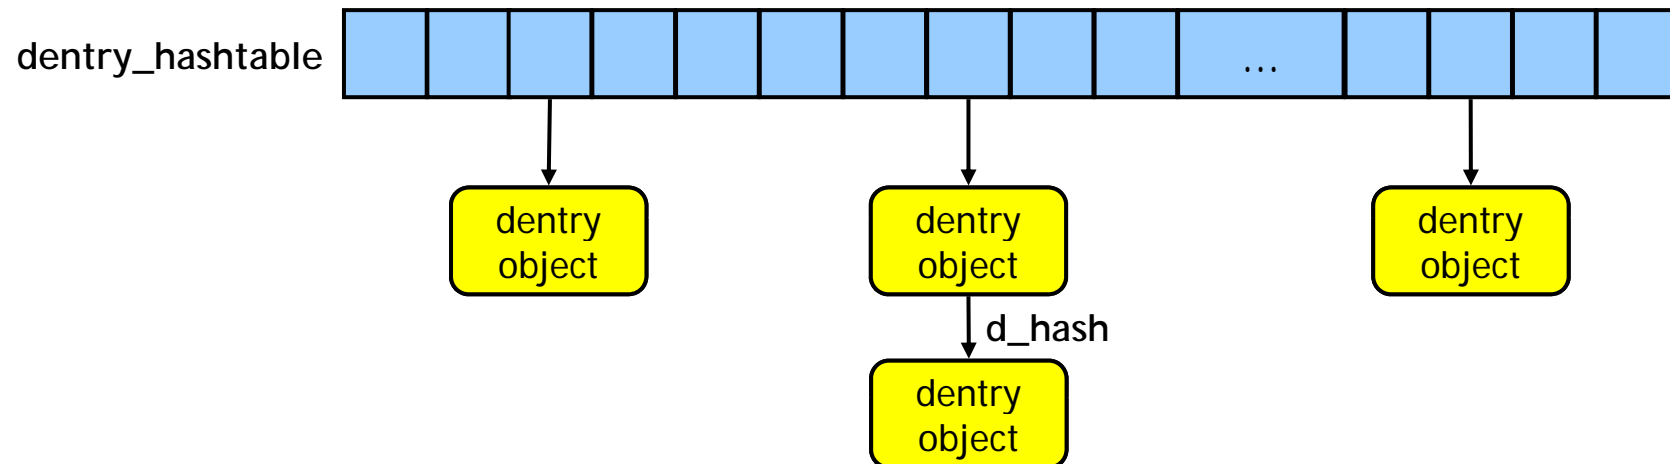
## ■ **struct dentry** <linux/dcache.h>

atomic_t	d_count;	Dentry object usage counter
unsigned int	d_flag;	Dentry cache flags
struct qstr	d_name;	File name
struct inode *	d_inode;	Inode associated with the dentry object
struct dentry *	d_parent;	Dentry object of parent directory
struct list_head	d_lru;	Pointers for the list of unused dentries
struct list_head	d_child;	For directories, pointers for the list of dentries in the same parent directory
struct list_head	d_subdirs;	For directories, head of the list of subdirectory dentries
struct list_head	d_alias;	Pointers for the list of dentries associated with the same inode (alias)
struct dentry_operations *	d_op;	Dentry methods
struct super_block *	d_sb;	Superblock object of the file
struct hlist_node	d_hash;	Pointer for list in hash table entry
int	d_mounted;	The number of file systems mounted on this dentry
...		

# Dentry Object (3)

## ■ The dentry cache

- Hash table: `dentry_hashtable` array
  - Derive the dentry object associated with a given filename and a given directory quickly.
  - Hash table size: 256 entries per megabyte of RAM
  - The adjacent elements associated with a single hash value are linked via `d_hash` field of the dentry object.

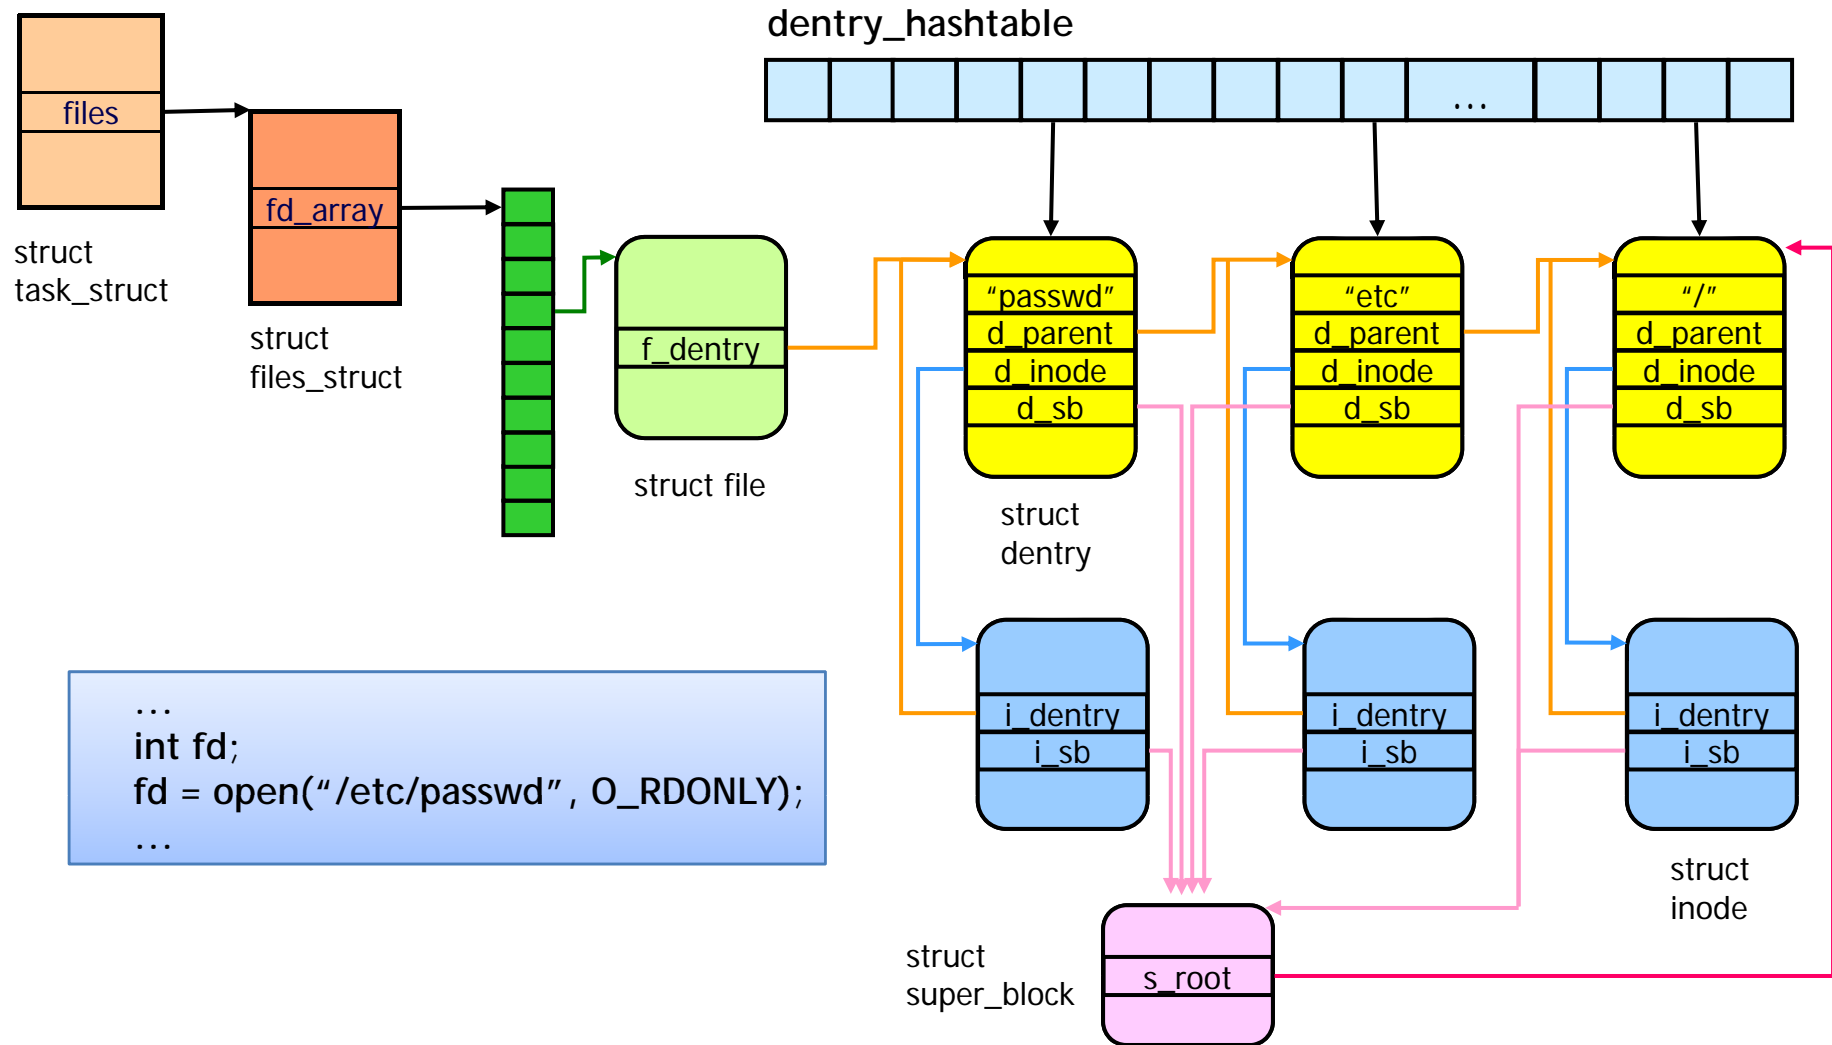


# Dentry Object (4)



- **struct dentry\_operations** <linux/dcache.h>
  - int **d\_revalidate**(struct dentry \*dentry, struct nameidata \*nd);
  - int **d\_hash**(struct dentry \*dentry, struct qstr \*name);
    - Create a filesystem-specific hash value for the dentry hash table
  - int **d\_compare**(struct dentry \*dentry, struct qstr \*name1, struct qstr \*name2);
  - int **d\_delete**(struct dentry \*dentry);
  - void **d\_release**(struct dentry \*dentry);
  - void **d\_iput**(struct dentry \*dentry, struct inode \*inode);
    - Called when a dentry object becomes negative (losing inode).
    - The default VFS function invokes iput() to release the inode object.

# The Big Picture





*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# FAT FS

**KAIST**



# FAT FS (1)



## ■ FAT filesystem

- Used in MS-DOS based OSes
  - MS-DOS, Windows 3.1, 95, 98, ...
- Originally developed as a simple file system suitable for floppy disk drives less than 500KB in size.
- FAT stands for File Allocation Table.
  - Each FAT entry contains a pointer to a region on the disk
- Currently there are three FAT file system types:  
FAT12, FAT16, FAT32

# FAT FS (2)

## ■ FAT filesystem organization



- FAT file system on disk data structure is all “little endian.”
- The data area is divided into clusters.
  - Used for subdirectories and files.
- Root directory region doesn't exist on FAT32.

# FAT FS (3)



## ■ Boot sector

- The first sector on the disk.
- Contains BPB (BIOS Parameter Block).
  - Sectors per cluster
  - The number of sectors on the volume.
  - Volume label.
  - The number of root directory entries.
  - File system type (FAT12, FAT16, FAT32)
  - and many more.
- If the volume is bootable, the first sector also contains the code required to boot the OS.

# FAT FS (4)



## ■ FAT (File Allocation Table)

- Starts at sector 1 (after the boot sector)
- The FAT defines a singly linked list of the clusters of a file.
- The first two entries in the FAT can be ignored.
  - The first entry available is entry 2.
- The individual entries in the FAT table define the “chains” of clusters that make up a file.
- There are two copies so that corruption of the FAT can be detected and repaired.

# FAT FS (5)



## ■ FAT12 example

- Each FAT12 entry is 12bits.
  - When designed, space was tight.
  - Pack 2 entries into 3 bytes.
  - 4096 possible clusters.
  - If a sector is 512bytes and cluster = 1 sector, can represent 2MB of data.
- FAT12 entry values:
  - 0 Unused cluster
  - 0xFF0-0xFF6 Reserved cluster
  - 0xFF7 Bad cluster
  - 0xFF8-0xFFFF End of Clusterchain mark
  - Other Next cluster in file

# FAT FS (6)

- Maximum partition size allowed

Block size	FAT-12	FAT-16	FAT-32
512 B	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1 GB	2 TB
32 KB		2 GB	2 TB



# FAT FS (7)




## ■ Directories

- The root directory is fixed in length and always located at the start of the volume (after the FAT).
  - FAT32 treats the root directory as just another cluster chain in the data area.
- A subdirectory is nothing but a regular file that has a special attribute indicating it is a directory.
  - No size restriction
- The data or contents of the “file” is a series of 32byte FAT directory entries.
  - Filename’s first character is usage indicator:
    - » 0x00      Never been used.
    - » 0xe5      Used before but entry has been released.

# FAT FS (8)

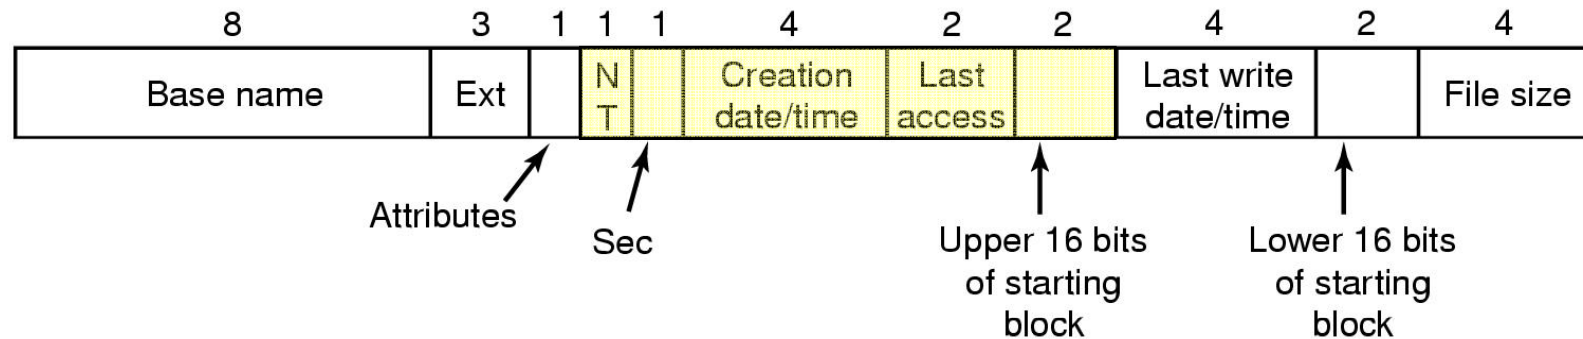
## ■ FAT directory entry




- Attributes:
  - Read Only, Hidden, System, Volume Label, Subdirectory, Archive

# FAT FS (9)

## ■ FAT32 directory entry



- An entry for a long file name



# FAT FS (10)

- Representing long file name in FAT32
  - The quick brown fox jumps over the lazy dog

68	d	o	g	A	0	C	K					0							
	3	o	v	e	r	A	0	C	K	t	h	e		l	a	0	z	y	
	2	w	n		f	o	A	0	C	K	x		j	u	m	p	0	s	
	1	T	h	e		q	A	0	C	K	u	i	c	k		b	0	r	o
	T	H E Q U I ~ 1					A	N	S	Creation	Last			Last			Low	Size	
Bytes										time	acc		Upp						



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# CRAMFS

**KAIST**



# CRAMFS (1)



## ■ Compressed ROM Filesystem

- Designed to be simple and small
- Uses the zlib routines to compress a file one page at a time
- Allows random page access
- The metadata is not compressed, but is expressed in a very terse representation to make it use much less disk space
- File size is limited to less than 16MB
- Maximum filesystem size is a little over 256MB
- Limited metadata: no timestamps, 8-bit gid, ...



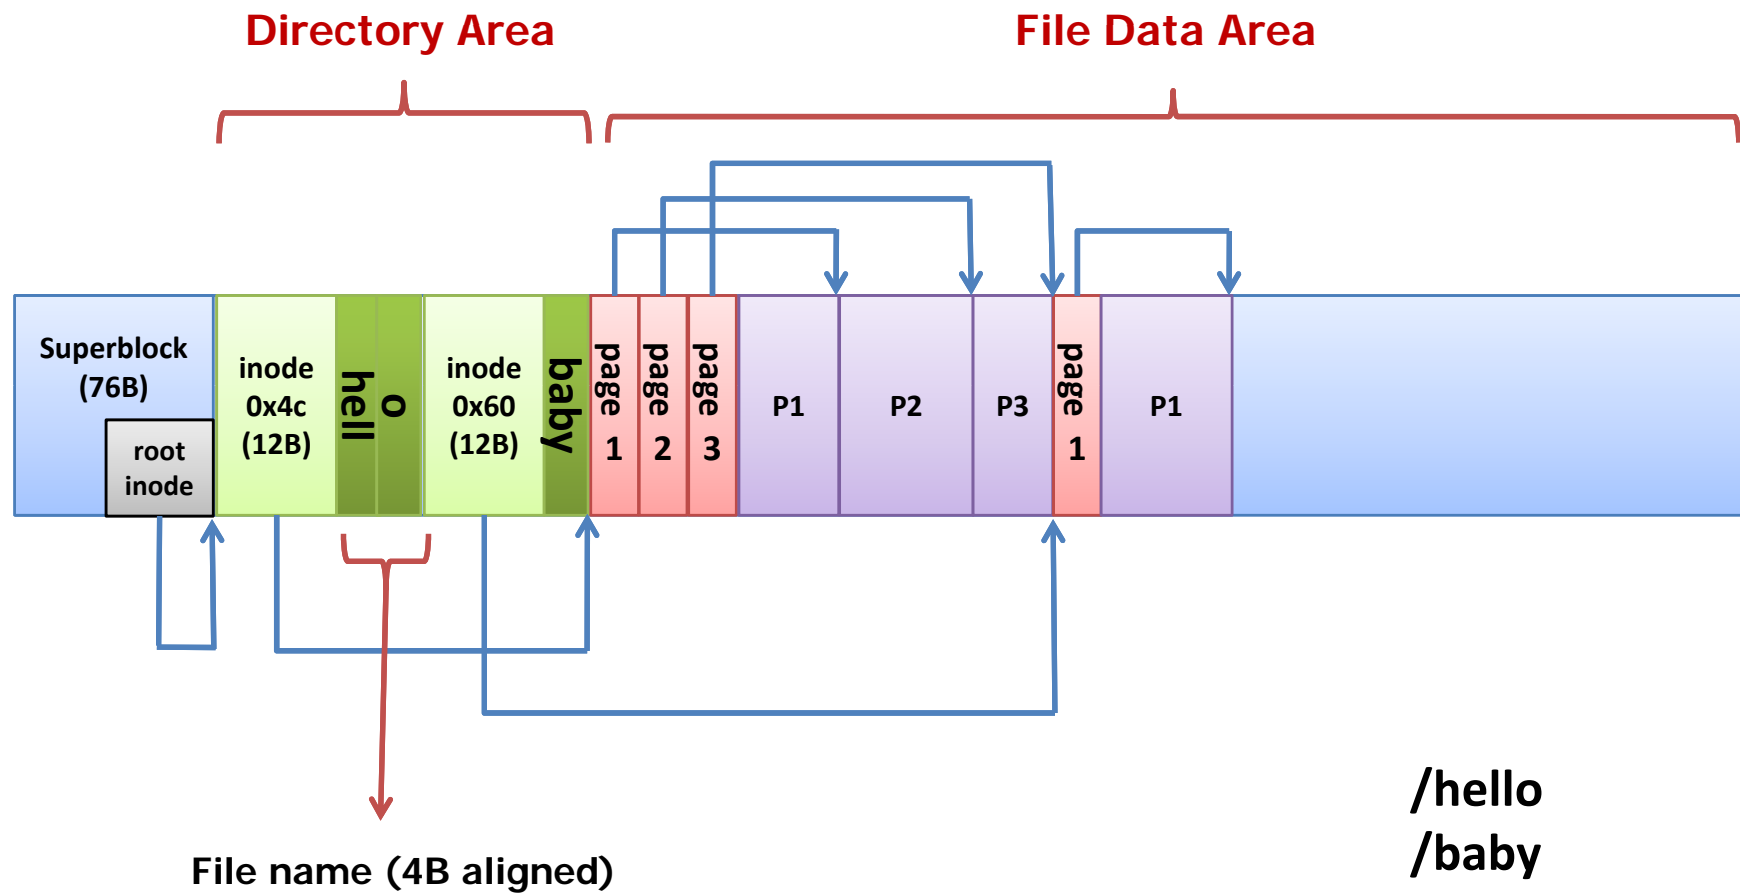
# CRAMFS (2)



## ■ File system layout

- Superblock
  - 76 bytes, fixed
- Directory area
  - One entry for each file or directory
  - Consists of inode (fixed size) + file name (variable size)
  - Offset of the directory entry is used as the inode number
- File data area
  - Offsets + compressed data chunks
  - All data blocks of a file are stored sequentially

# CRAMFS (3)





*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# JFFS/JFFS2

**KAIST**

# JFFS (1)

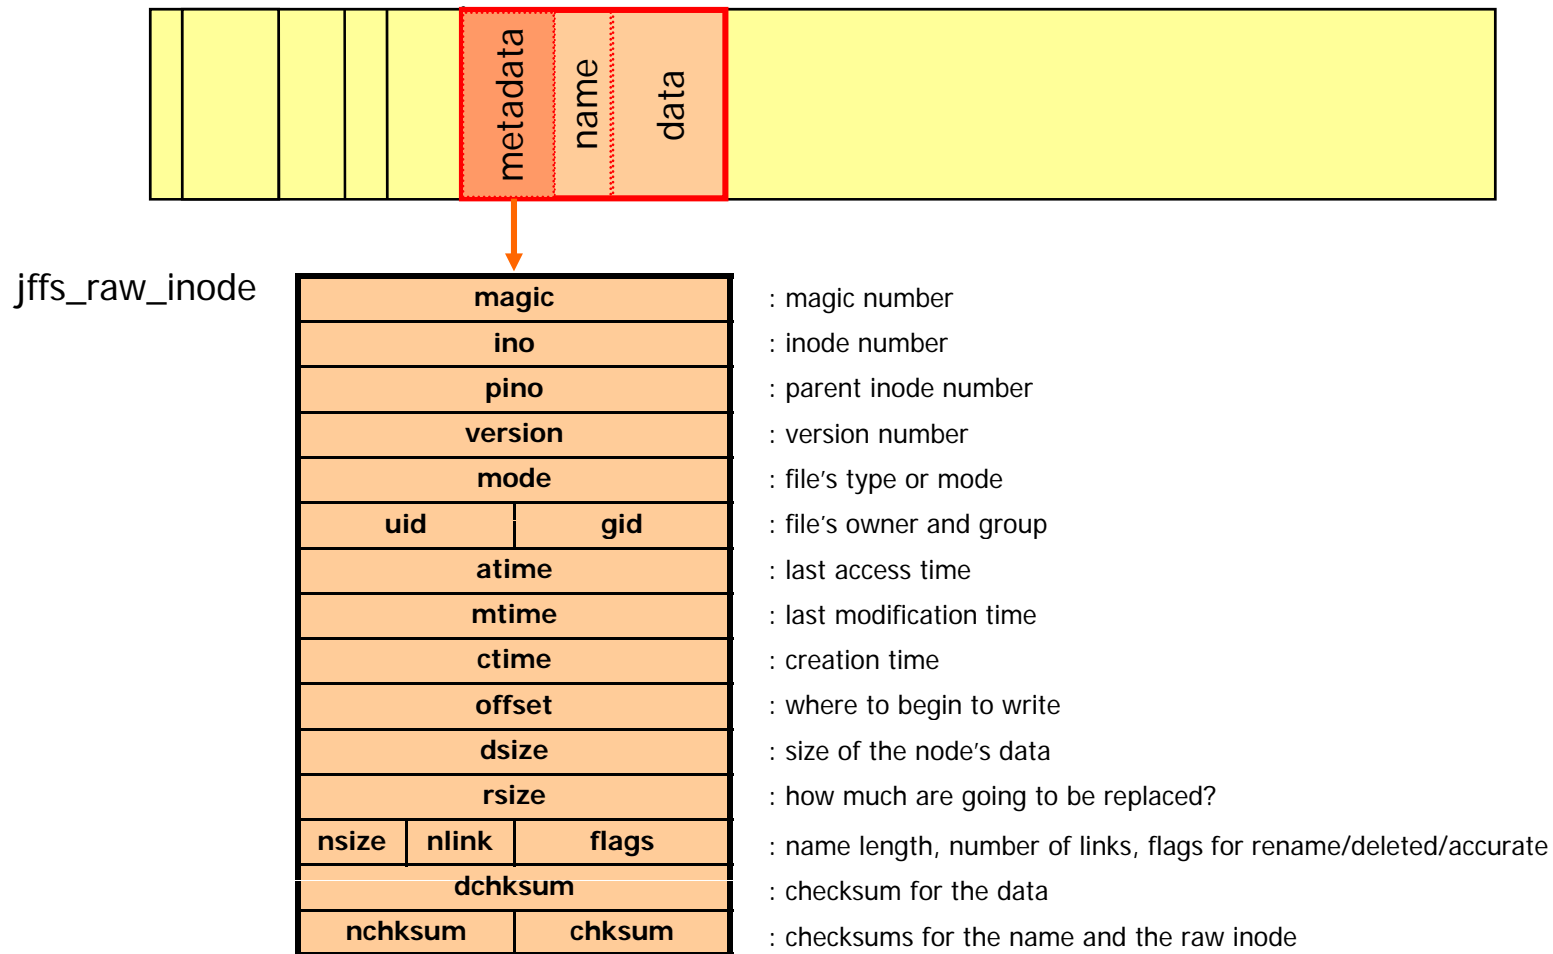


## ■ JFFS (Journaling Flash File Systems)

- Developed by Axis Communications, Sweden in 1999.
- Released under GNU GPL
- Designed for small NOR flashes
- A log-structured file system
  - Any file system modification is appended to the log.
  - The log is the only data structure on the flash media.  
Log = <metadata, (name), (data)>
  - A file is obsoleted by a later log in whole or in part.
  - Obsoleted logs are reclaimed via garbage collection.
- Rely on special in-core data structures for  
filename → metadata, metadata → data mappings.

# JFFS (2)


## ■ JFFS architecture



# JFFS (3)

## ■ Garbage collection

- The free space is eventually exhausted. Now what?
- Erase the oldest block in the log.



- Live nodes should be moved.
- Perfectly wear-leveled.



# JFFS2 (1)



## ■ JFFS2 limitations

- Poor garbage collection performance
  - A block is garbage collected even if it contains only clean nodes.
  - In many cases, there are static data. (libraries, program executables, etc.)
- No compression support
  - Flash memories are expensive.
- No support for hard links
  - File name and parent i-node are stored in each i-node.
- No support for NAND flashes

# JFFS2 (2)

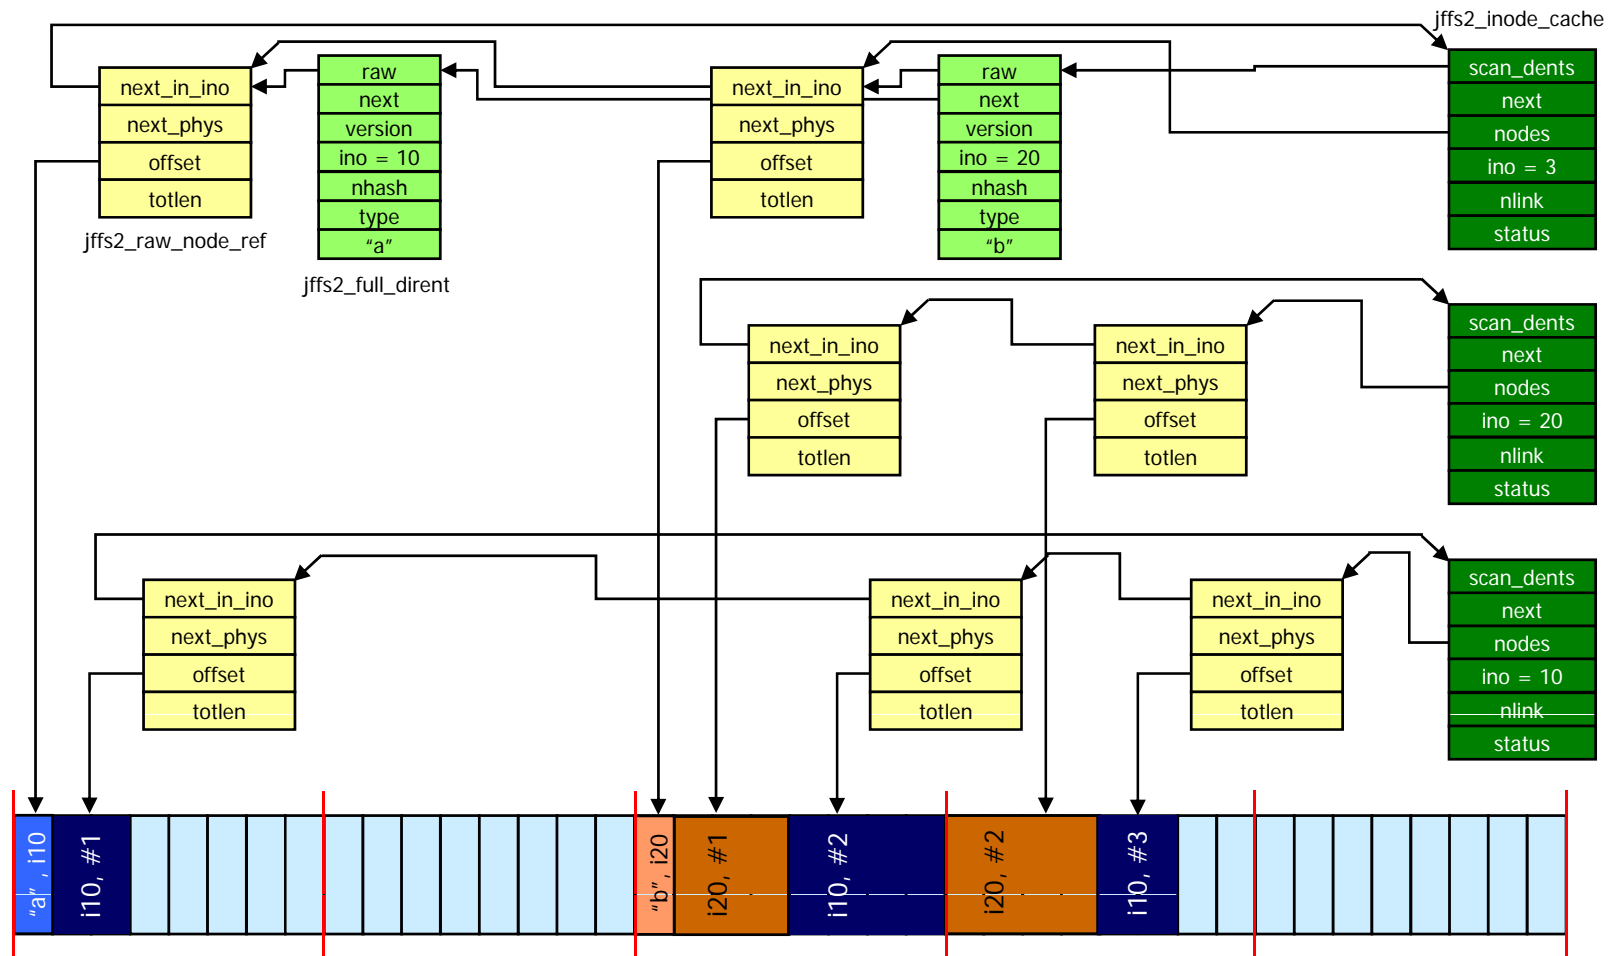


## ■ Node types

- JFFS2\_NODETYPE\_INODE
  - Similar to `jffs_raw_inode`
  - No filename, no parent i-node number
  - Compression support
- JFFS2\_NODETYPE\_DIRENT
  - Represent a directory entry, or a link
  - File name, i-node, parent i-node (directory's i-node), etc.
  - File name with i-node = 0: deleted file
- JFFS2\_NODETYPE\_CLEANMARKER
  - To deal with the problem of partially-erased blocks due to the power failure during erase operation

# JFFS2 (3)

## ■ JFFS2 architecture



# JFFS2 (4)



## ■ Block lists (old)

- free\_list: empty blocks
- clean\_list: blocks full of valid nodes
- dirty\_list: blocks containing at least one obsoleted node

## ■ Block lists (now)

- free\_list: empty blocks
- clean\_list: blocks full of valid nodes
- very\_dirty\_list: blocks with lots ( $>50\%$ ) of dirty nodes
- dirty\_list: blocks with some ( $<50\%$ ) dirty nodes
- erasable\_list: blocks which are completely dirty
- ...

# JFFS2 (5)



## ■ Garbage collection

- Invoked if the size of free\_list is less than the threshold.
- Small nodes can be merged by GC.
- Which blocks? (old)
  - 99% from dirty\_list ( $\text{jiffies} \% 100 \neq 0$ )
  - 1% from clean\_list (for wear-leveling)
- Which blocks? (now)
  - $n = \text{jiffies} \% 128$
  - If ( $n < 50$ ) GC from erasable\_list
  - else if ( $n < 110$ ) GC from very\_dirty\_list
  - else if ( $n < 126$ ) GC from dirty\_list
  - else if GC from clean\_list



# JFFS2 (6)



## ■ JFFS2 limitations

- Large memory consumption
  - In-core data structures
    - » `jffs2_raw_node_ref` (16bytes/node), `jffs2_inode_cache`
- Slow mount time
  - 4 sec for 4MB!
- Runtime overheads (space & time)
  - Build child directory entries from flash on directory access
  - Build node fragments on file access
  - All the inode's nodes should be examined (with CRC checked)
- Do not utilize NAND OOB area



# JFFS2 (7)



## ■ JFFS2 memory consumption example

- JFFS2 with 64MB NAND flash
  - Typical Linux root FS: 2.2MB  
(719 directories, 2995 regular files)
  - 64MB file with 512bytes/node: 6.7MB
  - 64MB file with 10bytes/node: 47.6MB
- JFFS2 with 1GB NAND flash (estimated)
  - Typical Linux root FS: 34.7MB
  - 64MB file with 512bytes/node: 104.2MB
  - 64MB file with 10bytes/node: 743.6MB

*(Source: JFFS3 Design Issues, June 4, 2005)*



*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

# Linux MTD Layer

**KAIST**

# MTD (1)



- **What is MTD (Memory Technology Device)?**
  - in <linux/drivers/mtd/Kconfig>

config MTD

tristate "Memory Technology Device (MTD) support"

help

Memory Technology Devices are **flash, RAM and similar chips**, often used for solid state file systems on embedded devices. This option will provide the **generic support** for **MTD drivers** to register themselves with the kernel and for potential **users** of MTD devices to enumerate the devices which are present and obtain a handle on them. It will also allow you to select individual drivers for particular hardware and users of MTD devices.

# MTD (2)



## ■ What is MTD? (cont'd)

- Generic device driver of memory mapped device
  - Flash memory, ROM, etc.
- Goals
  - To make it simple to provide a driver for new hardware by providing a generic interface between the hardware drivers and the upper layers of the system.
- MTD provides an “MTD device” abstraction.
- Added in Linux kernel 2.4.0
  - /usr/src/linux/drivers/mtd
- <http://www.linux-mtd.infradead.org>

# MTD (3)



## ■ MTD user module

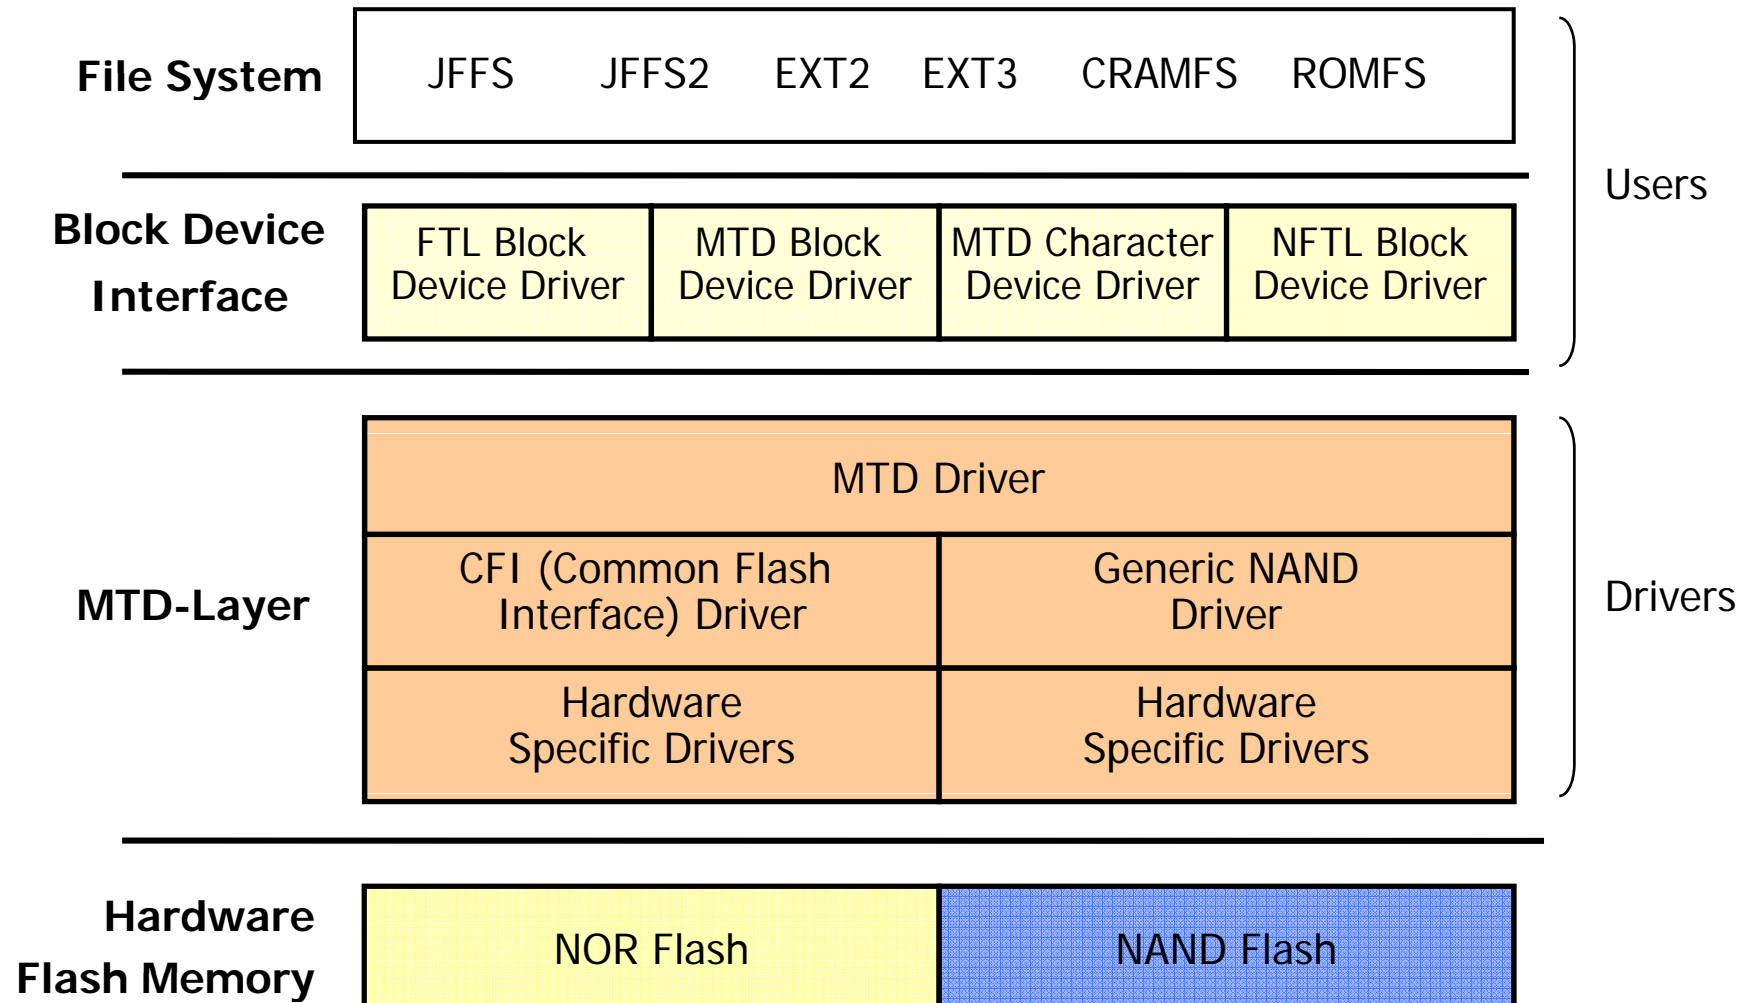
- Interfaces that can be used directly from userspace
  - Raw character access
  - Raw block access
  - Flash translation layer (FTL, NFTL, INFTL)
  - Flash aware file systems (JFFS2, YAFFS, ...)

## ■ MTD driver module

- Provide physical access to memory device and accessed through the user modules
  - On-board memory, On-board NAND flash
  - Common Flash Interface (CFI) on-board NOR flash
  - M-Systems' DiskOnChip 2000 and Millennium



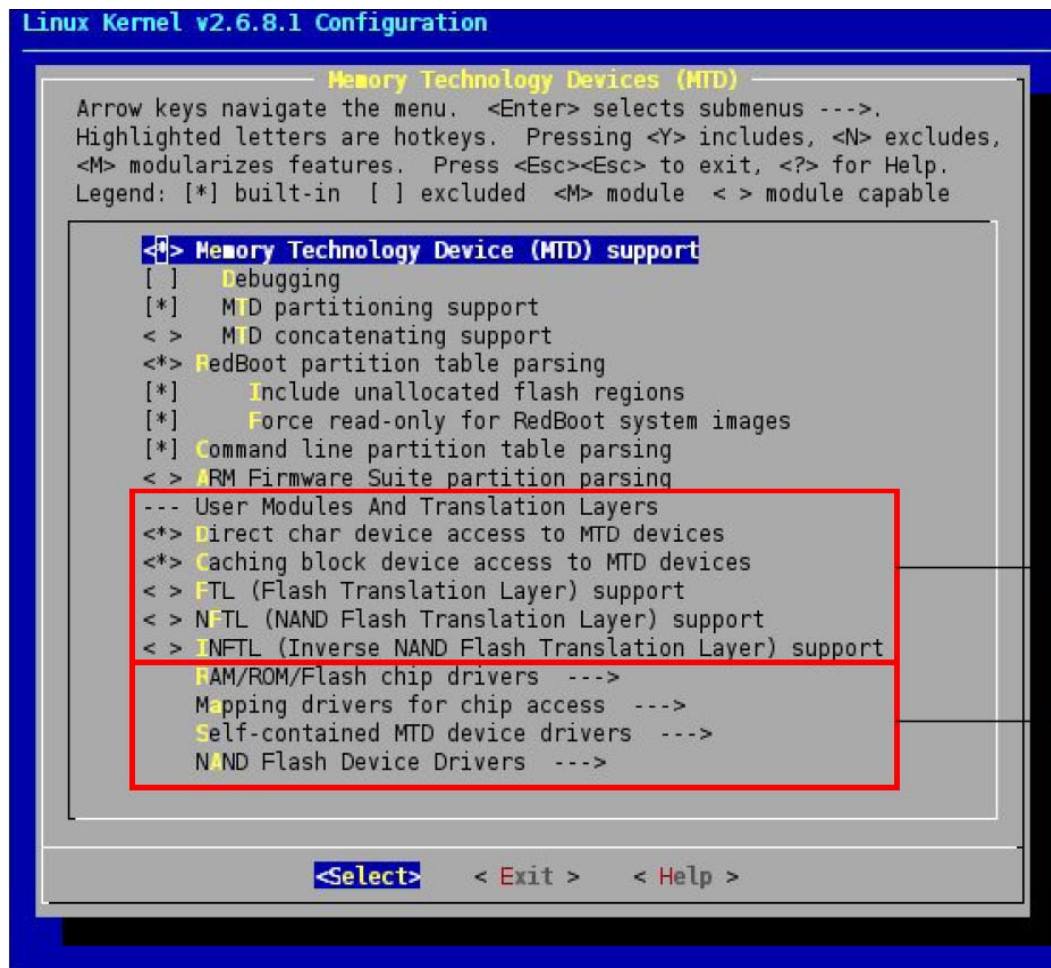
# MTD (4)





# Using MTD (1)

## ■ Kernel configuration



MTD user module

MTD driver module

# Using MTD (2)

## ■ Module loading

```
$ modprobe nandsim
$ cat /proc/mtd
dev:   size  erasesize  name
mtd0: 00800000 00002000 "NAND simulator partition"
$ modprobe jffs2
$ lsmod
Module                Size  Used by
nandsim                22948  0
nand                   32132  2 nandsim
nand_ids                4224  2 nandsim,nand
nand_ecc                2688  1 nand
mtdpart               10496  2 nandsim,nand
jffs2                 100528  1
zlib_deflate           21024  1 jffs2
mtdcore                 7108  4 nand,mtdpart,jffs2
$
$
```

# Using MTD (3)



## ■ Mounting

```
$ mknod /dev/mtdblock0 b 31 0
$ mount -t jffs2 /dev/mtdblock0 /mnt/
$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/hda2	19G	4.2G	14G	24%	/
tmpfs	506M	0	506M	0%	/dev/shm
/dev/hda5	46G	27G	18G	61%	/home
/dev/hdc1	22G	21G	1.8G	93%	/d
tmpfs	10M	136K	9.9M	2%	/dev
/dev/mtdblock0	8.0M	280K	7.8M	4%	/mnt

```
$
```

# Using MTD (4)



## ■ Main interface

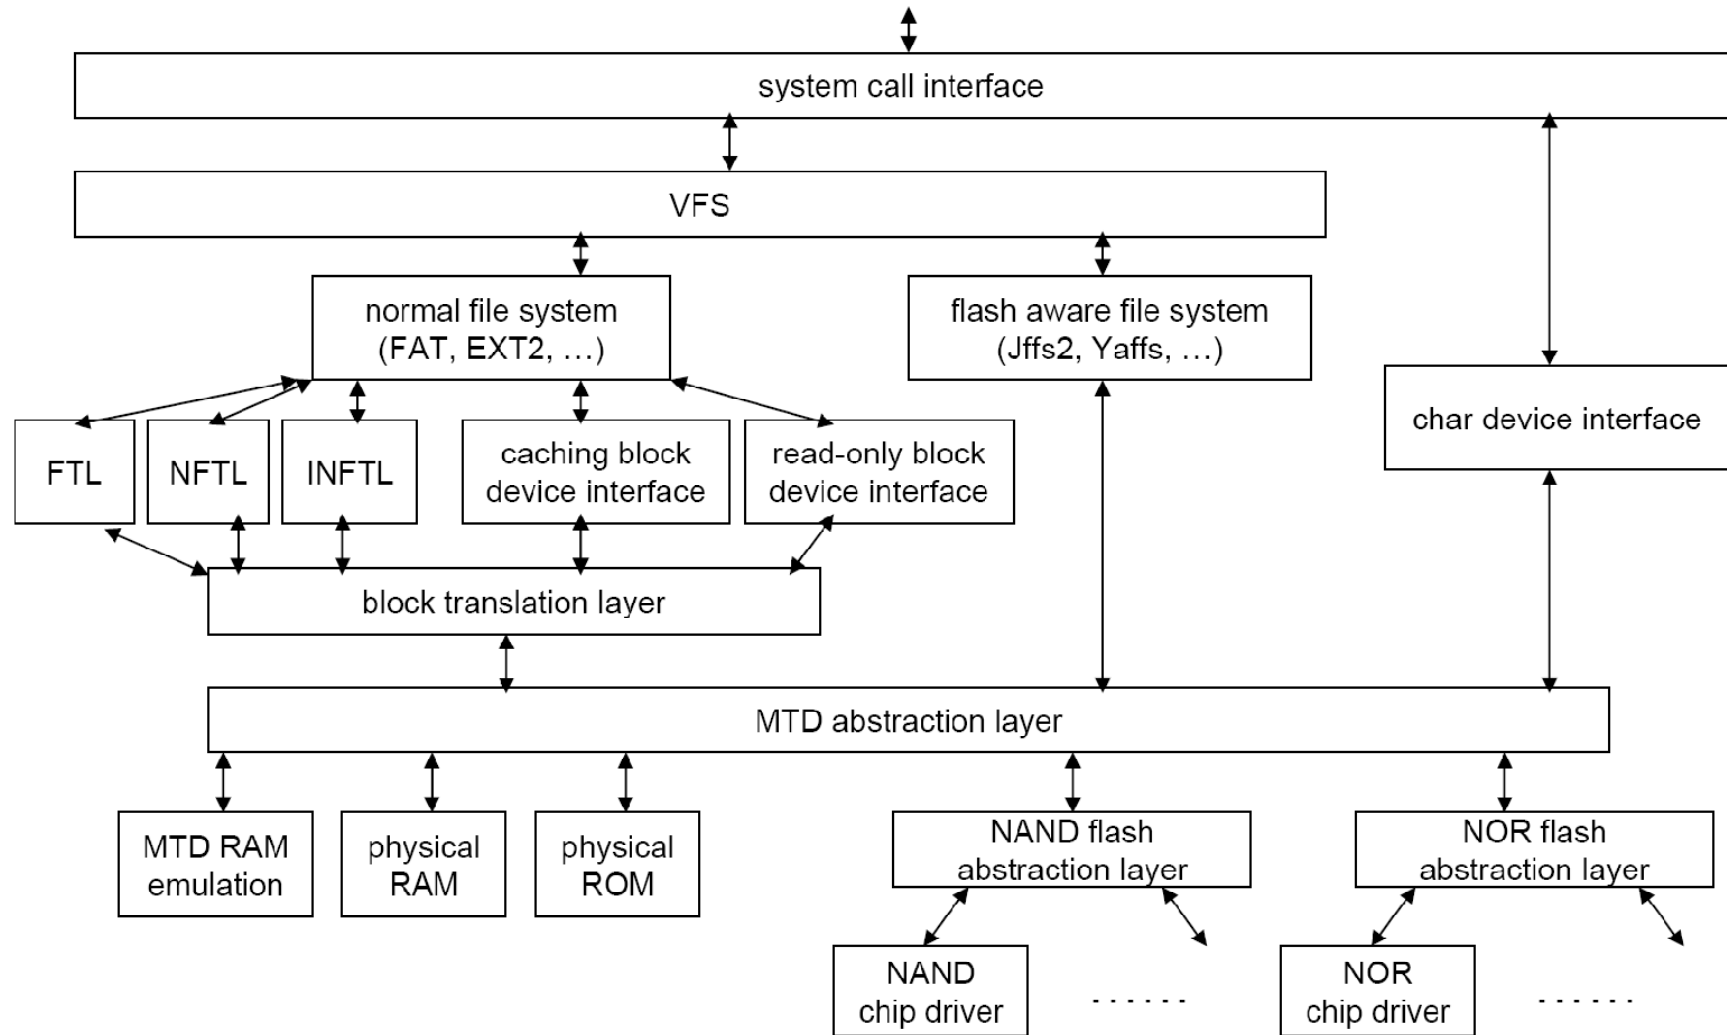
- int **read** (struct mtd\_info \*mtd, loff\_t from, size\_t len, size\_t \*retlen, u\_char \*buf);
- int **write** (struct mtd\_info \*mtd, loff\_t to, size\_t len, size\_t \*retlen, u\_char \*buf);
- int **erase** (struct mtd\_info \*mtd, struct erase\_info \*instr);

## ■ Misc. interface

- read\_ecc(), write\_ecc(), read\_oob(), write\_oob(), ...




# MTD Big Picture



# MTD Device (1)

## ■ MTD device

- MTD device consists of eraseblocks
- Eraseblock size varies, typically 32—128 KB
- Erasesblocks may be written to, but not re-written
- Whole erase block has to be erased first
- Then, it is possible to write there





# MTD Device (2)



## ■ Block device vs. MTD device

Block device	MTD device
<ul style="list-style-type: none"><li>• Consists of sectors</li><li>• Sectors are small (512, 1024 bytes)</li><li>• 2 operations: <b>read</b> and <b>write</b></li><li>• Bad sectors are hidden by hardware</li><li>• Sectors do not get worn out</li></ul>	<ul style="list-style-type: none"><li>• Consists of eraseblocks</li><li>• Eraseblocks are larger (32-128Kbytes)</li><li>• 3 operations: <b>read</b>, <b>write</b>, and <b>erase</b></li><li>• Bad eraseblocks are not hidden</li><li>• Eraseblocks get worn-out after <math>10^4</math>-<math>10^5</math> erasures</li></ul>


MTD device is more difficult to handle!

# MTD Device (3)



## ■ MTD partitions


- Flash chip may be split on several MTD partitions
- MTD partition is a set of consecutive eraseblocks
- MTD partition is a physical flash area



# MTD Device (4)

## ■ Drawbacks of MTD partitions

- MTD partitions are static
- Do not provide wear-leveling for the whole chip





*CS632/SEP564: Embedded Operating Systems (Fall 2008)*

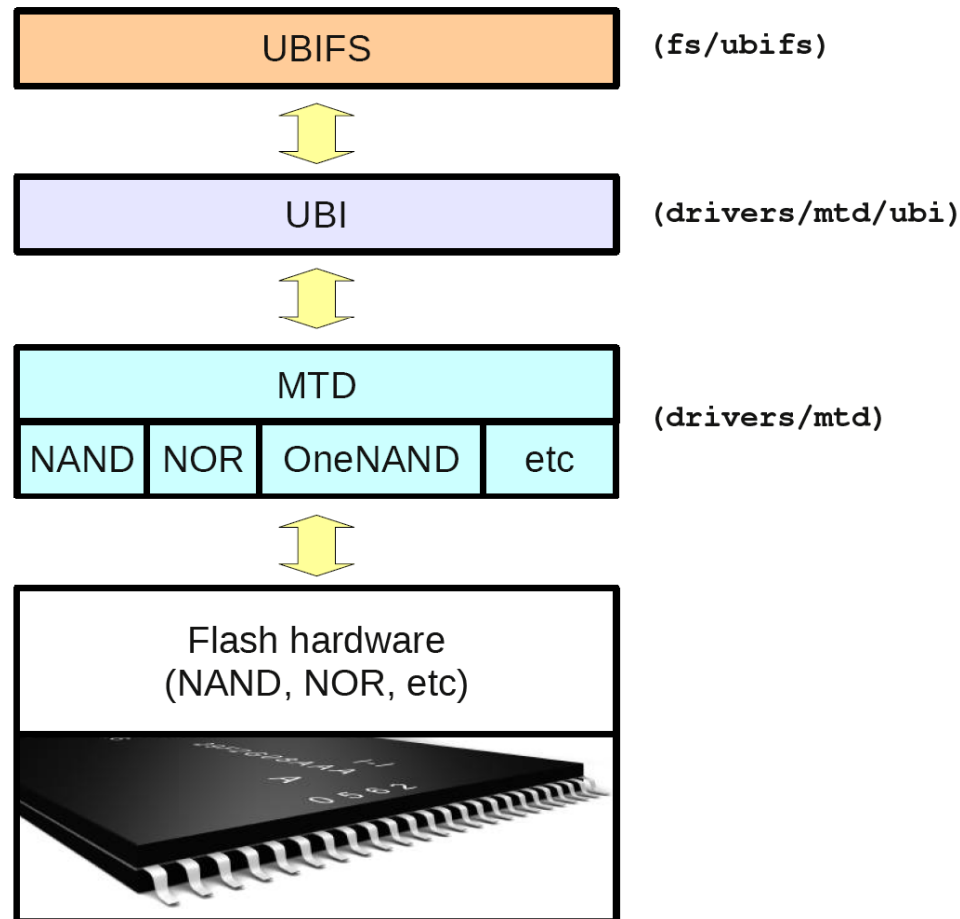
# UBI Layer

*Courtesy: Slides borrowed from “UBI - Unsorted Block Images ” by Artem Bityutskiy*



# UBI (1)

## ■ Unsorted Block Images




# UBI (2)

## ■ Logical volumes

- UBI provides logical volumes instead of MTD partitions
- UBI volumes are in a way similar to LVM volumes.
- UBI volumes may be dynamically created, deleted and re-sized.



© 2008 by UBI authors. All rights reserved. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike license.

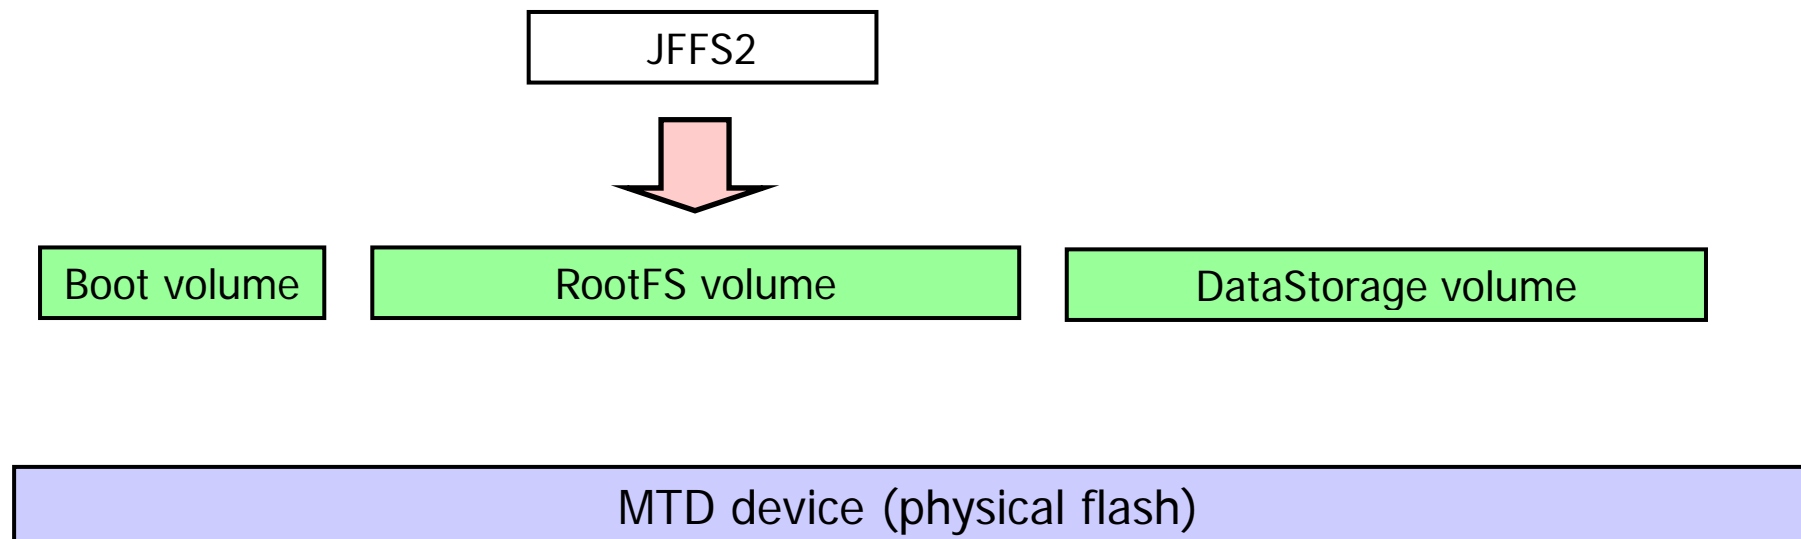




# UBI (3)

## ■ Wear-leveling

- UBI does wear-leveling across **whole** MTD device.
- Wear-leveling is done by UBI, not by the UBI user.



# UBI (4)



## ■ UBI volume vs. MTD partition

MTD partition	UBI volume
<ul style="list-style-type: none"><li>• Consists of physical eraseblocks (PEB)</li><li>• Does not implement wear-leveling</li><li>• Admits of bad PEBs</li></ul>	<ul style="list-style-type: none"><li>• Consists of logical eraseblocks (LEB)</li><li>• Implements wear-leveling</li><li>• Devoid of bad LEBs</li></ul>

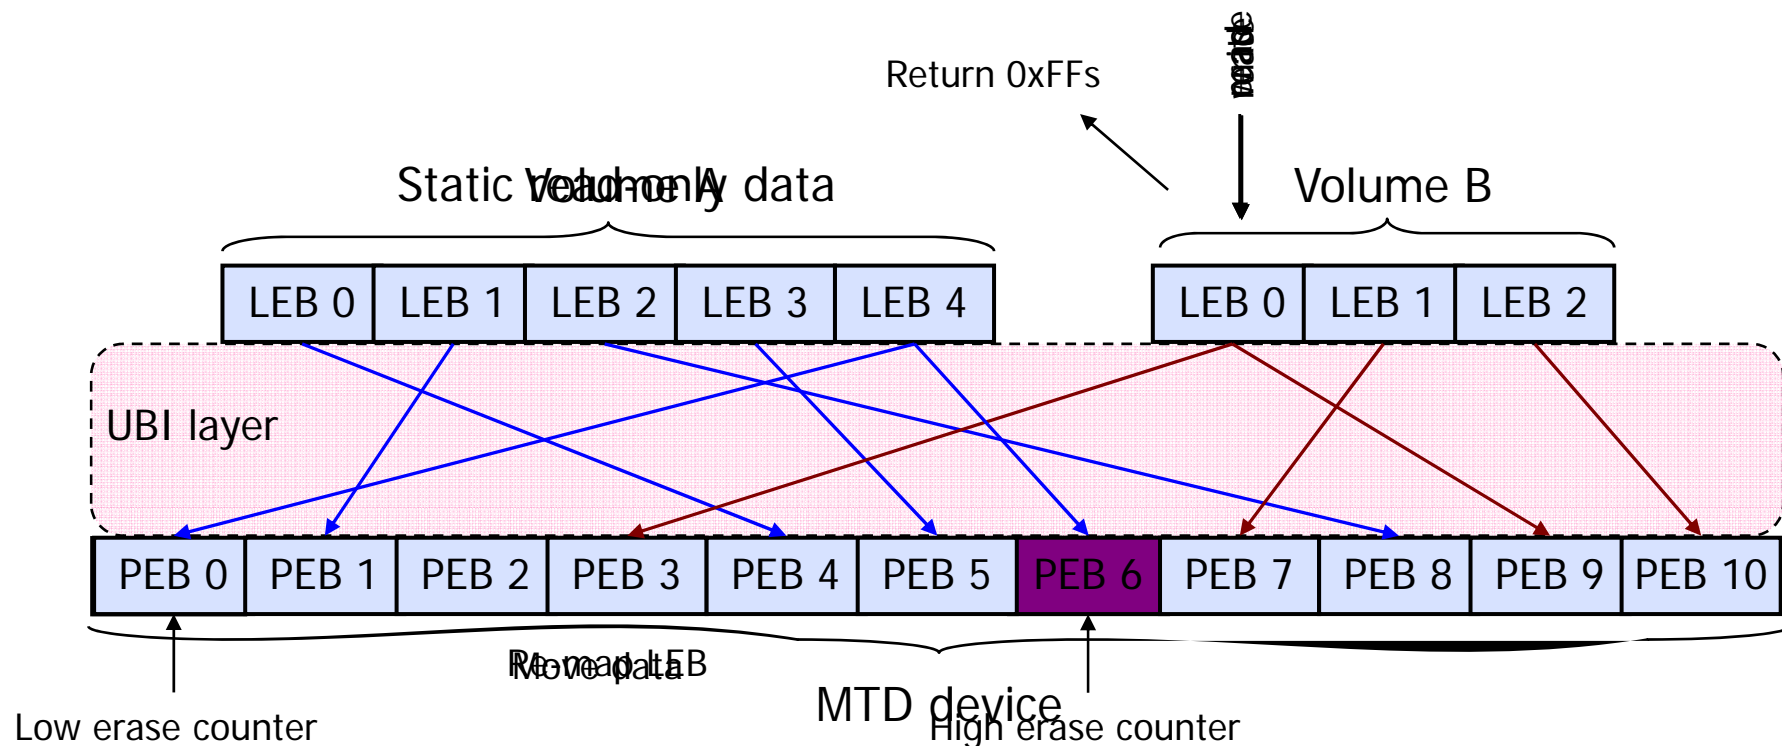
### Advantages of UBI

- Allows dynamic volume creation, deletion and re-sizing → more flexibility
- Eliminates the “wear” problem → simpler software
- Eliminates bad eraseblocks problem → simpler software

# UBI Internals (1)

## ■ How it works

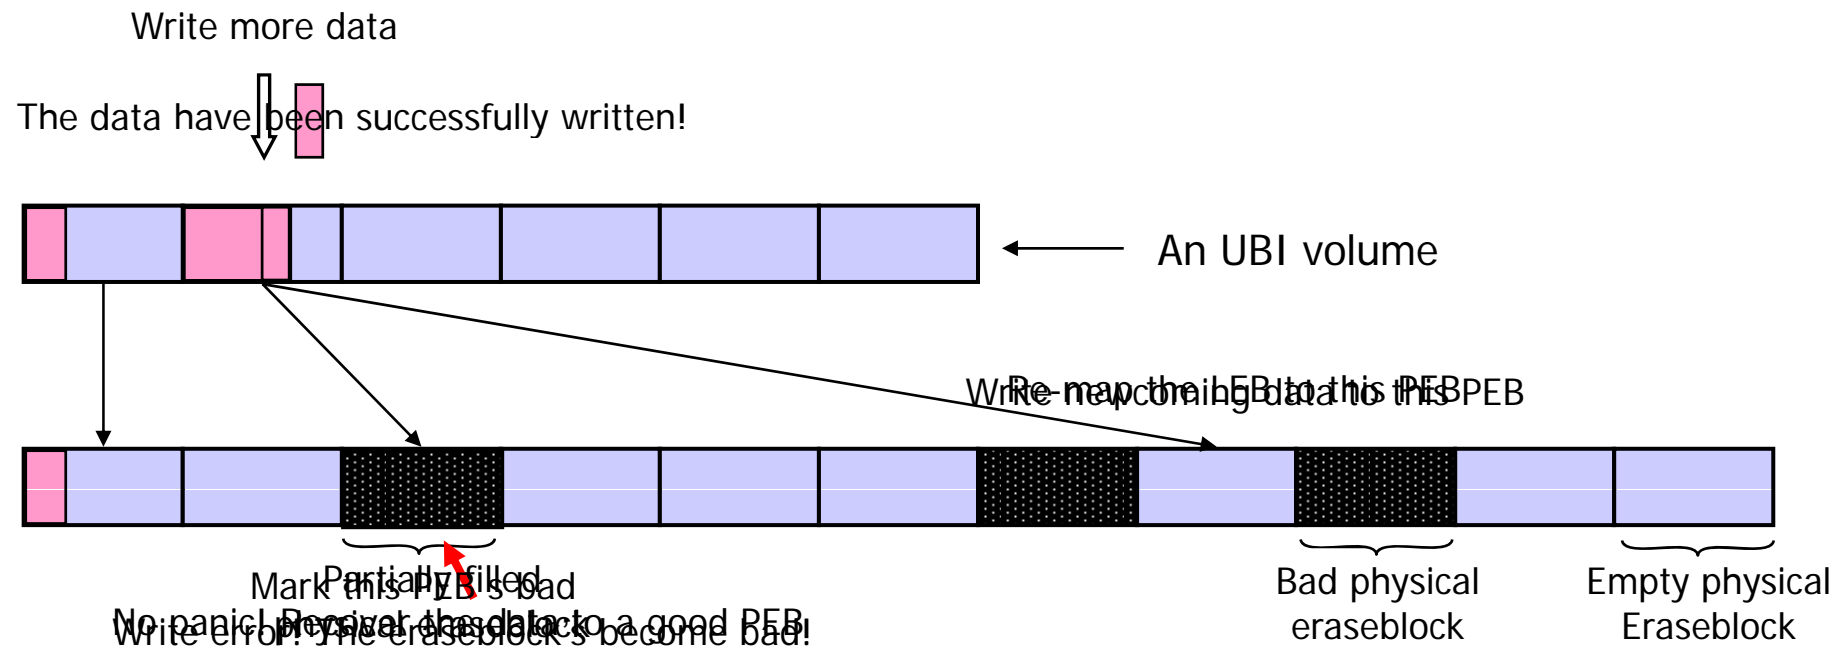
- LEBs are mapped to PEBs
- Any LEB may be mapped to any PEB



# UBI Internals (2)

## ■ Bad eraseblocks handling

- UBI volumes are devoid of bad eraseblocks
- UBI does proper error recovery transparently



# UBI Interfaces



## ■ UBI character devices

- UBI devices: `/dev/ubi0`, `/dev/ubi1`, ...
  - Volume create, delete, re-size, and get device description operations
- UBI volumes: `/dev/ubi0_0`, `/dev/ubi0_1`, ...
  - Read, write, update, and get volume description operations

## ■ UBI sysfs interface

- `/sys/class/ubi`

## ■ UBI in-kernel interface

- `Include/linux/mtd/ubi.h`