



CS632/SEP564: Embedded Operating Systems (Fall 2008)

Embedded Software Optimization

KAIST

Making Things Smaller



- Kernel
- Root filesystem
- Application

Making Things Smaller (1)

■ Downsizing the kernel

- Include only the necessary modules by a correct configuration
 - Device drivers, network stacks, file systems, etc.
- Remove unnecessary kernel functionalities
 - Turn off debugging and error logging flags
 - Remove support for `sysctl` (used to tweak kernel parameters at run time)
 - Get ~18KB by removing `msg*` (SysV) and `mq_*` (POSIX)
- Use the optimization flags when compiling the kernel
 - `-O1`, `-O2`, `-O3`, and `-Os`
 - Generally select `-O2` to achieve a balance between size and performance

Making Things Smaller (2)

■ Downsizing the kernel (cont'd)

- Decrease the size of the static buffer and array allocated in the kernel
 - Modify options in the menuconfig of the kernel
 - » The maximum number of supported peripherals, etc.
 - Use nm to identify the allocated size of each variable in an object file.
- Use strip to remove unnecessary symbols in the kernel and kernel modules
- Disable KALLSYMS which is used to print out symbolic crash information and symbolic stack backtraces

Making Things Smaller (3)

■ Linux-tiny patches

- http://elinux.org/Linux_Tiny
- A series of patches to reduce the memory and disk footprint of the kernel
- Originally started by Matt Mackall in Dec. 2003
- Stalled in late 2005 (Linux 2.6.14)
- Revived by Tim Bird (Sony) & CELF (Consumer Electronics Linux Forum) in 2007
- The latest release is for Linux 2.6.23
- Goals
 - ~300KB of a compressed kernel for simple configurations running comfortably with 2MB of RAM

Making Things Smaller (4)

■ Linux-tiny patches (cont'd)

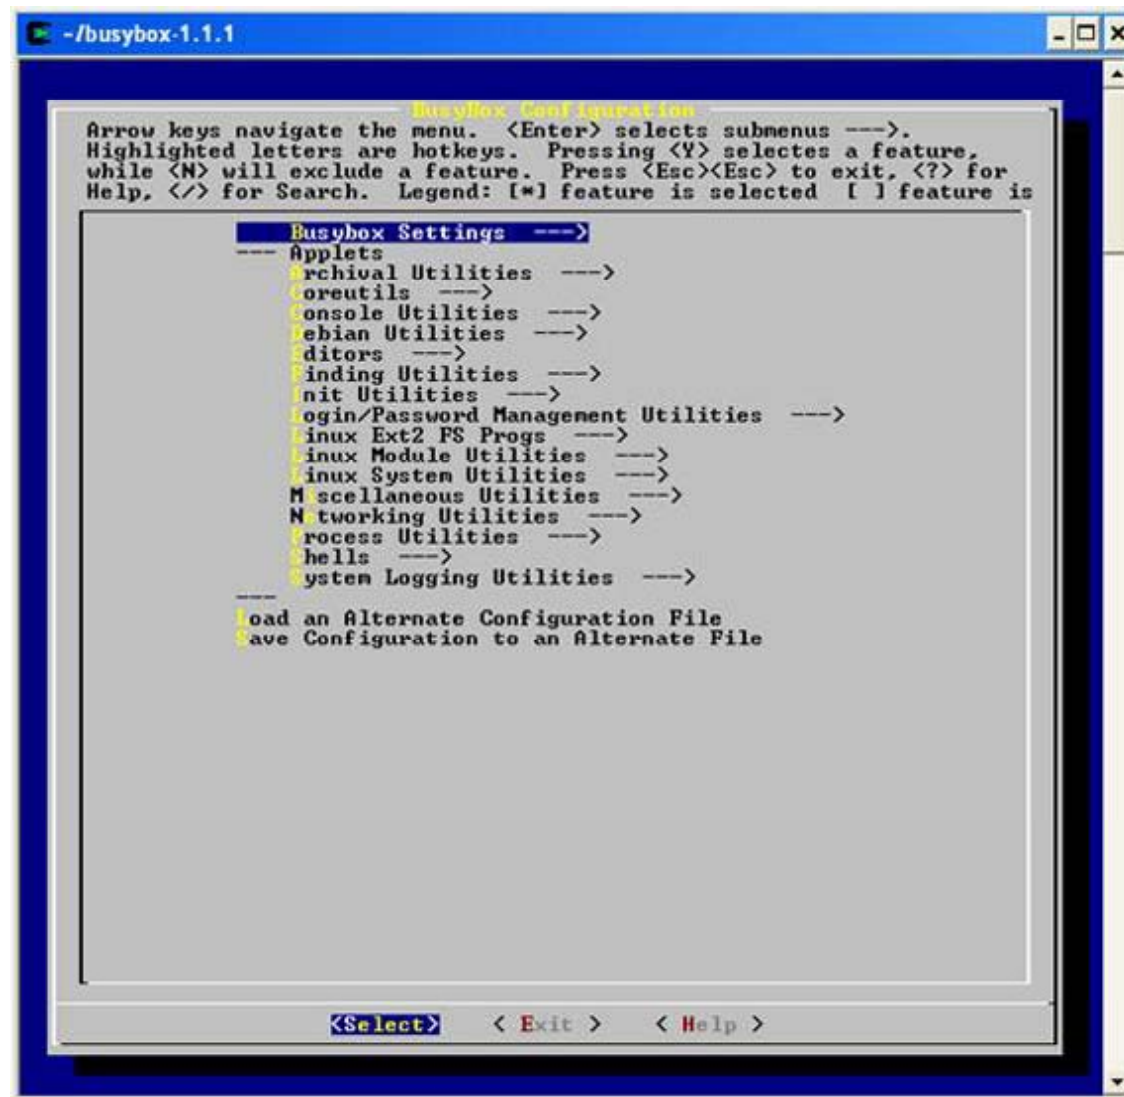
- Fine-grained control of printk() message compilation
 - ~ 60,000 printk() calls in Linux 2.6
 - All the static strings contribute to 5 ~ 10% of the kernel size
 - Control over the files that can use printk()
- Allow using function instead of table for CRC32 calculation
 - During packet processing in Ethernet driver
 - 2KB saving
- Network tweaking: buffer sizes, open sockets, ...
- Reduction of inlining
- Make some data structure sizes configurable, ...

Making Things Smaller (5)

■ Downsizing the root filesystem

- Use busybox (<http://busybox.net>)
 - BusyBox combines tiny versions of many common UNIX utilities into a single small executable file
 - Highly modular, allowing commands to be included and excluded at compile time
 - Each utility can be accessed by calling the single BusyBox binary with various names
 - » Supported by having a hard link or symbolic link
 - Reduces the overheads introduced by the executable file format (typically ELF)
 - Reduces internal fragmentation in memory and disk
 - Allows code to be shared between multiple applications without requiring a library

Making Things Smaller (6)



Making Things Smaller (7)

- **Downsizing the root filesystem (cont'd)**
 - Remove unnecessary libraries
 - Use ldd to identify the required shared libraries for each program
 - (cf.) readelf -d

```
jinsoo@csl: ~  
jinsoo@csl:~$ ldd /bin/ls  
linux-gate.so.1 => (0xb8002000)  
librt.so.1 => /lib/tls/i686/cmov/librt.so.1 (0xb7fd0000)  
libselinux.so.1 => /lib/libselinux.so.1 (0xb7fb6000)  
libacl.so.1 => /lib/libacl.so.1 (0xb7fad000)  
libc.so.6 => /lib/tls/i686/cmov/libc.so.6 (0xb7e4f000)  
libpthread.so.0 => /lib/tls/i686/cmov/libpthread.so.0 (0xb7e36000)  
/lib/ld-linux.so.2 (0xb7fe8000)  
libdl.so.2 => /lib/tls/i686/cmov/libdl.so.2 (0xb7e32000)  
libattr.so.1 => /lib/libattr.so.1 (0xb7e2d000)  
jinsoo@csl:~$ readelf -d /bin/ls |grep NEEDED  
0x00000001 (NEEDED) Shared library: [librt.so.1]  
0x00000001 (NEEDED) Shared library: [libselinux.so.1]  
0x00000001 (NEEDED) Shared library: [libacl.so.1]  
0x00000001 (NEEDED) Shared library: [libc.so.6]  
jinsoo@csl:~$
```

Making Things Smaller (8)

- **Downsizing the root filesystem (cont'd)**
 - Replace the standard C library with a small C library
 - uClibc, uClibc++, Newlib, or dietlibc
 - Such libraries remove the unused functions
 - Use a library optimizer Libopt to rebuild the libraries
 - Include only the necessary functions for the executable programs and shared libraries found in the root filesystem
 - Libopt utilizes objdump and nm to gather information about library object files, shared libraries and executable programs
 - Remove unnecessary commands and packages
 - Remove unnecessary directories and documents
 - /home, /mnt, /opt, /root, /boot, and /proc
 - Man, info, include, example directories

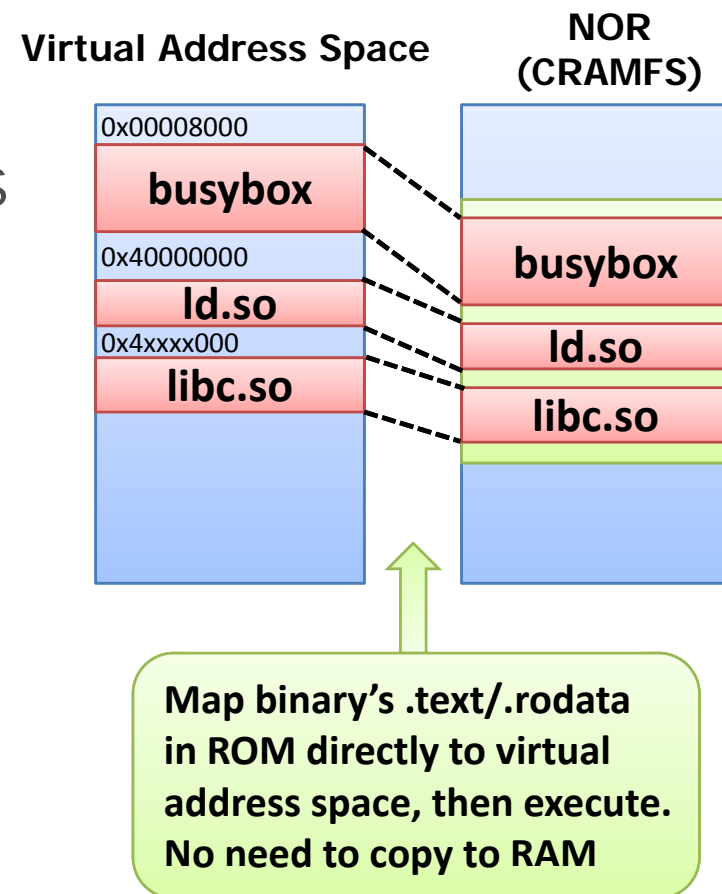
Making Things Smaller (9)

- **Downsizing the root filesystem (cont'd)**
 - Use compressed file system such as CRAMFS or SquashFS
 - Supports random read access with on-the-fly decompression
 - Eliminate unused code/data
 - CFLAGS += "-ffunction-sections -fdata-sections"
 - LDFLAGS += "-gc-sections"
 - Do not compile with -g
 - Use strip to remove all symbols
 - Compile with 16-bit instruction set (Thumb in ARM)

Making Things Smaller (10)

■ Reducing memory usage

- Application XIP (eXecute-In-Place)
 - Map code pages on NOR CRAMFS to Virtual Address Space directly.
 - Space vs. execution speed
 - Total ~45% reduction of page cache allocated to processes at stand by after bootup
 - The needed NOR flash size increases x 2 (Compression is omitted)



Making Things Smaller (11)

- **Reducing memory usage (cont'd)**
 - “Allocate on Write” for .data sections
 - Some pages in .data sections remain unwritten throughout process lifetime
 - The kernel allocates page cache to a process whenever it accesses pages in .data sections.
 - The page allocation can be deferred for .data until the process writes to the page
 - XIP + COW mechanisms
 - Total ~26% reduction of page cache allocated to processes at stand by idle state.
 - NOR flash access latency slows down process execution

Making Things Faster (1)



■ General bootup procedure

- Boot loader
 - System initialization
 - Kernel image loading from NOR/NAND/OneNAND Flash
- Kernel initialization
 - Kernel initialization
 - Device drivers initialization
 - Mount root filesystem
 - /sbin/init invocation
- Startup scripts (/etc/inittab, rc.*)
 - Load dynamic kernel modules
 - Mount filesystems
 - Run applications

Making Things Faster (2)



■ Boot loader

- Reduce the kernel image size
- Do not enable unnecessary devices (e.g., NIC)
- Compressed vs. uncompressed kernel image?
 - Depends on the flash speed, memory bandwidth, CPU performance, etc.
- Flash speed matters!
 - OneNAND synchronous burst mode?
- Kernel XIP
 - Allow kernel to be executed in-place in ROM or NOR Flash
- Use DMA for copying the kernel image
 - Save the time and CPU resources

Making Things Faster (3)



■ Kernel

- Preset loops_per_jiffy (LPJ)
 - Avoid a delay loop calculation (calibrate_delay())
- Avoid the RTC (Real-Time Clock) edge synchronization
- Disable console output
- Lazy device driver initialization
 - If a device driver is not necessary on system init, separate it as a kernel module
- Probe and initialize several devices in parallel
- Disable page-cache readahead option
- Bypass flash sector-mapping layer for read-only filesystems

Making Things Faster (4)



■ Startup scripts

- Insert kernel modules into the kernel
 - Reduce the time for kernel symbol resolution
 - Reduce the time for copying the entire image to memory
- Lazy loading for not-so-urgent kernel modules
 - e.g., USB (usbcore, usb-hub, usb-storage), networking, ...
- Remove unused services on init/rc.* script
- Run rc.* scripts in parallel
- Use static linking to avoid run-time overhead
 - Busybox, utilities, applications
- Directly call bootup commands inside busybox
- Compressed or uncompressed filesystem?

Making Things Faster (5)



■ Application

- Focus on the urgent things first
 - e.g., Screen & OSD display on DTV
 - Global constructors?
 - Lazy process/thread creation
- Preallocate memory used during bootup sequence
 - Avoid calling malloc()
- Use mmap() instead of read() for initial application data
 - Only the required pages are loaded on demand
 - More page faults, but less memory copy
 - e.g., Fonts, images, etc.

Making Things Faster (6)



■ Suspend/Resume

- Resume the snapshot image at boot time
- swsusp, Suspend2 (TuxOnIce), uswsusp,
 - Only support IA32, PPC, x86_64, IA64
- No official support for ARM and MIPS
 - (cf.) TP InstantBoot (<http://www.tripeaks.co.jp>)
- A snapshot image is created only once, and used repeatedly
- The snapshot image can be compressed
- The page cache can be populated in advance
- Resuming device states is tricky
- Who is resuming, kernel or boot loader?

Making Things Faster (7)

■ Dynamic loading of shared libraries

- A new application is loaded via `do_execve()`
- The type of executable is determined by parsing the ELF header
- For dynamically-linked executable, transfer the control to the dynamic linker (`ld.so`)
- `ld.so` loads all shared libraries needed by the application
- Relocate external references to shared libraries (Lazy binding of function calls)
 - Indirectly via GOT (Global Offset Table) and PLT (Procedure Linkage Table)
- Transfer the control to `libc` runtime

Making Things Faster (8)



■ Prelink

- Symbol lookup and relocation handling increases the application startup time significantly
- Originally developed for GUI programs with lots of shared libraries (OpenOffice, etc.)
- Prelinking steps
 - Identify executables to be prelinked and their dependent shared libraries
 - Assign non-overlapping virtual address space to each shared library
 - Apply the relocation
 - The dynamic linker is modified to handle prelinked executables/shared libraries

Making Things Faster (9)

■ Profile-driven functions reordering

- Profiling phase
 - CFLAGS += "-finstrument-functions"
 - `cyg_profile_func_enter()`: output the address of the called fn
- Analysis phase
 - Identify all the functions and libraries, if any, they belong to.
 - Generate a linker script
- Reordering phase
 - Relink the application with the new linker script

■ Profile-driven data reordering

- Data sections: `.rodata`, `.data`, `.bss`
- How can we identify data touched during startup?

Making Things Faster (10)

■ Profile-driven function/data reordering

- Reduces the number of pages faults
- Reduces the number of pages read from flash
- Better locality in TLB and cache
- Reliable as is done by the standard linker

- Requires the instrumented version of objects/libraries
- May require full source code and recompilation
- Standard libraries? System call stubs?
 - libc, libm, libpthread, libstdc++, libgcc, libgcc_eh, ...
- Data access profiling is not precise

Thank You!



■ Final exam.

- December 16, 2:30pm – 4:00pm
 - CS Building #4448
 - Scope: everything
 - Closed-book exam.
-
- Count as 30% of your final grade!
 - Submit your final project report on time. (another 70%)