

CS 6V81-05: System Security and Malicious Code Analysis Understanding the Implementation of Virtual Memory

Zhiqiang Lin

Department of Computer Science
University of Texas at Dallas

February 29th, 2012

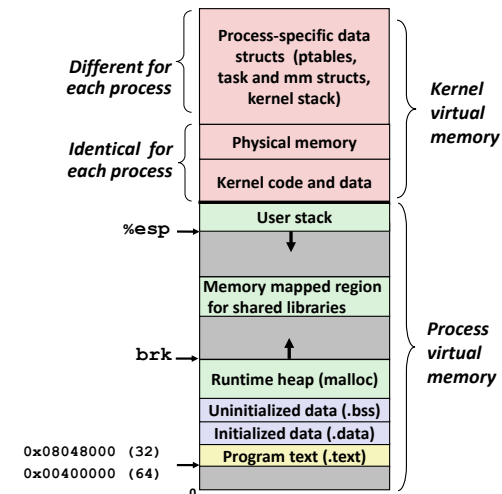
Outline

- 1 Overview
- 2 Implementation
- 3 Summary

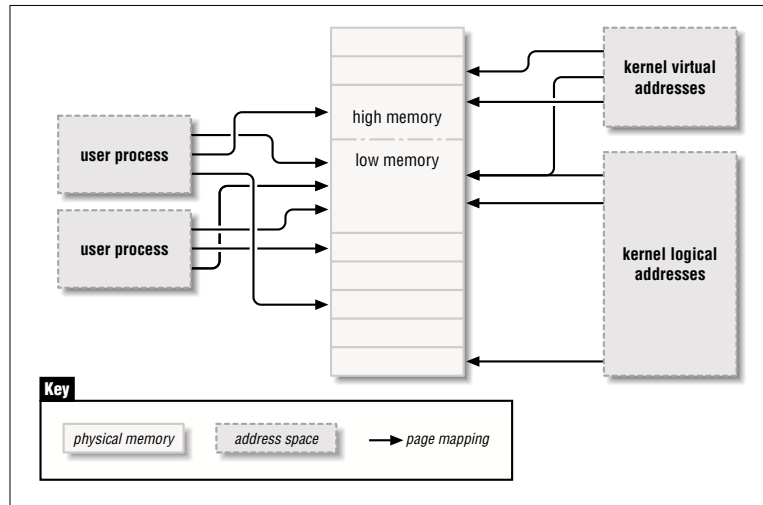
Outline

- 1 Overview
- 2 Implementation
- 3 Summary

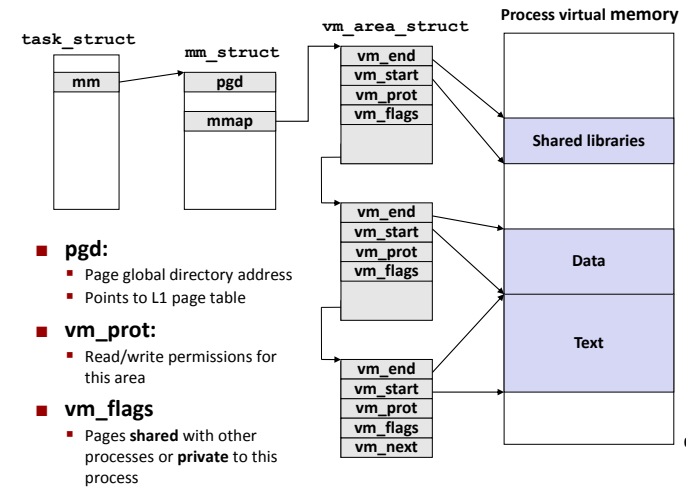
Virtual Memory of a Linux Process



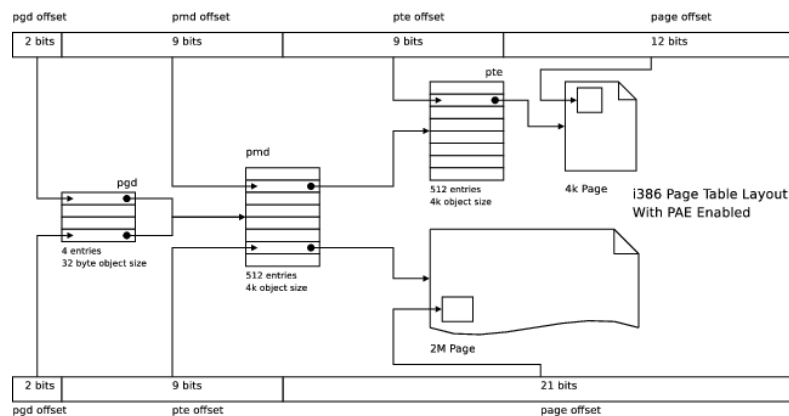
Address type in Linux



Linux Organizes VM as Collection of “Areas”

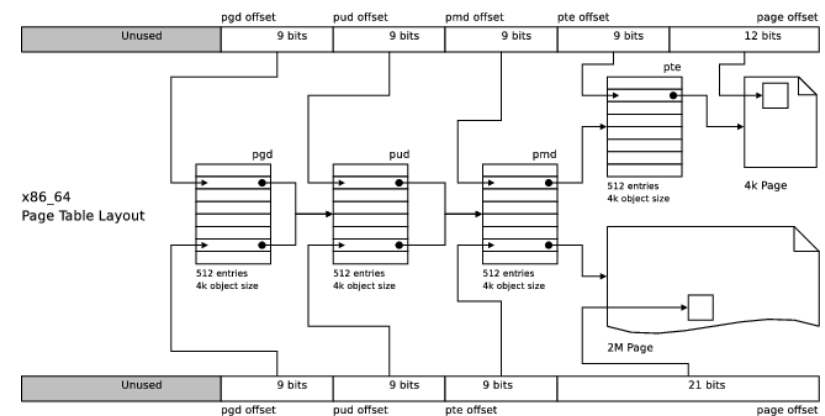


i386+PAE



sources: <http://linux-mm.org/PageTableStructure>

x86_64



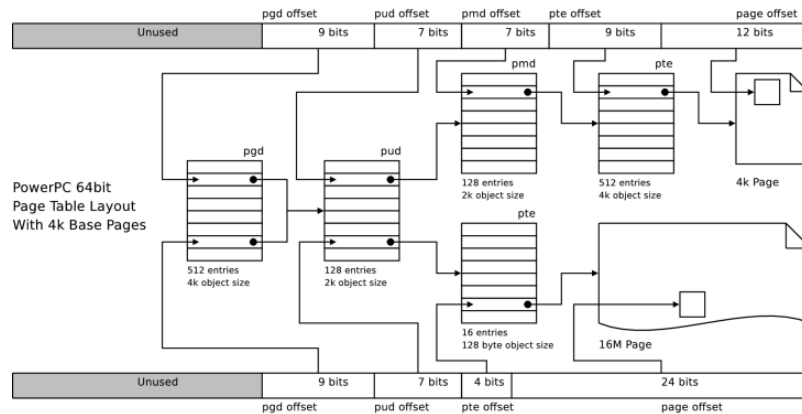
sources: <http://linux-mm.org/PageTableStructure>

Overview
○○○○○●○○○○○○○○

Implementation
○○○○○○○○○○○○

Summary
○○

Powerpc-4k



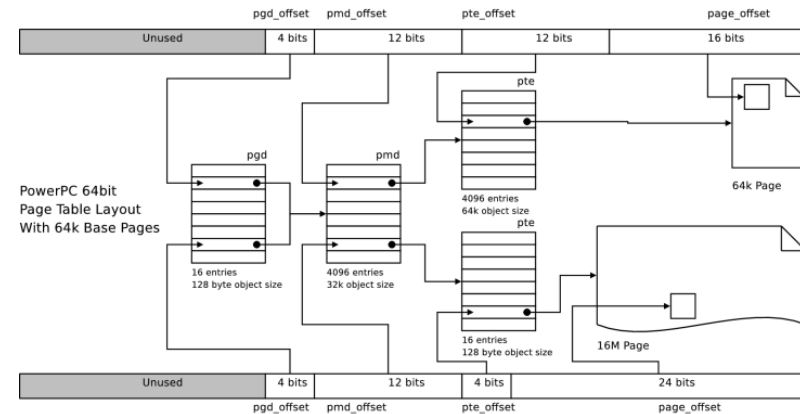
sources: <http://linux-mm.org/PageTableStructure>

Overview
○○○○○●○○○○○○○○

Implementation
○○○○○○○○○○○○

Summary
○○

Powerpc-64k




sources: <http://linux-mm.org/PageTableStructure>

Overview
○○○○○●○○○○○○○○

Implementation
○○○○○○○○○○○○

Summary
○○

Simplifying Linking and Loading



Segmentation fault:

accessing a non-existing page

Normal page fault

Overview
○○○○○●○○○○○○○○


Implementation
○○○○○○○○○○○○

Summary
○○

Demand paging


- **Key point:** no virtual pages are copied into physical memory until they are referenced!
 - Known as **demand paging**
- Crucial for time and space efficiency

Shared Objects



- Process 1 maps the shared object.
- Notice how the virtual addresses can be different.

Private Copy-on-write (COW) Objects

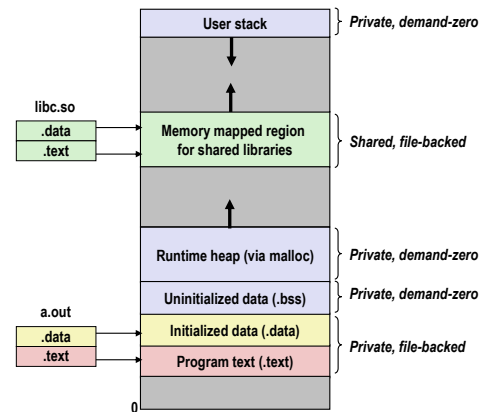


- Two processes mapping a **private copy-on-write (COW)** object.
- Area mapped as private copy-on-write.
- PTEs in private areas are flagged as read-only.
- Instruction writing to private page triggers protection fault.
- Handler creates new R/W page.
- Instruction restarts upon handler return.
- Copying deferred as long as possible!

Demand paging

- VM and memory mapping explain how fork provides private address space for each process.
- To create virtual address for new process
 - Create exact copies of current mm_struct, vm_area_struct, and page tables.
 - Flag each page in both processes as read-only
 - Flag each vm_area_struct in both processes as private COW
- On return, each process has exact copy of virtual memory
- Subsequent writes create new pages using COW mechanism.

The execve Function



- To load and run a new program a.out in the current process using execve:
- Free vm_area_struct's and page tables for old areas
- Create vm_area_struct's and page tables for new areas
 - Programs and initialized data backed by object files.
 - .bss and stack backed by anonymous files.
- Set PC to entry point in .text
 - Linux will fault in code and data pages as needed.


User-Level Memory Mapping

```
void *mmap(void *start, int len, int prot, int flags, int fd, int offset)
```

- Map len bytes starting at the offset of the file specified by file description fd, preferably at address start
 - **start:** may be 0 for “pick an address”
 - **prot:** PROT_READ, PROT_WRITE, ...
 - **flags:** MAP_ANON, MAP_PRIVATE, MAP_SHARED, ...
- Return a pointer to start of mapped area (may not be start)

User-Level Memory Mapping

```
void *mmap(void *start, int len, int prot, int flags, int fd, int offset)
```

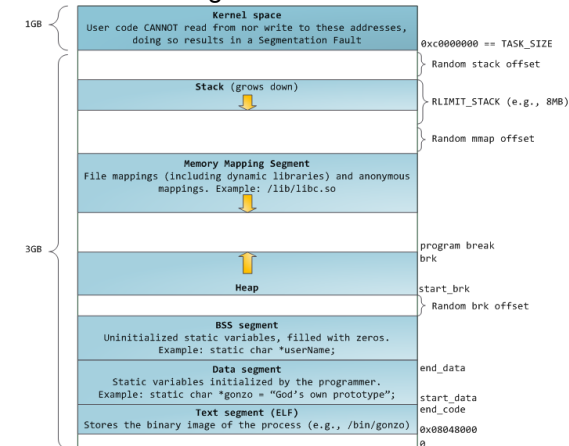


Outline

- 1 Overview
- 2 Implementation
- 3 Summary

Memory Management

Memory management is the heart of operating systems; Each process in a multi-tasking OS runs in its virtual address space.

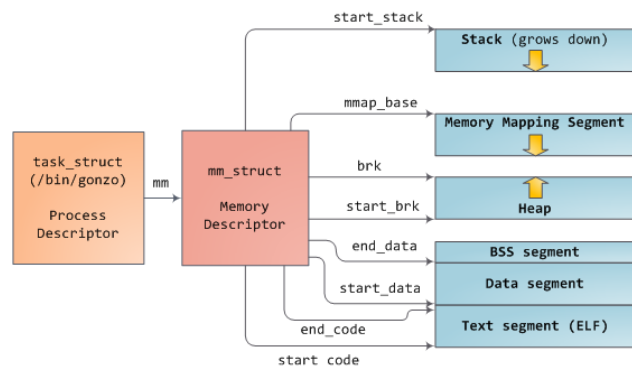


Overview
○○○○○○○○○○○○○○○○

Implementation
●○○○○○○○○○○○○

Summary
○○

How The Kernel Manages Process Memory



Credit: <http://duartes.org/gustavo/blog/post/how-the-kernel-manages-your-memory>

Overview
○○○○○○○○○○○○○○○○

Implementation
○○●○○○○○○○○○○

Summary
○○

mm_struct

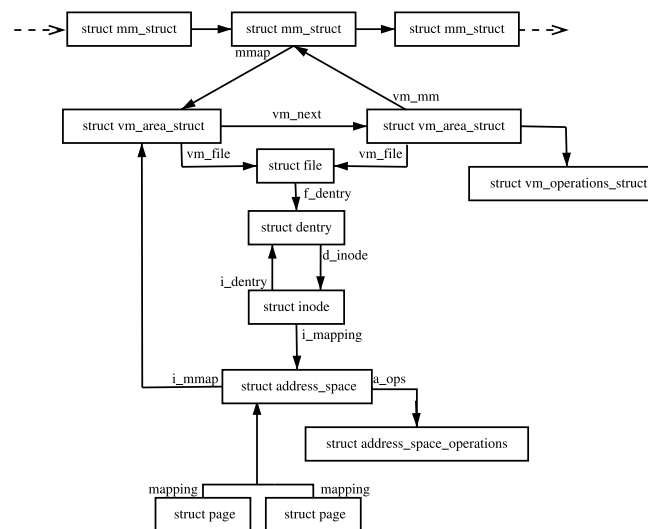
```
struct mm_struct {
    [0] struct vm_area_struct *mmap;
    [4] struct rb_root mm_rb;
    [8] struct vm_area_struct *mmap_cache;
    [12] long unsigned int (*get_unmapped_area)(struct file *, long uns
    [16] void (*unmap_area)(struct mm_struct *, long unsigned int);
    [20] long unsigned int mmap_base;
    [24] long unsigned int task_size;
    [28] long unsigned int cached_hole_size;
    [32] long unsigned int free_area_cache;
    [36] pgd_t *pgd;
    [40] atomic_t mm_users;
    [44] atomic_t mm_count;
    [48] int map_count;
    [52] struct rw_semaphore mmap_sem;
    [80] spinlock_t page_table_lock;
    [96] struct list_head mmlist;
    [104] mm_counter_t _file_rss;
    [108] mm_counter_t _anon_rss;
    [112] long unsigned int hiwater_rss;
    [116] long unsigned int hiwater_vm;
```

Overview
○○○○○○○○○○○○○○○○

Implementation
○○●○○○○○○○○○○

Summary
○○

Relevant data structures in address space



Overview
○○○○○○○○○○○○○○○○

Implementation
○○○○●○○○○○○○○

Summary
○○

mm_struct

```
struct mm_struct {
    [120] long unsigned int total_vm;
    [124] long unsigned int locked_vm;
    [128] long unsigned int shared_vm;
    [132] long unsigned int exec_vm;
    [136] long unsigned int stack_vm;
    [140] long unsigned int reserved_vm;
    [144] long unsigned int def_flags;
    [148] long unsigned int nr_ptes;
    [152] long unsigned int start_code;
    [156] long unsigned int end_code;
    [160] long unsigned int start_data;
    [164] long unsigned int end_data;
    [168] long unsigned int start_brk;
    [172] long unsigned int brk;
    [176] long unsigned int start_stack;
    [180] long unsigned int arg_start;
    [184] long unsigned int arg_end;
    [188] long unsigned int env_start;
    [192] long unsigned int env_end;
```

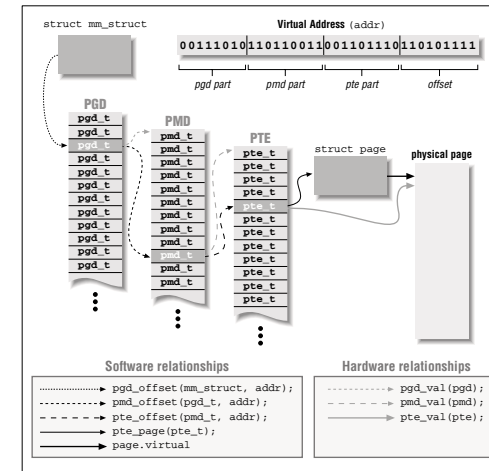
mm_struct

```

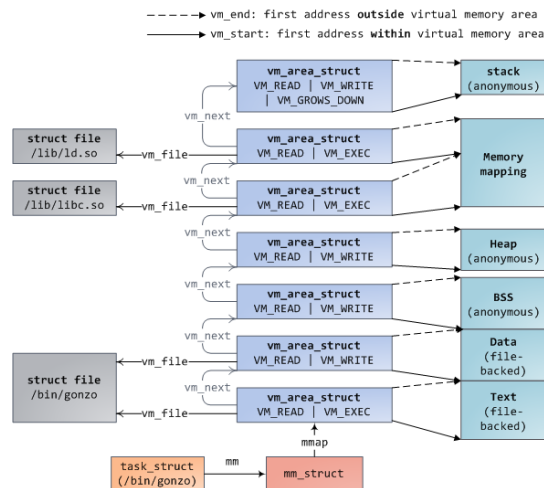
struct mm_struct {
[196] long unsigned int saved_auxv[44];
[372] unsigned int dumpable : 2;
[376] cpumask_t cpu_vm_mask
[380] mm_context_t context;
[424] long unsigned int swap_token_time;
[428] char recent_pagein;
[432] int core_waiters;
[436] struct completion *core_startup_done;
[440] struct completion core_done;
[468] rwlock_t ioctx_list_lock;
[484] struct kiocx *iocx_list;
}
SIZE: 488

```

Page Directory



vm_area_struct



Credit: <http://duartes.org/gustavo/blog/post/how-the-kernel-manages-your-memory>

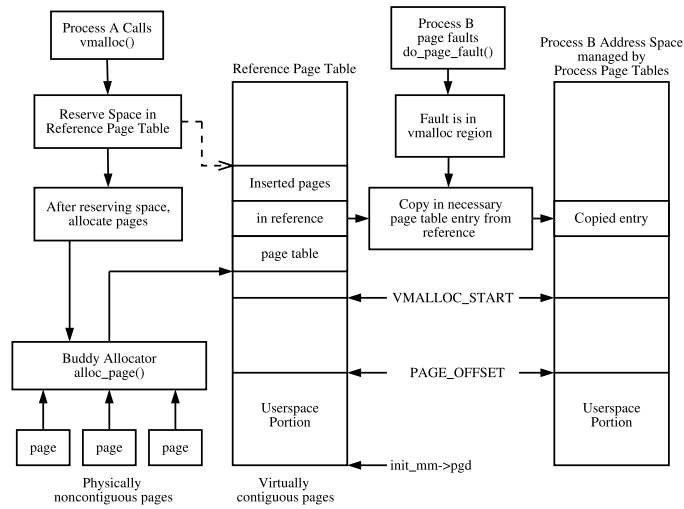
vm_area_struct

```

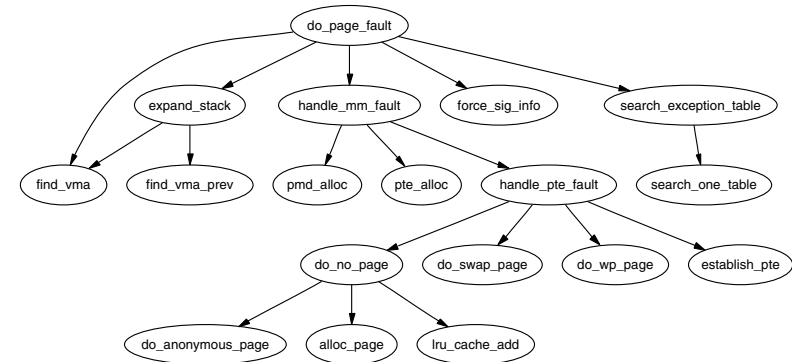
struct vm_area_struct {
[0] struct mm_struct *vm_mm;
[4] long unsigned int vm_start;
[8] long unsigned int vm_end;
[12] struct vm_area_struct *vm_next;
[16] pgprot_t vm_page_prot;
[20] long unsigned int vm_flags;
[24] struct rb_node vm_rb;
    union {
        struct {...} vm_set;
        struct raw_prio_tree_node prio_tree_node;
    }
[36] } shared;
[52] struct list_head anon_vma_node;
[60] struct anon_vma *anon_vma;
[64] struct vm_operations_struct *vm_ops;
[68] long unsigned int vm_pgoff;
[72] struct file *vm_file;
[76] void *vm_private_data;
[80] long unsigned int vm_truncate_count;
}
SIZE: 84

```

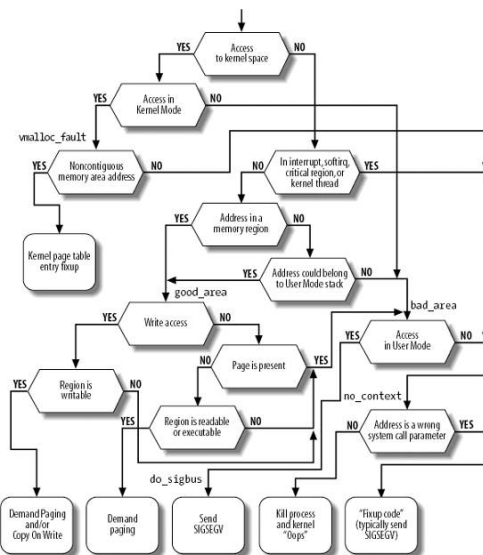
Page Fault



do_page_fault



do_page_fault



Outline

- 1 Overview
- 2 Implementation
- 3 Summary

Summary

- Memory management is crucial
 - Machine introspection
 - Memory forensics
 - Traversing kernel data structures to understand the memory
- Memory failures
- Exploits

References

- Dening BEFORE MEMORY WAS VIRTUAL
<http://cs.gmu.edu/cne/pjd/PUBS/bvm.pdf>
- Linux device drivers, Addison-wisely.
- <http://www.cs.cmu.edu/afs/cs/academic/class/15213-f10/www/lectures/16-vm-systems.pdf>
- Understanding Linux virtual memory manager
<http://www.kernel.org/doc/gorman/pdf/understand.pdf>
- Understanding Linux kernel (3rd edition)
- <http://linux-mm.org/PageTableStructure>